NP-Hardness and Additive Inapproximability of Deterministic Communication Complexity

Serge Gaspers* Tao Zixu He† Simon Mackenzie*

Abstract

Given the full truth table of a Boolean function whose input is split between two players, we study the complexity of determining the minimum worst-case number of bits they must communicate to compute it. This is a basic meta-complexity question, posed by Yao in 1979.

We prove NP-hardness of approximation within every fixed additive error for deterministic communication complexity, under deterministic polynomial-time reductions. The result holds both in the standard model and in Yao's original model, where the players alternate after every bit and only the receiver of the last bit must determine the answer. In particular, it resolves the exact computation question under Yao's original conventions. Earlier versions of this same paper established exact NP-hardness for the standard model, independently and concurrently with Hirahara, Ilango, and Loff.

In both models, a polynomial-time estimator with any fixed additive error would yield a polynomial-time 3-SAT algorithm, deterministic or bounded-error randomized according to the estimator. For the standard model, we also give quantitative implications for errors that grow with the number T of table entries. A bounded-error randomized polynomial-time estimator with additive error $c \log T / \log \log T$, for any fixed $0 \leq c < 1/8$, yields a bounded-error randomized 3-SAT algorithm on N variables and $O(N)$ clauses running in time $2^{O(N^\alpha)}$ for every fixed $8c < \alpha < 1$. A deterministic polynomial-time estimator with additive error $c_* \log \log T$, for an effective absolute constant $c_* > 0$, yields a deterministic algorithm running in time $2^{O((\log N)^2)}$ on such formulas.

*University of New South Wales
†University of Massachusetts Amherst

1 Introduction

Alice and Bob know a Boolean matrix M. Alice receives a row index and Bob receives a column index. They must determine the entry at their intersection, communicating as few bits as possible. Their local computation is free, and they must answer correctly for every pair of indices. The deterministic communication complexity $\mathrm{D}(M)$ is the smallest worst-case number of bits used by any such protocol.

We study how hard it is to estimate $\mathrm{D}(M)$ when the algorithm receives the whole table as input. The table is total, meaning every entry is specified, and we write T for its number of entries. An estimator runs in time measured in T and outputs a number, not a protocol. A randomized estimator must meet its error guarantee with probability at least $2/3$ on every table, while the protocols whose cost it estimates remain deterministic. Yao posed the problem of computing $\mathrm{D}(M)$ exactly from the table when he introduced communication complexity in 1979 [22].

We prove that computing deterministic communication complexity is NP-hard, and that this remains true when any fixed additive error is allowed. The exact NP-hardness result for the standard model was first presented in earlier versions of this work [6], independently and concurrently with Hirahara, Ilango and Loff [10].

We also prove exact and fixed-error NP-hardness in Yao's original model: the players alternate after every bit, and only the receiver of the last bit needs to determine the answer. The earlier exact-hardness results did not cover these speaking and output conventions. Our alternating construction therefore resolves Yao's question in the model in which it was posed.

The challenge of approximation. In the standard model, the exact-hardness reduction maps a 3-CNF formula to a table whose optimal cost is one bit larger if the formula is unsatisfiable than if it is satisfiable. An estimate within additive error a distinguishes the two cases only if their costs differ by more than $2a$. A natural way to enlarge the gap is to combine several such tables and hope that their one-bit gaps add up. However, a protocol for the combined problem may share communication among them, so their costs need not add, and the lower bound must hold against every such protocol. Combining problems also enlarges the truth table. Since the reduction writes out every entry, the table must have size polynomial in the size of the formula for each fixed gap. Given k, our reductions compute a table and a number B such that the optimal cost is $\leq B$ if the formula is satisfiable and $\geq B + k$ otherwise. We call this an *additive gap* of k bits. We do not establish hardness of approximation within any fixed multiplicative factor above one.

Two rules for speaking. In an *unrestricted* protocol, the bits sent so far determine who speaks next, and the same player may send several consecutive bits. In a *strictly alternating* protocol, the players take turns after every bit. Where an unrestricted protocol lets one player send two consecutive bits, an alternating protocol must let the other player speak in between, so the two optimum costs can differ.

Interlacing. Interlacing combines q copies of a matrix M through a pointer. In the q-fold interlace $\mathsf{I}_q(M)$, Alice receives a row x of M together with an index $i \in \{1, \ldots, q\}$, and Bob receives a list $(y_1, \ldots, y_q)$ of q columns of M, possibly with repetitions. The required output is $M(x, y_i)$. Alice's index acts as a one-step pointer into Bob's input, as in pointer chasing [19]. Figure 1 shows the case $q = 2$ for a 2×2 matrix.

There is a natural protocol: Alice sends i using $\log q$ bits, and the players then run an optimal protocol for M on x and y_i, for a total of $\mathrm{D}(M) + \log q$ bits. But is communicating the index an

unavoidable extra cost? Perhaps Bob could speak first, or the players could combine information about the index and the answer, saving bits.

The interlacing framework rules out such savings when M is robustly hard, that is, when every large subtable of M is nearly as hard as M itself. Then any protocol that does not start by revealing Alice's index leaves a subproblem that is still too hard for the bits that remain, so the natural protocol is optimal [6, 18].

The interlacing method for lower bounds was developed by Mackenzie and Saffidine [18] to refute the exact direct-sum conjecture, and underlies our original NP-hardness reduction [6]. If M has n columns, Bob has n^q possible lists. To keep its table polynomial, that reduction already replaced them by a polynomial-size family of lists. Repeating the construction needs more: the table produced by one round must again have lower bounds on its large subtables, so that the next round can use it as the base matrix M. An intermediate version obtained additive inapproximability under ETH using subexponential-time constructions [7].

Polynomial-size gap amplification. Our deterministic construction keeps explicit input sets of polynomial size at every amplification step, so every fixed gap is constructed in deterministic polynomial time. Our randomized construction instead keeps all tuples during the amplification. The resulting table is too large to print, but each entry can be computed from short descriptions of its row and column. At the end, the construction samples rows and columns of this table. The table has low rank, so short protocols can be correct on only few sets of columns, and a union bound over these sets shows that the sample preserves the gap with high probability (Theorem 9.4). In both constructions, the lower bounds for unrestricted protocols come from interlacing. Because the two constructions control the table size differently, they give different tradeoffs between the approximation error and the running time of the 3-SAT algorithm that an estimator would imply.

Strict alternation needs a separate construction. The short protocols in our unrestricted construction send multi-bit messages. Making them alternate adds turns, and these extra bits can exceed the gap. We therefore construct a different table, prove its lower bound using the fixed order of speaking, and again keep only a small explicit set of Bob's inputs.

Output and speaking conventions. By default, the conversation alone must determine the answer, which we call *public output*. We write $D_A^{\mathrm{alt}}(M)$ for the minimum cost with public output and strict alternation, with Alice starting. Yao's original convention instead allows the receiver of the last bit to use their own input as well as the conversation to determine the answer. We denote its optimum by $Y_{\mathrm{free}}(M)$, where the protocol may choose its starting player independently of the inputs. In both conventions, the bits exchanged must determine when the conversation has ended.

From approximation to satisfiability algorithms. For errors growing with the table size, we prove that an estimator for communication complexity would yield a 3-SAT algorithm with an explicitly bounded running time. If one assumes that such a 3-SAT algorithm does not exist, each implication becomes an inapproximability result. The fixed-error NP-hardness results need no such assumption.

1.1 Results

The following table summarizes the fixed-error hardness results and two examples of our general algorithmic implication. The full tradeoff is stated in Theorem 9.7.

Protocol model	Additive error	Implied 3-SAT algorithm	Algorithm type
Unrestricted	Every fixed error	Polynomial time	Same as estimator
Strict alternation	Every fixed error	Polynomial time	Same as estimator
Unrestricted	$c_* \log\log T$	$2^{O((\log N)^2)}$	Same as estimator
Unrestricted	$c \log T / \log\log T$, $c < 1/8$	$2^{O(N^\alpha)}$, every $8c < \alpha < 1$	Randomized

The last two rows concern formulas with N variables and at most $c_s N$ clauses, for any fixed density bound c_s. Each row describes the algorithm that a polynomial-time estimator would imply, and applies to both deterministic and randomized estimators. All randomized algorithms here have bounded error.

In each reduction below, Φ is a 3-CNF formula. The YES (satisfiable) case is $\Phi \in \mathsf{SAT}$, and the NO case is $\Phi \notin \mathsf{SAT}$. The reduction computes the threshold B from the formula and the requested gap, without deciding satisfiability.

Theorem 1.1 (Fixed gaps for unrestricted protocols)**.** *For every fixed integer $k \geq 1$, a deterministic polynomial-time algorithm maps a 3-CNF formula Φ to an explicit total Boolean table M and an integer B such that*

$$\Phi \in \mathsf{SAT} \implies \mathrm{D}(M) \leq B, \qquad \Phi \notin \mathsf{SAT} \implies \mathrm{D}(M) \geq B + k.$$

The polynomial may depend on k. The table can be taken square with power-of-two side length.

A different construction gives the same fixed-gap hardness under strict alternation.

Theorem 1.2 (Fixed gaps under strict alternation)**.** *For every fixed integer $k \geq 1$, a deterministic polynomial-time algorithm maps a 3-CNF formula Φ to an explicit total Boolean table M and an integer B such that*

$$\Phi \in \mathsf{SAT} \implies D_A^{\mathrm{alt}}(M) \leq B, \qquad \Phi \notin \mathsf{SAT} \implies D_A^{\mathrm{alt}}(M) \geq B + k.$$

The polynomial bounds on output size and running time may depend on k.

For additive error a, choose $k > 2a$ and compare the estimate with $B + k/2$. Thus, for every fixed error, estimating either cost in the preceding two theorems in deterministic polynomial time implies $\mathsf{P} = \mathsf{NP}$, and bounded-error randomized estimation implies $\mathsf{NP} \subseteq \mathsf{BPP}$. Changing the starting player increases the optimal cost by at most one bit, and so does requiring the receiver of the last bit to announce the answer (Lemma 10.1). By increasing the fixed gap, we therefore also obtain hardness under Yao's output convention.

Corollary 1.3 (Yao's output convention)**.** *For every fixed nonnegative additive error, a deterministic polynomial-time estimator of $Y_{\mathrm{free}}(M)$ on all explicit total Boolean tables implies $\mathsf{P} = \mathsf{NP}$, and a bounded-error randomized one implies $\mathsf{NP} \subseteq \mathsf{BPP}$. The same conclusions hold for estimating the public-output alternating cost with Alice first, Bob first, or a starter chosen independently of the inputs.*

Taking error zero in Corollary 1.3 answers Yao's original question by proving NP-hardness under his speaking and output rules. We do not establish a growing-error bound for the alternating model.

For unrestricted protocols, choosing the gap as a function of the formula size gives an explicit tradeoff between approximation error and the running time of the resulting 3-SAT algorithm (Theorem 9.7). Here N counts variables in the formula and T counts entries in the communication table. The following two consequences illustrate the randomized and deterministic regimes. The constant $1/8$ uses the sharper size bound for the randomized construction.

Corollary 1.4 (Randomized estimation and 3-SAT running time)**.** *Fix $0 \le c < 1/8$. Suppose a bounded-error randomized polynomial-time estimator approximates $\mathrm{D}(M)$ within additive error $c \log T / \log\log T$ on every sufficiently large explicit total Boolean table with T entries. Then, for every fixed $8c < \alpha < 1$ and every fixed c_s, sparse 3-SAT on N variables and at most $c_s N$ clauses has a bounded-error randomized algorithm with worst-case running time $2^{O(N^\alpha)}$. It suffices for the estimator to work on square tables whose side length is a power of two.*

For the smaller error scale below, the deterministic construction turns a deterministic estimator into a quasipolynomial-time algorithm for sparse 3-SAT.

Corollary 1.5 (Deterministic estimation and 3-SAT running time)**.** *There is an effective absolute constant $c_* > 0$ with the following property. If a deterministic polynomial-time estimator approximates $\mathrm{D}(M)$ within additive error $c_* \log\log T$ on every sufficiently large explicit total Boolean table, then, for every fixed c_s, 3-SAT on N variables and at most $c_s N$ clauses has a deterministic algorithm running in time $2^{O((\log N)^2)}$. It suffices for the estimator to work on square tables whose side length is a power of two.*

Section 9.4 gives the general tradeoff and explains how to choose the gap for other running-time bounds.

The proofs of Theorem 1.1 and Corollary 1.5 appear in Section 8.2. We prove Theorem 9.7 and Corollary 1.4 in Section 9.4, and Theorem 1.2 and Corollary 1.3 together at the end of Section 10. Section 2 explains the ideas behind these proofs.

1.2 Formal verification and proof dependencies

A companion Lean development provides machine-checked proofs of our constructions and their consequences for communication complexity and 3-SAT running time [13]. The formal statements retain explicit hypotheses for the cited source reductions. The public proof inspector [14] links manuscript statements to the corresponding Lean declarations and exposes their definitions and hypotheses. The repository [13] provides the source, build instructions, and verification records.

The earlier direct-sum construction of Mackenzie and Saffidine [18] has separately been formalized in Lean by Soltani Moakhar, Gholami, Springer, JafariRaviz, and Hajiaghayi [21], with statements and proofs available in a public inspector [2].

OpenAI GPT models used through ChatGPT Pro and Codex, and Anthropic's Claude models, assisted with proof development, correctness checking, writing, and the subsequent Lean formalization. The AI methodology statement describes their roles.

2 Technical overview

Our reduction encodes a graph so that failure of 3-coloring forces one extra bit of communication. A proper 3-coloring gives a protocol of an explicit cost $C - 1$ (Proposition 6.4). Conversely, interlacing forces every protocol of that cost to assign Alice's inputs to bins, in effect choosing a coloring, before Bob reveals which edge is tested. In the NO case, some tested edge is monochromatic, which forces an additional bit. The difficulty is to prove this lower bound against every protocol, and then to repeat the construction without losing the lower bound or polynomial size.

We explain three steps: the one-bit gap, the lower bounds under input deletion that make the gap repeatable, and the explicit short input lists that make every fixed number of rounds a polynomial-time reduction. Strict alternation requires a different encoding, described last.

2.1 Constructing the source matrix

Our source reduction maps a formula to a graph of maximum degree $\Delta = O(1)$ that is 3-colorable if the formula is satisfiable (the YES case), and otherwise leaves at least a constant fraction $\gamma > 0$ of its edges monochromatic under every assignment of three colors to its vertices (the NO case, Theorem 5.2). This NO condition survives small deletions of vertices. The graph has no isolated vertices, so $|E| \geq |V|/2$, and deleting fewer than $\gamma|V|/(4\Delta)$ vertices destroys fewer than $\gamma|E|/2$ edges, so at least $\gamma|E|/2$ monochromatic edges remain.

The construction takes a base communication matrix R of cost $d = \mathrm{D}(R)$ and adds rows encoding the graph, described below, to interlaced copies of R. In the first round, R is a simple matrix of known cost (Lemma 5.4). We allow a general base matrix so that the construction can be repeated, with the output of one round as the base matrix of the next. We also need lower bounds that survive deletion of some inputs. We measure the retained rows and columns using fixed positive rational weights, summing to one on each side. These act as fractional counts, which let us control how much each input family contributes, independently of its number of labels. Later, adjusting weights also lets us shorten tuple lists while keeping the weights induced on small sets of coordinates (Section 7). The weights do not change communication cost, since every protocol must still compute every entry correctly. The deterministic construction never samples from them, and its final output is an ordinary Boolean table with the weights discarded.

A proper 3-coloring places the vertices in three groups so that the endpoints of every edge lie in different groups. We call these groups *bins*. The interlacing construction uses a power-of-two number $s \geq 4$ of bins. We treat each graph vertex as an *item* to be placed in a bin, and add $s - 3$ auxiliary items that occupy the extra bins, so that the vertices cannot use them as extra colors. At an edge, its two endpoints and the auxiliary items are *active*, and all other items are *inactive*. We require that each bin contain at most one active item at every edge. The auxiliary items are active at every edge, so no two of them share a bin. Since every vertex of our source graph lies on an edge, no vertex can share a bin with an auxiliary item, and the vertices share the remaining three bins, exactly as in 3-coloring.

Each item i gives Alice one possible input (i, u) for every row u of R. A coloring partitions these inputs by the bin of their item. Naming the bin need not identify Alice's item or her row of R.

The full construction uses $4s$ labeled groups, each containing a copy of all items, so that the proof can discard damaged groups. For each edge in each group, a binary vector over all items marks the active items of that group with 1 and every other item with 0. We index these vectors by α. In the NO case, every assignment of items to bins has a constant fraction of indices at which one bin contains two active items and a third item that is inactive there, even after deleting a small fraction of the items (Lemma 6.3). We say there is a *collision* at such an index. The inactive third

item can come from another group. There are $4s$ groups but only s bins, and at most s groups can have a bin containing items of no other group. Hence for most groups the bin holding an active pair also holds an item of another group, which is inactive at that index because each index activates items of one group only. The proof extracts an assignment from any short protocol.

Alice's inputs are of two kinds. An *item row* consists of an item i and a row u of the base matrix R. The other inputs, called *template rows*, are grouped into s bins and are obtained from R by nested interlaces, described below. Their purpose is to force every protocol of cost at most $C-1$ to begin with $\log s$ consecutive bits from Alice (Theorem A.3). They name a bin before Bob reveals any index α.

Bob's columns have two families. The first supplies a separate index and list of base columns for each bin, together with a list used by the item rows, and enforces the initial $\log s$ bits. The second uses the same index α and the same list of base columns in every bin. At this index, the template rows use $s-2$ positions of the list, active items use distinct additional positions, and inactive items use one further common position. Thus two active items and one inactive item in the same bin expose $s+1$ positions.

Defining the table without a coloring. A template row (p, z) consists of its bin label p and a tuple z that specifies, for each index α, a pair $z_\alpha = (r, u)$ of a position $r \in \{1, \ldots, s-2\}$ and a base row u. Bob's first-family column supplies (α_p, γ_p) for each bin, where γ_p is a list of $2s-1$ base columns. The entry is $R(u, \gamma_{p,r})$, with $(r, u) = z_{\alpha_p}$. Thus each template row uses only its own bin's data. Letting the data of different bins vary independently gives the outer interlace, with one block per bin, that forces Alice's initial messages.

If a template row carried just one pair (r, u) instead of the tuple z, its entry after the bin was chosen would not depend on Bob's index α_p. The template subproblem would no longer require that index to be communicated. With the tuple, Bob's index selects which pair of Alice's input matters, and then Alice's position r selects which entry of Bob's list matters. The players are left with one instance of R. These successive selections are the nested interlaces, with transposition exchanging the players' roles between levels.

Once Alice's first $\log s$ bits are forced, they split the item rows into the same s parts as the template rows, and this yields a bin assignment, as explained in the next subsection.

An item row (i, u) has no bin label, because the reduction cannot use a coloring to assign one. Bob therefore also holds a list $\xi = (\xi_1, \ldots, \xi_{s+1})$ of base columns shared by all bins. For each index α, the at most $s-1$ active items receive distinct positions $\pi_\alpha(i) \in \{1, \ldots, s\}$ in a fixed public order. If all Bob's indices agree on α and item i is active there, the entry is $R(u, \xi_{\pi_\alpha(i)})$, and otherwise it is $R(u, \xi_{s+1})$. These rules use no coloring. In the coloring protocol, Bob checks privately whether his indices agree, and if Alice's item is active, her bin and the index identify it.

The second family sets every bin's index to the same α and every list to the same γ, with $\xi = (\gamma_{s-1}, \ldots, \gamma_{2s-1})$. The template and item rows now read positions of one common list. A single protocol must handle both families, and before Bob speaks, Alice's messages depend only on her row. Consequently, the second family detects collisions in the partition forced by the first.

The second family has separate labels and a fixed positive share of the column weight. In the full-product construction, its columns already occur in the first family but may have very small total weight there. The separate weight prevents a small deletion from removing all collision tests. Definition 6.2 gives the full entry rules and weights.

A coloring gives a protocol. In the YES case, fix the bin assignment given by a proper 3-coloring. Alice names her bin p, using $\log s$ bits, and Bob sends that bin's index α_p, using $\log q_2$ bits, where

the number q_2 of indices is a power of two. Alice then sends one of

$$\underbrace{s-2}_{\text{template positions}} \quad + \quad \underbrace{1}_{\text{active item}} \quad + \quad \underbrace{1}_{\text{inactive item}} \quad = s$$

possible messages. If Alice's item is active, her bin and the index determine it, since a proper coloring puts at most one active item in each bin. Bob can therefore identify the selected base column from his input, and the players finish by running a protocol for R.

The two Alice messages each cost $\log s$ bits, giving total cost $C - 1 := d + 2\log s + \log q_2$. For an arbitrary graph, Alice can instead name her active item's position explicitly. This uses at most $2s - 1$ message values and costs one additional bit. Thus the table always has a cost-C protocol, and a coloring saves one bit (Proposition 6.4).

2.2 Why the reduction rules out every short protocol

We prove that the NO-case lower bound survives deletion of a small fraction of the row and column weight, so that it can be reused in later rounds. Suppose a protocol of cost $C - 1$, with any pattern of interaction, computes the remaining matrix. From such a protocol we extract a single bin assignment. Its first $\log s$ bits are Alice's, sent before Bob can reveal any index. They split Alice's rows into s parts, and we assign each item to a part that receives at least a $1/(2s)$ fraction of the weight of its rows. For a NO graph, this fixed assignment must contain collisions.

A short protocol fixes a bin assignment. Group the template rows by their bin p. In the first column family, Bob holds a separate index and list of base columns for every bin, and in the full matrix every combination occurs. The entry of a template row in bin p depends only on the data for that bin. On template rows the matrix is therefore an s-fold interlace whose blocks are the bins.

Within one bin, the template rows already require $d + \log q_2 + \log s$ bits, by two applications of the interlacing lower bound (Lemma 6.6). Of the $C - 1$ available bits, only $\log s$ remain to tell the bins apart. Suppose Bob sends a bit while ℓ bins are still possible. His heavier child keeps every Alice row and at least half the column weight, so the interlacing bound for ℓ bins still demands $\log \ell$ extra bits, one more than remain. Similarly, if one of Alice's bits kept half the weight of more than $\ell/2$ bins on one side, that side would need more bits than remain. The separation theorem (Theorem A.3) therefore forces the first $\log s$ bits to be Alice's and to split the bins evenly. Grouping her rows by these $\log s$ bits gives s parts, one for each bin. Each part contains almost all the template-row weight of one bin and very little of the weight of the others.

All items of one type have rows of equal total weight, so a small deletion makes few items lose more than half of the weight of their rows (i, u). We discard these and assign each remaining item to a part containing at least a $1/(2s)$ fraction of the original weight of its rows. Such a part exists by averaging over the s parts. The rows of one item need not stay in one part, since different base rows u may follow different paths. However, the assignment is fixed before Bob has spoken, so it cannot depend on the edge tested by his index. For a NO graph, Lemma 6.3 guarantees many indices at which one assigned bin contains two active items and an inactive item.

A collision gives an extra position. Restrict Bob to the second column family at a collision index α. All bins now use the same index and the same list of base columns. Within the collision's bin, the template rows and the three items read

$$\underbrace{s-2}_{\text{template positions}} \quad + \quad \underbrace{2}_{\text{distinct active positions}} \quad + \quad \underbrace{1}_{\text{inactive position}} \quad = s+1$$

distinct positions. The inactive item is essential: two active items alone would give only $(s-2)+2 = s$ positions. As s is a power of two, naming one of $s+1$ positions costs one more bit than one of s.

Why collisions force more communication. Counting positions alone does not prove a lower bound against an interactive protocol. Bob's replies restrict which lists of base columns remain possible, and a lower bound for the original column set need not survive this restriction. In the second family, every bin has the same index α. After Alice's $\log s$ bits, Bob may send bits that narrow down the possible values of α before Alice speaks again, so that each branch of the protocol covers only some indices. If collisions occurred at only a few indices, Bob could isolate them on a branch where naming α costs far fewer than $\log q_2$ bits, leaving room for the extra position. We therefore need collisions at a constant fraction of all indices.

For each index, we track Bob's retained second-family inputs with that index through his bits, and stop when Alice next speaks. Each index at which the deletion leaves most of these inputs gives one tracked set, and a collision marks its set, recording the extra bit. An unmarked set follows a child retaining at least half its weight. A marked set splits into two unmarked sets if both children keep enough weight, and otherwise keeps its mark. Counting marked sets twice, a split replaces two by one plus one, preserving the count.

A stopping node reached after r Bob bits can accommodate a count of at most $q_2 2^{-r}$ within its remaining communication budget. The stopping nodes are prefix-free, so Kraft's inequality bounds their total capacity by q_2. Initially, however, the count exceeds q_2: collisions occur at a constant fraction of the indices (Lemma 6.3), while each index carries the same weight, so the small deletion damages far fewer. This contradiction rules out cost $C-1$ in the NO case (Lemma 6.8 and Theorem 6.9). Together with the coloring protocol, it yields the one-bit gap against every protocol, whatever its pattern of interaction.

2.3 Making the one-bit gap repeatable

A one-bit gap does not automatically yield a larger gap. Using the output as the next base matrix requires lower bounds on the restrictions that remain after messages, and a hard matrix can have easy restrictions. For example, restricting Bob's tuples in an interlace of R to those with all coordinates equal removes the need to learn Alice's index.

We therefore maintain a stronger invariant in the NO case. Measuring inputs by the fixed positive rational weights, we require, for a threshold ρ fixed across rounds, that every restriction retaining the weights below needs the stated communication:

retained row weight	retained column weight	communication lower bound
$\geq \rho$	$\geq 3/5$	d
$\geq \rho$	$\geq 3/10$	$d-1$.

The threshold ρ may depend on the source size. We call these the *deletion bounds* (Definition 5.1). The second column threshold is half the first: after Bob sends one bit, the heavier branch retains at least half the column weight, and the allowed lower bound drops by one. The YES upper bounds do not require this invariant.

The source construction gives its extra bit only on restrictions retaining almost all the weight on both sides (Theorem 6.9). The deletion bounds needed for another round allow most of the row weight and a constant fraction of the column weight to disappear. Two further interlaces close this difference, each handling one side.

Preserving the extra bit after input deletion. Transpose the source matrix to N, so its two column families become row families, and form the w-fold interlace $U = \mathsf{I}_w(N)$. Alice holds a copy index and a row of that copy, and Bob holds one column for each copy. Consider a restriction retaining almost all row weight but only a small fixed fraction θ of the tuple weight. Suppose this restriction has a protocol of cost $C + \log w - 1$, one bit less than naming a copy and running the cost-C protocol for N.

Stop each path just before Bob's first bit or just after Alice's $(\log w)$-th bit, whichever comes first, padding paths that end earlier with fixed Alice bits. The first-family rows force almost all their weight in most copies to reach distinct stopping nodes at depth $\log w$, one assigned to each such copy (Lemma 6.13). At these nodes $C - 1$ bits remain. Bob has not spoken, so his tuple set is still the initial restriction. By the conditioning lemma (Lemma 6.10), almost every copy still retains almost all its column weight. If its assigned node also kept almost all its rows, the source lower bound would require C further bits. So the node must miss a substantial part of its copy's rows.

Almost all first-family weight is already at the assigned nodes, and little weight was deleted initially. Summing over the assigned copies therefore places substantial surviving second-family weight outside the node assigned to its copy. Since only Alice has spoken, every surviving row reaches exactly one stopping node. Each such displaced row reaches either another copy's assigned node or a node with no assigned copy.

Neither group can receive enough weight (Figure 3). An assigned node already contains the hard first-family rows of its own copy, so much second-family weight from another copy would need more than the remaining $C - 1$ bits (Lemma 6.15). The unassigned nodes have small total capacity (Lemmas 6.13 and 6.14). The two bounds sum to less than the displaced weight, a contradiction (Theorem 6.16). Transposing U exchanges the two retained-weight conditions, yielding W with its full lower bound when a small fraction of rows and almost all columns remain.

Recovering the bit in the deletion bounds. For constant retained column weight, W satisfies the deletion bounds only one bit below its full cost. A final interlace combines this guarantee with the full-cost bound just proved. For full products, with product tuple weights, a tuple set of retained weight v cannot have small projections at too many coordinates. If its projection onto coordinate i has original weight p_i, containment in the product of its projections gives

$$v \le \prod_i p_i, \qquad \text{hence} \qquad \sum_i \log(1/p_i) \le \log(1/v).$$

Among sufficiently many copies retaining enough row weight, at least one therefore retains almost all its column weight, so the full-cost bound for W applies.

Finding one such copy would recover only the cost of W, and we must also account for the bits needed to identify a copy. We do so by induction over messages. At an Alice bit, one branch keeps enough copies with substantial row weight. At a Bob bit, keeping the heavier tuple set increases $\log(1/v)$ by at most one. Choosing enough copies at the outset leaves room for these increases. The deletion bounds control the intermediate restrictions, and the full-cost bound closes the remaining one-bit deficit (Theorem 6.17). For the explicit lists, the conditioning argument described next replaces the full-product inequality (Theorem 7.2).

Consequently, the three steps above form one round $R \mapsto R^+$, which has a public overhead h and satisfies the following guarantee. In the YES case,

$$\mathrm{D}(R^+) \le \mathrm{D}(R) + h.$$

In the NO case, if $\mathrm{D}(R) = d \ge 3$ and R satisfies the deletion bounds at value d, then

$$\mathrm{D}(R^+) = d + h + 1,$$

and R^+ satisfies the same deletion bounds at value $d + h + 1$ (Theorem 5.3). This restoration is what makes iteration possible. Starting from the initial matrix of known cost $d_0 \geq 3$ (Lemma 5.4), which satisfies the deletion bounds, the matrix R_k after k rounds has

$$\text{YES: } \mathrm{D}(R_k) \leq d_0 + kh, \qquad \text{NO: } \mathrm{D}(R_k) = d_0 + kh + k.$$

The threshold $d_0 + kh$ is computable without deciding whether the graph is colorable.

2.4 Keeping the deterministic construction polynomial

Allowing every tuple of q columns from a set Y creates $|Y|^q$ inputs. We replace each product by an explicit list of polynomially many tuples with positive rational weights. For an absolute constant D, the list has exactly the same joint weights as the full product on every set of at most D coordinates. This is *limited independence*. We build and enumerate the entire list deterministically, and the weights do not represent random choices made by the reduction.

We extend tuples one coordinate at a time and remove tuples by adjusting their weights. Associate with each tuple a zero–one vector indicating its assignment on every set of at most D coordinates, plus a final entry one. The required joint weights and the total weight specify the weighted sum of these vectors. If the list is longer than the vector dimension, a linear dependence lets us change the weights without changing that sum. Since total weight is fixed, the change has both positive and negative components, so we can move along it until one tuple has weight zero and remove it, keeping all other weights nonnegative. Compressing after each new coordinate keeps the intermediate lists short, avoiding enumeration of an exponentially large product. With q coordinates over m labels, the list has at most $(1 + qm)^D$ tuples with polynomial-bit weights (Lemma 7.1). The original entry rules still apply to every listed tuple, so the upper-bound protocols are preserved.

The lower bound is subtler, because conditioning on earlier messages can destroy independence. For example, retaining only tuples whose first two coordinates agree makes those coordinates dependent. Our conditioning lemma shows that, after conditioning on any event of weight at least v, only $\mathrm{poly}(1/v)$ coordinates are exceptional. On every set of at most eight other coordinates, the conditional distribution is close in total variation to the original product distribution (Lemma 6.10). By a moment bound under the original weights, too many disjoint coordinate sets with substantially changed distributions would contradict the weight of the conditioning event.

We choose the interlacing widths and deletion thresholds so that enough coordinates remain outside the exceptional set. Projecting onto these coordinates loses at most the total-variation error in column weight, so the lower bounds of the base matrix still apply (Theorems 6.12 and 7.2).

The widths are fixed polynomials in n, so each round takes time polynomial in the source graph size n and the encoding length of the current matrix, with an exponent independent of the round. After k rounds, size and running time are at most $n^{O(A^k)}$ for an absolute constant $A > 1$. Every fixed gap therefore gives a deterministic polynomial-time reduction (Theorems 1.1 and 8.4).

For sparse formulas, $n = N^{O(1)}$, so a suitable $k = \Theta(\log \log N)$ makes the construction quasipolynomial in N, while $\log \log T = O(k)$. A deterministic polynomial-time estimator with additive error $c_* \log \log T$, for a small enough constant c_*, would therefore decide sparse 3-SAT in quasipolynomial time (Corollary 1.5).

2.5 Two further consequences of keeping fewer inputs

Sampling at the end tolerates larger errors. For the randomized result, we keep full products of tuples throughout the amplification. The table is then too large to print, but we can compute single entries from short descriptions of the inputs. After all k rounds, we use the input weights as

sampling probabilities: we draw m rows and m columns independently and print the m^2 entries at their intersections. Randomness is used only by the reduction; the protocols whose cost is measured remain deterministic.

In the NO case, a protocol may be chosen after the sample is drawn, so the sample must defeat every protocol of cost below the NO cost simultaneously. Fix a protocol tree with Alice's messages and the leaf outputs, but not Bob's bits. A column is *compatible* with it and a row set if one choice of Bob's bits, depending only on the column, makes the tree correct on all those rows. Over all such trees of depth t and all row sets, a matrix of real rank r has at most $2^{O(r \log(r+2)\, 2^t)}$ distinct sets of compatible columns (Theorem 9.3). The count is independent of the number of input labels. For a tree of depth below the NO cost and a row set of weight $\geq 3/5$, the deletion bounds (with $\rho \leq 3/5$) give the compatible columns weight $< 3/5$. A union bound over these sets then shows that, with high probability, $m = \Theta(r \log(r+2)\, 2^t)$ random columns are not all contained in any one of them. Once such a column sample is fixed, every row set compatible with a short protocol on the sampled columns has weight $< 3/5$, so the same union bound applies to the sampled rows (Theorem 9.4). The sample need only preserve the gap, since no further amplification follows. An interlace stacks one block per copy, so it multiplies the rank by at most its width. One round thus multiplies it by at most 2^h, and $\operatorname{rank}(R_k) \leq 2^{d_0 + kh}$ (Proposition 5.5). For a sparse formula on N variables the graph has $N^{1+o(1)}$ vertices and $h = (1 + o(1)) \log N$, so taking $k \approx N^\alpha / \log N$ gives $\log T = 2 \log m = (4 + o(1)) k \log N$ and $k = (\alpha/4 + o(1)) \log T / \log \log T$. Sampling and evaluation take time $n^{O(k+1)} = 2^{O(N^\alpha)}$ (Lemma 9.5). An estimate within additive error $c \log T / \log \log T$ distinguishes the two cases whenever $\alpha > 8c$, giving Corollary 1.4.

Strict alternation needs a different encoding. The coloring protocol above uses consecutive bits from the same player. Inserting dummy turns can cost more than the gap, so it does not give hardness under strict alternation.

We instead start from the hypergraph-coloring reduction of Guruswami and Lee [8], whose balanced YES colorings give a short alternating protocol. For a power of two s, each edge has s^2 vertices. In the YES case there is an s-coloring in which every edge contains at least $s - 1$ vertices of every color. In the NO case, every coloring with at most 2^s colors has a monochromatic edge (Proposition 10.2).

Alice holds a vertex and a tuple of $s/2$ addresses. Bob holds an edge, a coordinate selecting one address, and a Boolean array indexed by an address and an edge position. The selected address and the vertex's position in Bob's edge identify the array bit to output. The output is zero if the vertex is outside the edge (Definition 10.3). Write $s = 2^t$, and let b bits name an edge. In the YES case, the players exchange the color, address and edge information in alternating turns. The coloring leaves at most $2s - 1$ candidate vertices of Alice's color in Bob's edge, giving total cost $2b + 4t + 2$ (Proposition 10.4).

For the NO lower bound, fix a short protocol and, for each row, record Alice's responses to *every* possible sequence of Bob's earlier bits during her first t turns. There are at most $2^{2^t - 1} \leq 2^s$ such response rules. Assign each vertex its most frequent rule, used by at least a 2^{-s} fraction of its address tuples. The NO promise gives a monochromatic edge, so many rows sharing one early response rule must still identify many different array positions. Counting only the transcript of one interaction would not suffice, because Bob's messages may depend on his entire array.

We construct a polynomial-size explicit list of arrays in which each pattern on any j positions occurs in at most a $2^{-j} + \varepsilon$ fraction of the arrays. Following a hypothetical short protocol while retaining large sets of rows and arrays leaves many arrays. By the list property, the few output patterns of the last turns on the retained rows match fewer arrays than survive (Corollary 10.6,

Lemma 10.5, and Proposition 10.7). This gives a NO lower bound of $2b + 6t - 6$, and hence a gap of $2t - 8$. Allowing either starter and receiver-only output loses at most two bits (Lemma 10.1). Choosing t as a sufficiently large constant therefore gives a deterministic polynomial-time reduction for every fixed gap in Yao's convention.

3 Related work

Which optimization problem is hard? For explicitly given total tables, Kushilevitz and Weinreb proved that approximating deterministic communication complexity is hard under cryptographic assumptions [16]. Ilango, Loff and Oliveira proved NP-hardness and strong inapproximability for *partial* matrices, where the protocol is required to be correct only on specified input pairs [12]. Our tables are *total*: every pair of a row and a column has a specified bit, and the protocol must be correct on all of them.

For total matrices, Hirahara, Ilango and Loff proved additive hardness for the number of bits in a three-message protocol, and hardness of approximating the number of leaves of a protocol tree for every fixed number $k \geq 3$ of messages [9]. A message is a block of consecutive bits from one player, and a leaf is a possible end of the conversation. Neither restricting the number of messages nor counting leaves gives the same optimization problem as minimizing bits under strict alternation or unrestricted speaking order. The exact-hardness results for the unrestricted model [6, 10] left open both hardness with one bit of additive error and Yao's original alternating model, as Hirahara, Ilango and Loff noted [10]. Our fixed-gap theorems resolve both.

Ingredients of our reductions. For the unrestricted reductions, we obtain hard graph-coloring instances by combining reductions of Dinur and of Bogdanov, Obata and Trevisan [3, 5]. In these graphs, each vertex meets only a constant number of edges. The alternating reduction uses the hypergraph-coloring promise of Guruswami and Lee [8] described in Section 2.

Both reductions also replace large input sets by short explicit lists that keep the statistics their lower-bound proofs use. For the deterministic interlacing construction, we follow Koller and Megiddo [15, Section 3]: we extend a list of tuples one coordinate at a time and delete tuples without changing the joint weights of the values on every set of a bounded number of coordinates. For strict alternation, we need a short list of Boolean arrays on which no bit pattern at a few positions is much more frequent than for random arrays, and we take it from the small-bias construction of Alon, Goldreich, Håstad, and Peralta [1].

Bounds on the number of column sets on which short protocols can be correct, expressed in terms of matrix rank, appear in graph algorithms [4, 20] and in work on hard submatrices [11]. Our count (Theorem 9.3) uses the row-sum method of Oum, Sæther and Vatshelle [20]: the integer sum of some Boolean rows is zero exactly in the columns where all those rows are zero. Mackenzie's hardness condensation [17] also counts the input sets on which short protocols succeed, to shrink a communication problem while keeping its cost high relative to the new input length. Our sampling theorem (Theorem 9.4) has a different goal: a sampled subtable whose cost stays above a given bound.

Reading guide. After the preliminaries (Section 4), the formal sections follow the order of Section 2, and the appendices contain the remaining estimates and parameter checks.

4 Preliminaries

4.1 Communication and estimation

An estimator receives the whole table, whereas each player knows only its own input.

Definition 4.1. A table $M : X \times Y \to \{0,1\}$ is total, and its input sets X and Y are finite and nonempty sets of labels. Equal rows or columns may have different labels. Its size is $T = |X||Y|$. A deterministic public-output protocol is a binary tree whose internal nodes specify a speaker and a bit-valued function of that player's input and the preceding transcript. Each leaf has a fixed output, and the protocol must compute every entry correctly. Its cost is its maximum depth, charging every bit, including fixed bits. Local computation is free. Write $\mathrm{D}(M)$ for the least cost with unrestricted speaking order.

An additive-$a(T)$ estimator returns $\widehat{D}(M)$ satisfying $|\widehat{D}(M) - \mathrm{D}(M)| \le a(T)$. Its output must be comparable in polynomial time with any binary rational of bit length polynomial in T. A bounded-error randomized estimator satisfies this guarantee with probability at least $2/3$ on every input. Polynomial time means worst-case time polynomial in T. Estimators for the alternating costs below follow the same conventions.

In the alternating model, the starting player and the output convention can each affect the cost.

Definition 4.2 (Alternating costs)**.** In strict alternation the players speak one bit at a time, in turn. Let D_A^{alt} and D_B^{alt} be the minimum public-output costs when Alice or Bob starts, respectively, and let $D_{\mathrm{free}}^{\mathrm{alt}} = \min\{D_A^{\mathrm{alt}}, D_B^{\mathrm{alt}}\}$. The starter is chosen independently of the inputs. Every bit is charged, and stopping is public: the transcript determines whether communication has ended.

Yao's receiver-output cost Y_{free} permits the same choice of starter, but the receiver of the last bit determines the answer from that player's input and the transcript. A protocol of depth zero instead designates one player, who determines the answer from that player's own input alone. Stopping is still public and correctness is required on every input.

Our quantitative statements use sparse 3-SAT: the input has N variables and at most $c_s N$ clauses, where c_s is fixed in advance. Running times are worst-case bounds, including all random tapes for randomized algorithms; bounded error means success probability at least $2/3$ on every formula. Constants may depend on c_s and on the estimator. We do not apply a sparsification step with an uncharged exponential cost, and all stated running times refer directly to these sparse inputs.

All communication protocols in this paper are deterministic and exact. Only estimators and randomized reductions use random bits. All logarithms are binary unless written ln. For an integer $q \ge 1$ write $[q] = \{1, \dots, q\}$, and for $a \ge 1$ write $\mathrm{pow}_2(a) = 2^{\lceil \log a \rceil}$ for the least power of two at least a.

Every $n \times n$ table has $\mathrm{D}(M) \le \lceil \log n \rceil + 1$: Alice sends her row index and Bob sends the entry.

4.2 Relabeling and input weights

We keep the matrices small in two ways, by explicit short lists and by sampling. In both, each new row or column label decodes to an original one.

Let $F : X \times Y \to \{0,1\}$, $f : X' \to X$ and $g : Y' \to Y$, and define F' on $X' \times Y'$ by $F'(x', y') = F(f(x'), g(y'))$. Alice computes $f(x')$, Bob computes $g(y')$, and they run a protocol for F. Hence F' has at most the communication cost of F in every model above, and at most its rank. Costs and ranks are equal when f and g are onto, since then F is obtained from F' in the same way. Transposition preserves rank and unrestricted cost, and exchanges D_A^{alt} and D_B^{alt}.

In particular, duplicating rows and columns makes a table square with power-of-two side length without changing its cost. We may also require the side length to be at least a prescribed bound b polynomial in the input size. The side length is then less than $2\max\{b, |X|, |Y|\}$, so the enlarged table can be printed in polynomial time.

For the unrestricted construction, we assign the labels in each input set positive rational weights summing to one. For row weights μ_X and column weights μ_Y, a set's weight is the sum of its labels' weights, for example $\mu_X(A) = \sum_{x \in A} \mu_X(x)$. Uniform weights give ordinary fractions of labels. Unequal weights let us prescribe the total weight of each input family and shorten tuple lists while preserving their marginals (Section 7). The weights are auxiliary data for the construction and its lower-bound proof. Communication cost depends only on the table. The final deterministic reduction discards the weights. Distinct labels keep separate weights, even when their rows or columns agree.

If X' has weights μ', a decoder $f : X' \to X$ induces weights on X by adding the weights of labels with the same image, and labels of X outside the image get weight zero. A set $A \subseteq X$ therefore has induced weight $\mu'(f^{-1}(A))$. In particular, a retained set $A' \subseteq X'$ has image of induced weight at least $\mu'(A')$, since $A' \subseteq f^{-1}(f(A'))$.

Under product weights on Y^q, a tuple $(y_1, \ldots, y_q)$ has weight $\prod_{j=1}^{q} \mu_Y(y_j)$. Some counting arguments are shorter in probability notation. There we regard a label as drawn according to its weights and use the standard terms distribution, marginal, and independence. Such a draw only serves the calculation and is not a random choice by a communication protocol. An event is a set of labels. Conditioning on a retained set A of positive weight means renormalizing: $(\mu \mid A)(B) = \mu(A \cap B)/\mu(A)$. When we project retained tuples onto some coordinates, we specify whether we measure the projection with the original or the conditional weights.

Counting surviving blocks. Later proofs need many blocks of a partition each to retain a given fraction of their own weight, which we derive from the total retained weight as follows. Suppose we partition an input set into q blocks, each of weight $1/q$, such as the blocks $\{i\} \times X$ of rows of $\mathsf{I}_q(R)$ in Definition 4.3 below. If a retained subset has total weight β, the fractions of their own weight that the blocks retain sum to $q\beta$. We set aside at most e blocks and suppose that at most r of the other blocks retain fraction $\geq t$. Bounding these $r + e$ fractions by one and every remaining fraction by t gives

$$q\beta \leq r + e + qt.$$

Thus, if $q\beta > r + e + qt$, more than r blocks not set aside retain fraction $\geq t$. If the q blocks instead partition only a part of the input set, of total weight λ, that part retains fraction at least $(\beta - (1 - \lambda))/\lambda$ of its weight, since the rest of the input set contributes at most $1 - \lambda$ to β.

4.3 Interlacing and restrictions

In an interlace, a row is a row of the base matrix together with an index i, a column is a tuple of base columns, and the entry is read from the ith coordinate of the tuple.

Definition 4.3 (Interlacing)**.** For a Boolean matrix $R : X \times Y \to \{0, 1\}$, the full q-fold interlace $\mathsf{I}_q(R)$ has rows $(i, u) \in [q] \times X$, columns $(y_1, \ldots, y_q) \in Y^q$, and entry

$$\mathsf{I}_q(R)((i, u), (y_1, \ldots, y_q)) = R(u, y_i).$$

When R has normalized input weights μ_X, μ_Y, give row (i, u) weight $\mu_X(u)/q$ and column $(y_1, \ldots, y_q)$ weight $\prod_{j=1}^{q} \mu_Y(y_j)$.

Why column deletion matters. For a fixed base row u, the rows (i, u) with different indices i read different coordinates of Bob's tuple, but these rows may coincide after columns are deleted. For example, consider the 2×2 base matrix

$$R = \begin{array}{c|cc} & 1 & 2 \\ \hline a & 0 & 1 \\ b & 1 & 0 \end{array}.$$

The left-hand table in Figure 1 is its two-fold interlace. We obtain the right-hand table by keeping only columns $(1, 1)$ and $(2, 2)$. For each fixed $u \in \{a, b\}$, the rows $(1, u)$ and $(2, u)$ then coincide.

Rows may depend on i

u	i	(1,1)	(1,2)	(2,1)	(2,2)
a	1	0	0	1	1
a	2	0	1	0	1
b	1	1	1	0	0
b	2	1	0	1	0

Shaded columns are deleted.

keep $y_1 = y_2$ →

Rows no longer depend on i

u	i	(1,1)	(2,2)
a	1	0	1
a	2	0	1
b	1	1	0
b	2	1	0

identical rows for fixed u

Figure 1: Deleting columns with $y_1 \neq y_2$ removes the dependence on the row index i. For each fixed u, the rows $(1, u)$ and $(2, u)$ agree on every retained column.

This collapse occurs for every base matrix R. On the columns $(y, \ldots, y)$, the entry of $\mathsf{I}_q(R)$ is $R(u, y)$ for every row index i. Hence a lower bound for the full interlace need not survive arbitrary column deletion. For this reason, Definition 5.1 requires lower bounds only for restrictions that retain a prescribed weight of rows and of columns.

Keeping fewer column tuples. Full interlacing creates $|Y|^q$ columns, too many for an explicit polynomial-time reduction. We therefore replace Y^q by a short list with rational weights. For every assignment of values to at most D coordinates, we preserve the total weight of tuples matching that assignment. The definition below calls these exact counting identities D-wise independence. Lemma 7.1 constructs such lists deterministically (Sections 7 and 8).

Definition 4.4 (Weighted lists of tuples)**.** Fix integers $q, D \geq 1$, the tuple length and the independence order. For a finite alphabet Y with prescribed normalized rational weights μ, an indexed tuple list consists of a finite nonempty label set $\mathcal{L}$, positive normalized rational weights ν on its labels, and a decoder $z : \mathcal{L} \to Y^q$. For $J \subseteq [q]$, write $z(\ell)_J$ for the restriction of $z(\ell)$ to J. The list is *exactly D-wise independent* if, for every J with $|J| \leq D$, the map $z_J : \ell \mapsto z(\ell)_J$ induces the product weights μ^J. These induced weights are the marginal on J. No independence on larger sets is required. A set of labels may contain one label but not another with the same decoded tuple.

For a base matrix $R : X \times Y \to \{0, 1\}$ with column weights μ, the interlace on this list has a column for each label ℓ, with entry $R(u, z_i(\ell))$ at row (i, u). Rows keep the weights of Definition 4.3, and columns get the weights ν. For both kinds of interlace, *block i* is the row set $\{i\} \times X$, and we measure row weight within a block by the base row weights, without division by q.

The full product is the list with labels Y^q, identity decoder and product weights. For a label set $K \subseteq \mathcal{L}$ and a set $J \subseteq [q]$, write $\pi_J K = \{z(\ell)_J : \ell \in K\} \subseteq Y^J$ for the projection of its decoded tuples.

To apply a lower bound to a set J of surviving blocks, we project the retained column tuples onto the coordinates in J. By the next observation, the resulting matrix is a submatrix of the restriction.

Observation 4.5 (Restrictions projected to a set of blocks)**.** *Let G be a restriction of an interlace of R, on the full product or on a tuple list, with retained label set K, and let J be a set of blocks. For each $i \in J$, the rows in block i read only coordinate i of the decoded tuple. For each $w \in \pi_J K$, choose a label $\ell_w \in K$ with $z(\ell_w)_J = w$. Identifying ℓ_w with w, the restriction of G to these labels and to the retained rows of the blocks in J is then the full-product interlace on J, restricted to the columns $\pi_J K$ and to the same rows. This matrix is a submatrix of G, so every communication lower bound for it holds for G. In that interlace, its column set has weight $\mu^J(\pi_J K)$, and each block keeps its base row weights. If the marginal on J is μ^J, as for full products and for exactly $|J|$-wise independent lists, then $\mu^J(\pi_J K) \geq \nu(K)$, since $K \subseteq z_J^{-1}(\pi_J K)$. We never use the weights of the chosen representatives, and we assert no independence of the retained labels.*

5 A reusable gap construction

In this section we state the one-round amplifier $R \mapsto R^+$ and then iterate it. In the NO case the amplifier raises the cost by one bit more than in the YES case, and its output again satisfies the deletion bounds (Definition 5.1) that the next round needs.

5.1 Lower bounds that survive restrictions

The deletion bounds use two column-weight thresholds, 3/5 and 3/10. The second is half the first because, after Bob sends one bit, the heavier branch retains at least half the column weight. At the lower threshold we allow the communication lower bound to drop by one bit.

Definition 5.1 (Deletion bounds)**.** Let $F : X \times Y \to \{0,1\}$ have positive normalized rational input weights. For $0 < \rho \leq 1$ and an integer d, F satisfies the *deletion bounds* $\mathcal{P}(F; \rho, d)$ if, for every $j \in \{0,1\}$ and all $U \subseteq X$, $V \subseteq Y$,

$$\mu_X(U) \geq \rho \text{ and } \mu_Y(V) \geq (3/5)2^{-j} \quad \Longrightarrow \quad \mathrm{D}(F[U,V]) \geq d - j.$$

Here $F[U,V]$ is the restriction of F to $U \times V$. No matching upper bound is required.

Whether the source graph is 3-colorable determines whether the constructed matrix has the smaller (YES) or the larger (NO) cost. For an unsatisfiable formula, we need the graph to stay non-3-colorable after a small fraction of its vertices is deleted, since otherwise deleting inputs could erase the lower bound. The reductions of Dinur and of Bogdanov, Obata and Trevisan provide such graphs [3, 5].

Theorem 5.2 (Bounded-degree coloring source)**.** *There are constants $\Delta \geq 1$ and rational $\gamma > 0$, and a deterministic polynomial-time reduction from 3-SAT to nonempty loopless multigraphs $G = (V, E)$ with no isolated vertices and maximum degree at most Δ, such that:*

1. *a satisfiable formula produces a 3-colorable graph;*

2. *an unsatisfiable formula produces a graph for which every* 3*-coloring has at least* $\gamma|E|$ *monochromatic edge occurrences.*

For sparse formulas on $N \geq 1$ *variables, there is a fixed positive integer* p *such that the output satisfies, for* $n = |V|$*,*

$$N \leq |E|, \qquad n, |E| \leq N(\log(N+2))^p, \qquad n/2 \leq |E| \leq \Delta n/2.$$

The proof in Section G.1 composes the two cited reductions and checks the size bounds and graph conventions stated above.

We call graphs from satisfiable formulas YES and those from unsatisfiable formulas NO. We build the amplifier in two ways, which satisfy the same guarantee. In the *full-product case*, used for the randomized reduction, every interlace uses all column tuples, and Lemma F.2 fixes the widths and thresholds. In the *explicit case*, used for the deterministic reduction, the interlaces use the small tuple lists of Definition 8.3, and Lemma F.3 fixes the widths and thresholds. In each case a final interlace restores the deletion bounds at the full output cost for the next round, by Theorem 6.17 in the full-product case and by Theorem 7.2 in the explicit case.

Theorem 5.3 (Amplification under input deletion)**.** *Let* $G = (V, E)$ *be a YES or NO graph produced by Theorem 5.2, and put* $n = |V|$*. In either case there are a public row-weight threshold* ρ*, a public integer* h*, and a deterministic, implicitly represented construction* $R \mapsto R^+$ *on base matrices* R *with positive normalized rational input weights. In the full-product case,* n *must exceed an effectively computable constant,* $\rho = 2^{-L}$ *with* $L = \lceil \log n \rceil$*, and* $h = L + O_{\Delta,\gamma}(\log L)$*. In the explicit case,* $\rho = 1/\mathsf{N}$ *and* $h = 10 \log \mathsf{N}$*, with* $\mathsf{N} = s^4$ *for the power of two* s *of Lemma F.3. In both cases:*

(i) On either source promise,

$$\mathrm{D}(R^+) \leq \mathrm{D}(R) + h + 1,$$

and, if G *is YES, the stronger bound* $\mathrm{D}(R^+) \leq \mathrm{D}(R) + h$ *holds.*

(ii) If G *is NO and the base matrix* R *satisfies* $\mathrm{D}(R) = d \geq 3$ *and* $\mathcal{P}(R; \rho, d)$*, then*

$$\mathrm{D}(R^+) = d + h + 1, \qquad \mathcal{P}(R^+; \rho, d + h + 1).$$

(iii) On either source promise, $\mathrm{rank}(R^+) \leq 2^h \, \mathrm{rank}(R)$*.*

The proof of Theorem 5.3 is assembled at the end of Section 6 (p. 32). It uses the interlacing estimates of Section A and, in the explicit case, the small lists of Section 7. The rank bound (iii) follows from Lemma 6.18.

5.2 Iteration

One round adds h bits to the YES upper bound and $h + 1$ bits to the NO lower bound. Starting from the initial matrix of Lemma 5.4, of known cost d_0, the two bounds therefore differ by i bits after i rounds.

Lemma 5.4 (Initial matrix)**.** *For* $L \geq 1$*, let* $p_0 = \mathrm{pow}_2(64(L+1))$ *and* $R_0(z, j) = z_j$ *on* $\{0,1\}^{p_0} \times [p_0]$*, with uniform input weights. Then*

$$d_0 := \log p_0 + 1, \qquad \mathrm{D}(R_0) = d_0, \qquad \mathrm{rank}(R_0) = p_0 = 2^{d_0 - 1}, \qquad \mathcal{P}(R_0; 2^{-L}, d_0).$$

For $L = \lceil \log n \rceil$ *and* $n \geq 4$*, the initial matrix has polynomial size in* n *and can be printed in polynomial time.*

The proof is in Section A.4 (p. 55).

Proposition 5.5 (Iterated amplification)**.** *Fix one case of Theorem 5.3, whose threshold is ρ and whose overhead is h, and a source graph to which it applies. Iterate the theorem from the matrix R_0 of Lemma 5.4 with $L = \log(1/\rho)$, an integer in both cases, using the same parameters in every round, and let R_i be the matrix after $i \geq 0$ rounds. Put $B_i = d_0 + ih$, $H_i = B_i + i$, and $r_i = 2^{B_i - 1}$. A YES graph gives $\mathrm{D}(R_i) \leq B_i$. A NO graph gives $\mathrm{D}(R_i) = H_i$ and $\mathcal{P}(R_i; \rho, H_i)$. On either promise, $\mathrm{rank}(R_i) \leq r_i$ and $\mathrm{D}(R_i) \leq H_i$.*

Proof. Lemma 5.4 gives the case $i = 0$, with $d_0 \geq 3$. By Theorem 5.3, each round raises the YES bound by h, the NO cost and the general upper bound by $h + 1$, and rank by a factor at most 2^h. In the NO case, $H_i \geq 3$ and $\mathcal{P}(R_i; \rho, H_i)$ give the hypotheses of (ii), which restores the deletion bounds at threshold ρ, so the induction continues. □

6 Proof of amplification under input deletion

We now prove Theorem 5.3, for the matrices N, U, W of Definition 6.2. By Lemma 6.6 and Theorem A.3, a depth-$(C-1)$ protocol for a near-full restriction of N begins with $\log s$ Alice bits, and the collision argument then gives Theorem 6.9, the lower bound for N after near-full deletion. For U, the template bounds and Lemma 6.10 and Theorem 6.12 give the stopping nodes of Lemma 6.13. The arrival bounds Lemmas 6.14 and 6.15 and Theorem 6.9 give Theorem 6.16: W has full cost when little column weight is deleted, and deletion bounds one bit below full cost. Finally, Theorem 6.17 for full products and Theorem 7.2 for explicit lists combine these two conclusions to restore the deletion bounds at full cost.

6.1 Widths and tolerances

We need enough coordinates to encode the graph, and enough interlaced copies that the lower bounds survive the restrictions left by a protocol's messages. Both cases use one recipe with two scales: ρ, the retained fraction of base rows for which hardness must survive between rounds, and s, the number of bins used to group inputs within a round.

Definition 6.1 (Amplification parameters)**.** Let $n = |V|$ be the number of vertices of the source graph. Choose powers of two s, ρ^{-1} with $n \leq \rho^{-1}$, $\rho \leq s^{-4}$, and $s \geq s_0$, the effective constant of Lemma F.1. Set

$$q_1 = s - 2, \qquad q_2 = s^2/\rho, \qquad w = s^8, \qquad Q = s^{24},$$
$$x = \rho^8, \qquad \eta = s^{-2}, \qquad \theta = (512Q)^{-1}, \qquad \delta_U = \eta/(16w).$$

The source construction of Definition 6.2 has q_1 template positions and q_2 indices, and the two further interlaces have widths w and Q. The deletion bounds of Lemma 6.6 for the matrices A, R_{temp} and N_{ref} of the source construction use row-weight threshold x, and the lower bound for the source matrix N survives deleting an η fraction of each input set's weight (Theorem 6.9). For the transpose W of the first further interlace (Definition 6.2), θ is the row-weight threshold of its deletion bounds, and restrictions retaining row weight at least θ and column weight at least $1 - \delta_U$ keep the full cost (Theorem 6.16).

In the full-product case, we take $L = \lceil \log n \rceil$, $\rho = 2^{-L}$, and $s = \mathrm{pow}_2(2^{16}(L+1))$. Lemma F.2 shows that these choices meet the constraints above once n exceeds an effective constant. In the explicit case, we take the least dyadic $s \geq s_0$ with $s^4 \geq n$, after enlarging s_0 if necessary, and set $\rho = s^{-4}$ (Lemma F.3). Section F verifies the numerical bounds (10) shared by both cases and the additional conditions of each case. All these choices are fixed across rounds.

6.2 Encoding the graph

We construct the source matrix of Section 2 from interlaced copies of the base matrix R and rows encoding the graph. In a NO instance, every assignment of the items defined below to bins has a *collision* (Lemma 6.3), that is, a coordinate at which one bin contains two active items and an inactive item. We show in the rest of the proof that a collision forces the extra bit of the NO case. Because of the template rows, defined below, every protocol cheaper than the NO-case cost must begin with $\log s$ bits from Alice, and these bits distinguish the s bins (Theorem A.3). The item rows encode the graph.

Definition 6.2 (Source construction)**.** Let $R : X \times Y \to \{0,1\}$ have positive normalized rational input weights, and let $s \geq 4$ be dyadic, meaning a power of two. Coordinates ranging over integer sets such as $[q_2]$ are uniform below, and product sets carry the product of the specified factor weights.

Active items. We repeat the graph in several groups so that the proof can discard groups that lose too many items. An item is a vertex or auxiliary label in one of $4s$ groups: $\mathcal{I} = [4s] \times (V \sqcup [s-3])$. Each edge occurrence in each group has a coordinate at which its two endpoints and the $s-3$ auxiliary items of that group are *active*. All other items are inactive there. Pad these $4s|E|$ coordinates with inactive coordinates to length $q_0 = \mathrm{pow}_2(4s|E|)$, then repeat each coordinate q_2/q_0 times, where q_2 is a dyadic multiple of q_0. Write $v_i(\alpha) \in \{0,1\}$ for the indicator that item i is active at index $\alpha \in [q_2]$, and π_α for the map numbering the at most $s-1$ active items $1, 2, \ldots$ in a fixed order on $\mathcal{I}$. Neither depends on any assignment of items to bins.

Template interlaces. Recall $q_1 = s-2$, which leaves two of the s values of Alice's last message in Proposition 6.4 for active and inactive items. Define

$$\Gamma = Y^{2s-1}, \qquad J = [q_1] \times X, \qquad S = J^{q_2}, \qquad J_2 = [q_2] \times \Gamma.$$

The template rows use the first q_1 positions of a list $\gamma \in \Gamma$, and we reserve its last $s+1$ positions for the item rows, since $q_1 + s + 1 = 2s - 1$. The two successive interlace matrices are

$$A(\gamma, (r,u)) = R(u, \gamma_r), \qquad R_{\mathrm{temp}}(z, (\alpha, \gamma)) = A(\gamma, z_\alpha).$$

Thus A has rows indexed by Γ and columns by J, and R_{temp} has rows indexed by S and columns by J_2. Moreover, A^{T} is a q_1-fold interlace of R with repeated column labels, and $R_{\mathrm{temp}}^{\mathsf{T}}$ is a q_2-fold interlace of A. A template label $z \in S$ is a tuple of pairs $z_\alpha = (r,u) \in J$, one for each index α. Against a list γ, the pair gives the entry $R(u, \gamma_r)$.

Rows and entries of B_{raw}. The matrix B_{raw} has template rows (p, z), indexed by a bin $p \in [s]$ and a template label, and item rows (i, u), indexed by an item and a base row:

$$\mathcal{K} = ([s] \times S) \sqcup (\mathcal{I} \times X).$$

We give the template rows most of the weight, so that deleting an η fraction of all row weight removes at most an η/λ fraction of theirs. Set $\lambda = 255/256$ and $\mu = 1/256$. Template rows have total weight λ, and vertex rows and auxiliary rows each have total weight $\mu/2 = 1/512$. Within each type, the bin or item is uniform and z or u has the weights of S or X. For a column $Z = (\xi, j_1, \ldots, j_s)$, where $\xi \in Y^{s+1}$ and $j_p = (\alpha_p, \gamma_p) \in J_2$, set

$$B_{\mathrm{raw}}((p,z), Z) = R_{\mathrm{temp}}(z, j_p),$$

$$B_{\mathrm{raw}}((i,u), Z) = \begin{cases} R(u, \xi_{\pi_\alpha(i)}), & \alpha_1 = \cdots = \alpha_s = \alpha,\ v_i(\alpha) = 1, \\ R(u, \xi_{s+1}), & \text{otherwise.} \end{cases}$$

Thus an item row uses its assigned position when the indices $\alpha_1, \dots, \alpha_s$ agree and the item is active there, and otherwise it uses the last entry of the shared list.

Two column families. The first column family $\Omega = Y^{s+1} \times J_2^s$ has product weights. The second column family is

$$\mathcal{D} = \{(\mathrm{tail}(\gamma), j, \dots, j) : j = (\alpha, \gamma) \in J_2\}, \qquad \mathrm{tail}(\gamma) = (\gamma_{q_1+1}, \dots, \gamma_{2s-1}),$$

with the weights induced from J_2. We give the two families the weights

Family	Weight	Role
Ω	λ	Force Alice's initial messages
$\mathcal{D}$	μ	Read the three tail positions of a collision.

The two families have distinct labels even where their columns coincide, so each keeps its own weight. At every bin p, the label j_p has the weights prescribed on J_2 in either family. We call the columns of the two families Ω-*inputs* and $\mathcal{D}$-*inputs*. After transposition they become rows, and we then also call them Ω-*rows* and $\mathcal{D}$-*rows*.

The transpose and two further interlaces. Transpose B_{raw}, and restrict the transpose to the rows from the first column family:

$$N = B_{\mathrm{raw}}^{\mathsf{T}}, \qquad N_{\mathrm{ref}} = N[\Omega, \mathcal{K}],$$

where the weights on Ω are normalized. The matrix N_{ref} is used only in the lower-bound analysis. Finally set

$$U = \mathsf{I}_w(N), \qquad W = U^{\mathsf{T}}, \qquad R^+ = (\mathsf{I}_Q(W))^{\mathsf{T}}.$$

As in Section 2, in a protocol below the NO-case cost the Ω-inputs force Alice's first $\log s$ bits to name a bin (Theorem A.3), and this fixes an assignment of the items to bins. Against a $\mathcal{D}$-input with label (α, γ), the rows of two items active at α and of one inactive item read three distinct positions of the same list γ. The interlace U, transposed to W, extends the lower bound for N to restrictions retaining only row weight θ. For W we obtain the deletion bounds only one bit below full cost (Theorem 6.16), and the final interlace restores that bit.

Throughout, we write $d = \mathrm{D}(R)$. With the parameters of Definition 6.1, we define the costs and the threshold used in the lower-bound analysis (in the explicit case, we drop the subscript g that Section 8 attaches to its parameters):

$$H_1 := d + \log s, \qquad H_2 := H_1 + \log q_2, \qquad C := H_2 + \log s + 1, \qquad x_{\mathrm{cap}} := 2s\rho. \tag{1}$$

In the NO case, we will show $\mathrm{D}(A) \geq H_1$, $\mathrm{D}(R_{\mathrm{temp}}) \geq H_2$ and $\mathrm{D}(N) \geq C$. The collision bound of Lemma 6.6 requires the q_1 template positions and the three tail positions to retain base row weight at least x_{cap} within their blocks.

The following example shows the $s + 1$ positions exposed by a collision. Fix $s = 4$ and a coordinate α at which u, v are active and h, from another group, is inactive. Suppose these three items lie in the same bin, with $\pi_\alpha(u) = 1$ and $\pi_\alpha(v) = 2$. On a $\mathcal{D}$-input, their rows use $\gamma_3, \gamma_4, \gamma_7$, respectively. Together with the bin's $s - 2 = 2$ template positions, they expose $s + 1 = 5$ positions of the same list.

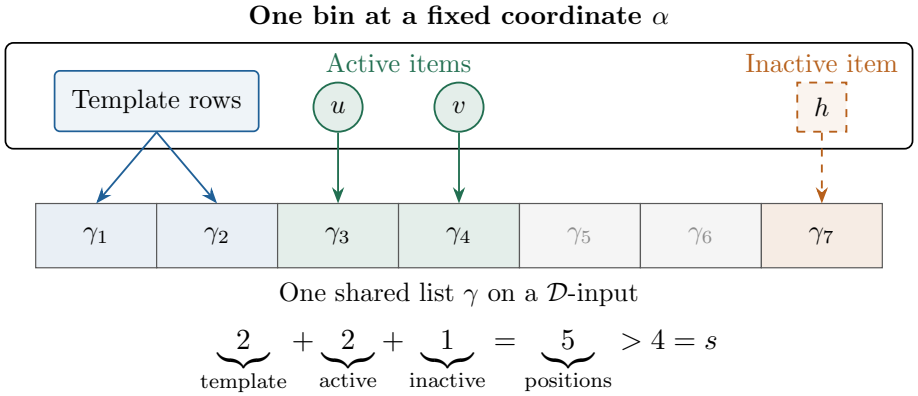


Figure 2: A collision exposes $s+1$ positions, illustrated for $s=4$. On a $\mathcal{D}$-input with label (α, γ), the template rows in the depicted bin use positions $1, 2$, the active items u, v use positions $3, 4$, and the inactive item h uses position 7. Activity and the positions read are determined by α, independently of the bin assignment. Only the positions used by the displayed rows are highlighted, and other items are omitted.

The next lemma shows that, for a NO graph, every assignment of the items remaining after a few deletions still has collisions as in Figure 2 at a constant fraction ζ of the coordinates.

Lemma 6.3. *A YES graph admits an assignment of $\mathcal{I}$ to $[s]$ with at most one active item in each bin at every coordinate. For a NO graph, delete fewer than $(s/2)\delta_v n$ vertex items and fewer than $s/2$ auxiliary items, where*

$$\delta_v = \frac{\gamma|E|}{2\Delta n}, \qquad \zeta = \frac{\gamma}{24}.$$

Every assignment of the remaining items to the s bins then has at least ζq_2 coordinates at which one bin contains two active items and an inactive item.

We prove the collision guarantee in Section 6.3. First, the YES assignment gives a short protocol. With $d = \mathrm{D}(R)$ and C as in (1), its cost is at most $C - 1$.

Proposition 6.4. *If $\mathrm{D}(R) \le y$, then $\mathrm{D}(N) \le y + 2\log s + \log q_2 + 1$. For a YES graph, the final $+1$ can be omitted. Both bounds remain valid after restricting either input set.*

Proof. For a YES graph, fix a public bin assignment χ from Lemma 6.3. We describe a protocol for B_{raw}. Alice holds a template row (p, z) or an item row (i, u), and Bob holds $Z = (\xi, j_1, \dots, j_s)$ with $j_{p'} = (\alpha_{p'}, \gamma_{p'})$ for each $p' \in [s]$. They exchange:

Alice	bin p (use $p = \chi(i)$ for an item row)	$\log s$ bits,
Bob	α_p	$\log q_2$ bits,
Alice	one of q_1 template positions, active, inactive	$\log s$ bits.

The third message has $q_1 + 2 = s$ possibilities. For a template row, Alice writes $z_{\alpha_p} = (r, u)$ and sends code r, and the answer is $R(u, (\gamma_p)_r)$, where Bob knows $(\gamma_p)_r$. For an item row, she sends active precisely when $v_i(\alpha_p) = 1$. Bob can test from his own input whether $\alpha_1 = \cdots = \alpha_s$. If so and the code is active, Bob identifies the active item in bin p, unique by the choice of χ, and uses $\xi_{\pi_{\alpha_p}(i)}$. Otherwise he uses ξ_{s+1}. For both row types, the players then finish by running a cost-y protocol for R once.

For an arbitrary graph, we use any public assignment and replace the active code by the s possible values of $\pi_{\alpha_p}(i)$. There are now $q_1 + s + 1 = 2s - 1$ message values, so the third message costs one additional bit. These rules apply to Ω-inputs and $\mathcal{D}$-inputs alike, since both have the form $(\xi, j_1, \ldots, j_s)$. Exchanging the players gives the protocol for $N = B^{\mathsf{T}}_{\mathrm{raw}}$, and restriction preserves its cost. □

6.3 A collision survives the conversation

The NO-case lower bound for N uses the NO promise only through Lemma 6.3. We need that lemma's collision count for every assignment of the retained items, because Alice's first messages fix the assignment before any index is revealed.

Proof of Lemma 6.3. For a YES graph, let $c : V \to [3]$ be a proper 3-coloring, which exists by Theorem 5.2. Assign each vertex v to bin $c(v)$, and assign the $s-3$ auxiliary labels to the remaining bins. Every edge coordinate activates that edge's two endpoints and those auxiliaries in one group, so all its active items occupy distinct bins. Padded coordinates activate no items.

Now suppose that G is NO. By the deletion hypotheses in Lemma 6.3, fewer than $s/2$ groups lose an auxiliary item and fewer than $s/2$ groups lose more than $\delta_v n$ vertices. Hence more than $3s$ groups retain every auxiliary item and all but at most $\delta_v n$ vertices.

Fix such a group. If two auxiliaries share a bin, they are active together at every edge coordinate. Otherwise the auxiliaries occupy $s-3$ distinct bins, leaving three free bins. Color the retained vertices in those free bins by their bins, and extend the coloring arbitrarily to all other vertices. The argument below works for every such extension. At least $\gamma|E|$ edge occurrences are monochromatic by the NO promise, and at most $\Delta\delta_v n = \gamma|E|/2$ meet a deleted vertex. Every remaining monochromatic edge occurrence gives two active items in one bin. If an endpoint lies in the bin of an auxiliary, it shares that bin with the auxiliary. Otherwise both endpoints lie in free bins and, having the same color, share a bin. In both cases this group has at least $\gamma|E|/2$ edge coordinates with two active items in one bin.

It remains to find an inactive item in the same bin. For each counted coordinate, we choose one bin containing the two active items. We discard a group if one of its chosen bins contains retained items of that group only, and call such a bin exclusive to the group. A bin is exclusive to at most one group, since it contains items of that group, so at most s groups are discarded. In each of the more than $2s$ remaining groups, every chosen bin also contains a retained item from another group. That item is inactive at the chosen coordinate, which activates items of one group. We have therefore found more than $\gamma s|E|$ coordinates with a collision. Since $q_0 < 8s|E|$, their fraction among the q_0 padded coordinates exceeds $\gamma/8 > \zeta = \gamma/24$. Repeating each coordinate q_2/q_0 times preserves this fraction. □

The next estimate is used to prove the template-layer deletion bounds below for full-product interlaces. For explicit lists we instead use Lemma 6.10 and Theorem 6.12. The proof of the deletion bounds in both cases is in Section A.4, and the collision argument uses only the resulting bounds.

Lemma 6.5. *Suppose $\mathcal{P}(M; x, H)$ holds for a rationally weighted matrix M, where $H \geq 3$. Let $u \geq 16$ be dyadic with $ux \leq 1$. Every restriction of a full-product interlace of M retaining at least $9u/16$ blocks, each of base row weight at least ux, and a column set of weight at least $2^{-u/1024}$, has cost at least $H + \log u$.*

The proof is in Section A.3, on p. 54.

For the collision argument, we need deletion bounds for A, R_{temp} and N_{ref}, and a *collision bound*, which says that the template positions together with the three positions of a collision require at least $H_1 + 1$ bits.

Lemma 6.6 (Template-layer deletion bounds)**.** *Consider either case of Theorem 5.3: the full-product construction of Definition 6.2 with the parameters of Lemma F.2, where $\rho = 2^{-L}$, or the explicit construction of Definition 8.3 with the parameters of Lemma F.3, where $\rho = \rho_{\mathrm{g}} = 1/\mathsf{N}$ and the subscript* g *is suppressed. Suppose $\mathrm{D}(R) = d \geq 3$ and $\mathcal{P}(R; \rho, d)$, let $x = \rho^8$ be the block threshold, and define $H_1, H_2, C, x_{\mathrm{cap}}$ by* (1)*. Then*

$$\mathcal{P}(A; x, H_1), \qquad \mathcal{P}(R_{\mathrm{temp}}; x, H_2), \qquad \mathcal{P}(N_{\mathrm{ref}}; x, C - 1).$$

In the interlace of R on Γ (on Γ_{g} in the explicit case), the first q_1 positions together with any three distinct tail positions require at least $H_1 + 1$ bits whenever every position retains base row weight at least x_{cap} within its block and the retained set of list labels has weight greater than $1/2$.

The proof is in Section A.4, after the initial-matrix proof (p. 55).

Three proofs below (Lemmas 6.8 and 6.13 and Claim B.2) follow retained input sets along a protocol prefix and count them where the prefix stops. A mark lets a set count twice until it splits into two sets. In the collision application, the mark records the extra bit forced by a collision.

Lemma 6.7 (Kraft count)**.** *Let q be a power of two, and let V be a set of nodes of a binary protocol tree, none an ancestor of another, all at depth at most $\log q$. For $v \in V$ at depth r_v put $\kappa_v = q2^{-r_v}$. Suppose each node $v \in V$ holds h_v objects, $z_v \leq h_v$ of them marked, and $\sum_{v \in V}(h_v + z_v) > q$. Then some v has $h_v + z_v > \kappa_v$. For this v, if $\kappa_v \geq 2$ then $h_v \geq \kappa_v/2 + 1$, and if $\kappa_v = 1$ then $h_v \geq 2$ or $h_v = z_v = 1$.*

Proof. Kraft's inequality gives $\sum_v \kappa_v \leq q$, so some v has $h_v + z_v > \kappa_v$, and $z_v \leq h_v$ gives $2h_v > \kappa_v$. If $\kappa_v \geq 2$, then κ_v is even and $h_v \geq \kappa_v/2 + 1$. If $\kappa_v = 1$, then $h_v + z_v \geq 2$, so $h_v \geq 2$ or $h_v = z_v = 1$. □

In the applications, the objects are retained input sets. We follow these sets through the bits of one player, starting from sets of weight at least t, by the following *tracking rule*, which keeps every tracked set at weight at least $t2^{-r}$ after r such bits. An unmarked set follows a child retaining at least half its weight. After r bits, a marked set splits into two unmarked sets if both children retain weight at least $t2^{-r-1}$. Otherwise it follows a child meeting that threshold, which exists since one child retains half its weight, and keeps its mark. Thus every mark splits at most once. Before it splits, a marked set loses less than $\sum_{r \geq 0} t2^{-r-1} = t$ of its weight. A split replaces the contribution 2 of a marked set to $\sum_v (h_v + z_v)$ by $1 + 1$. Hence $\sum_v (h_v + z_v)$ equals the number of initial sets plus the number of initial marks.

Lemma 6.8 (Following collision witnesses)**.** *Let $q \geq 2$ be dyadic, let $0 \leq b < q$ be an integer, and let $x > 0$ satisfy $2qx \leq 1/8$. Suppose a protocol with at most $H + \log q$ bits left has $q - b$ indices, each with a retained set of list inputs of weight at least $3/4$. Each set consists of inputs with that fixed index. Mark $w_0 > b$ of these sets. Assume the following lower bounds, where all set weights are computed from the original input weights.*

(i) Until Alice speaks, a set of retained lists at one index with weight at least x, against the unchanged set of Alice's inputs, costs at least $H > 0$.

(ii) At the first Alice bit, one child retains a set of Alice's inputs on which, for every dyadic $u \geq 2$ with $ux \leq 1$, retained sets at any $u/2 + 1$ distinct indices, each of weight at least ux, cost at least $H + \log u$.

(iii) A marked set retaining weight greater than $1/2$ *costs at least* $H+1$ *before Alice speaks.*

Then no such protocol exists.

Proof. We track the retained list sets of each index through Bob bits by the tracking rule with $t = 2qx$, and we stop at the first Alice bit. A marked set loses less than $2qx \le 1/8$ before it splits, so an unsplit mark retains weight greater than $1/2$. By (i), a tracked set cannot end at a leaf, nor reach Bob depth $\log q + 1$: it would still have weight at least x but only $H-1$ bits left. The stopping nodes therefore satisfy the hypotheses of Lemma 6.7, where v holds h_v tracked sets and z_v unsplit marks, and $\sum_v (h_v + z_v) = q - b + w_0 > q$. Let v be the node given by that lemma, with $\kappa = \kappa_v$. Its sets have distinct indices, because descendants of a split never rejoin, and each has weight at least $2\kappa x$.

If $\kappa \ge 2$, then (ii) with $u = \kappa$ requires $H + \log \kappa$ bits after Alice's bit, but only $H + \log \kappa - 1$ remain. If $\kappa = 1$ and $h_v \ge 2$, then (ii) with $u = 2$ and the available weights $2x$ requires $H+1$ bits after Alice's bit, but only $H-1$ remain. If $h_v = z_v = 1$, then by (iii) the unsplit marked set needs $H+1$ bits at v, where only H remain because v has depth $\log q$. □

Theorem 6.9 (Source hardness after deletion)**.** *Use either the full-product construction of Definition 6.2 or the explicit-list construction of Definition 8.3, with a NO graph and* $\mathrm{D}(R) = d \ge 3$*, and define* $H_1, H_2, C, x_{\mathrm{cap}}$ *by* (1)*. Suppose the deletion bounds and the collision bound stated in Lemma 6.6 hold. Assume the numerical inequalities* (10) *of Lemma F.1, that the weights on* S *give exact* 8*-wise independence with the prescribed weights on* J *as marginal, and that the bin tuple* $(j_1, \dots, j_s)$ *of the* Ω*-inputs has exactly* 8*-wise independent weights with the prescribed* J_2 *marginal. Then* $\mathrm{D}(N) = C$*, and every restriction retaining at least* $1-\eta$ *of each input set's total weight has cost at least* C*. In either case of Theorem 5.3 with* $\mathcal{P}(R; \rho, d)$*, these hypotheses follow from Lemma 6.6, Lemma F.1, and the exact independence of the constructed weights.*

Proof. In outline, Alice's first $\log s$ bits assign the items to bins. The collisions of this assignment give the marked sets of Lemma 6.8, and Lemma 6.6 gives its hypotheses (i)–(iii). Lemma 6.8 then rules out depth $C-1$.

Suppose a restriction of N retaining at least $1-\eta$ of each input set's weight has a protocol of depth at most $C-1$. We transpose it to a protocol of the same depth for the corresponding restriction of B_{raw}. Each template bin has row weight λ/s, so it loses at most $\eta s/\lambda < 1/1000$ of its total S weight. The first column family Ω retains normalized weight at least $1 - \eta/\lambda > 99/100$. Template entries do not depend on the component ξ of a column. We therefore project the retained Ω-inputs onto their bin tuple $(j_1, \dots, j_s)$ (in the explicit case, onto their $\mathcal{J}_{\mathrm{g}}$ label) and choose one retained completion for each image label. These representatives form a restriction of an interlace of R_{temp} on the template rows, with column weight at least $99/100$ (Observation 4.5).

We apply Theorem A.3 to the protocol inherited on these representatives, without pruning its tree. We take $M = R_{\mathrm{temp}}$, $d = H_2$, $u = s$, x in place of its ρ, and row threshold $999/1000$ in place of its ξ. Its hypotheses follow from the deletion bounds $\mathcal{P}(R_{\mathrm{temp}}; x, H_2)$, the depth $C - 1 = H_2 + \log s$, and the numerical inequalities $999/1000 > 2sx$ and $999/1000 - 2(s-1)x > 998/1000$, with $99/100 > 18/25$. By that theorem, every binary history of length below $\log s$ is reached by template inputs, and every such node is Alice's. They are actual nodes of the original protocol, so Alice sends the first $\log s$ bits on every retained input pair, including pairs with item rows or $\mathcal{D}$-inputs. At the resulting s nodes, the respective template bins retain sets $A_1, \dots, A_s \subseteq S$, each of weight greater than $998/1000$. We needed the representatives only to establish this prefix. From now on we use all originally retained columns, with $C - 1 - \log s = H_2$ bits left.

We next fix the item assignment before choosing an index. We set aside items whose rows retain less than half of their base row weight in the restriction. By the row weights in Definition 6.2,

each item type has total weight $1/512$, uniform over $4sn$ vertex or $4s(s-3)$ auxiliary items. Each set-aside item loses more than half its weight, and at most η is deleted, so we set aside fewer than $4096\eta sn$ vertex items and fewer than $4096\eta s(s-3)$ auxiliary items. These numbers are below the deletion thresholds of Lemma 6.3, because (10) gives $4096\eta sn \le \delta_v n/16$ and $4096\eta s(s-3) < s/2$, as checked in the proof of Lemma F.1. By averaging over these s nodes, we can assign every remaining item to a node that receives at least $1/(2s)$ of its base weight. The parameter choices give $4s^2\rho \le 1$, so

$$\frac{1}{2s} \ge 2s\rho = x_{\mathrm{cap}}.$$

Thus the assigned part retains at least x_{cap} of the item's base row weight. This assignment is fixed before the index is chosen. By Lemma 6.3, at least ζq_2 coordinates contain two active items and an inactive item in one bin. We assign each such coordinate to one node $j \in [s]$ at which its collision occurs and write $U_j \subseteq [q_2]$ for the resulting disjoint sets.

For each j, we discard two kinds of indices α: those at which the template labels in A_j project onto too little base row weight, and those at which too little $\mathcal{D}$-weight is retained. First, we discard α for A_j if one of the q_1 blocks of the J-projection of A_j at α retains base row weight less than x_{cap}. Since $x_{\mathrm{cap}} \le 1/8$ by (10) and each block has weight $1/q_1$ in J, the J-projection then misses more than $7/(8q_1)$ of the weight of J. If there were at least $8q_1$ such indices, we could choose $8q_1$ of them and let Z count those at which z_α lies outside the projection at α, for $z \in S$ drawn from its weights. Then Z would satisfy $\mathbb{E}Z > 7$ and, by pairwise independence, $\operatorname{Var} Z \le \mathbb{E}Z$. Every label in A_j has $Z = 0$, so Chebyshev's inequality would give $\Pr(A_j) \le \Pr(Z = 0) < 1/7$, contrary to $\Pr(A_j) > 998/1000$. Thus fewer than $8q_1$ indices are discarded for A_j. Second, we discard α if the $\mathcal{D}$-inputs with index α retain less than $3/4$ of their weight. The $\mathcal{D}$-inputs have total weight μ, and their index is uniform on $[q_2]$. Each such α therefore loses more than $\mu/(4q_2)$ of the column weight. Since the restriction deletes at most η of the column weight, we discard fewer than $4\eta q_2/\mu$ such indices. Let B_j be the union of these two sets of discarded indices and put $b_j := |B_j|$. The inequalities in (10) give

$$\sum_{j=1}^{s} b_j < 8sq_1 + 4s\eta q_2/\mu < \zeta q_2/4.$$

Since $\sum_j |U_j| \ge \zeta q_2$, some j satisfies $|U_j| > 4b_j$. Fix this j, write $b := b_j$, and define $w_0 := |U_j \setminus B_j|$. Then $w_0 \ge |U_j| - b > 3b$.

Fix a collision index $\alpha \in U_j \setminus B_j$. Let i_1, i_2 be two active items and i_3 an inactive item in one bin at α, as given by Lemma 6.3. We choose retained template rows for the first q_1 projected blocks and retained rows of i_1, i_2, i_3. By the entry rule in Definition 6.2, against a $\mathcal{D}$-input with label (α, γ) these rows read the following entries, where u is the base row of the respective row:

$$R(u, \gamma_\ell) \quad (\ell \in [q_1]), \qquad R(u, \gamma_{q_1 + \pi_\alpha(i_1)}), \qquad R(u, \gamma_{q_1 + \pi_\alpha(i_2)}), \qquad R(u, \gamma_{q_1+s+1}).$$

The three tail positions are distinct by the collision definition and the tail convention in Definition 6.2. At each of these positions, the selected rows have base row weight at least x_{cap}, for template positions because $\alpha \notin B_j$, and for item positions because of the threshold in the item assignment above. The selected rows depend only on the projected labels and not on γ, so they form a restriction of the interlace of R on Γ in Lemma 6.6. The collision bound of Lemma 6.6 therefore requires $H_1 + 1$ bits whenever the retained $\mathcal{D}$-inputs at index α have weight greater than $1/2$ in Γ.

We now apply Lemma 6.8 with $q = q_2$ and $H = H_1$. Its $q_2 - b$ indices are those outside B_j, and its $w_0 > 3b$ marked sets are those at the indices in $U_j \setminus B_j$. The numerical bounds (10) give $q_2 x \le 10^{-4}$, hence $2q_2 x \le 1/8$. We now verify its hypotheses (i), (ii) and (iii).

Until Alice speaks again, the set A_j of template labels is unchanged. Its projection onto each coordinate α has weight at least $\Pr(A_j) > 998/1000 > 3/5$ in J, so (i) follows from $\mathcal{P}(A; x, H_1)$ (Observation 4.5). At the first Alice bit, one child retains a part of A_j of weight greater than $499/1000 > 9/25$. After transposition, the entries of that child at indices with retained list sets are $A(\gamma, z_\alpha)$, so they form the interlace of A with those sets of lists as row blocks and tuple list S. Since S is exactly 8-wise independent and this child retains template-label weight greater than $9/25$, Lemma 6.11 with $M = A$, $d = H_1$ and $2^a = u$ gives cost at least $H_1 + \log u$ for any $u/2 + 1$ retained index sets of weight at least ux. This gives (ii) for the template labels retained at that child. The collision bound established above gives (iii), since Alice's inputs are unchanged before that bit. Hence Lemma 6.8 excludes depth $C - 1$. The upper-bound protocol of Proposition 6.4 has depth C, which proves equality. □

6.4 Lower bounds after restricting tuple inputs

Restricting the tuple inputs can destroy the independence among their coordinates (Section 2), so we use independence only under the original tuple weights. For $0 < v, \tau \le 1$, the next lemma shows that, after conditioning on any event of weight at least v, every joint marginal on at most eight coordinates is within total variation distance $\varepsilon(v, \tau)$ of the product weights, provided these coordinates avoid a set of at most $e(v, \tau)$ exceptional coordinates. Set

$$\varepsilon(v, \tau) := \min\{\tau, 2^{-30} v^{1/256}\},$$
$$e(v, \tau) := 8 \left\lceil 2^{20} \varepsilon(v, \tau)^{-2} v^{-1/512} \right\rceil, \qquad e(v) := e(v, 1).$$

We use independence order $8192 = 8 \cdot 1024$ so that any 1024 events on pairwise disjoint sets of at most eight coordinates are independent under the original weights, as a 1024th-moment bound requires.

Lemma 6.10 (Conditioning outside exceptional coordinates)**.** *Let ν be the normalized weights of an indexed exactly 8192-wise independent tuple list (Definition 4.4) on Z^q with marginal μ, and let $0 < v, \tau \le 1$. For every event Y of labels with $\nu(Y) \ge v$, there is a set $E \subseteq [q]$, with $|E| \le e(v, \tau)$, such that*

$$d_{\mathrm{TV}}\big((\nu \mid Y)_J, \mu^J\big) \le \varepsilon(v, \tau) \qquad (J \subseteq [q] \setminus E,\ |J| \le 8).$$

Here total variation is the maximum absolute difference between the two distributions' probabilities of an event. The reduction need not compute the event or the exceptional set.

The proof is in Section C.1.

Combined with Observation 4.5, Lemma 6.10 bounds the column weight of projected restrictions after conditioning. If $K \subseteq Y$ and J consists of at most eight coordinates outside E, then the inverse image of $\pi_J K$ contains K, so

$$\mu^J(\pi_J K) \ge (\nu \mid Y)_J(\pi_J K) - \varepsilon(v, \tau) \ge (\nu \mid Y)(K) - \varepsilon(v, \tau).$$

Thus the projected restriction of Observation 4.5 has column weight at least $(\nu \mid Y)(K) - \varepsilon(v, \tau)$, and we need no product assumption on $\nu \mid Y$.

The next two results apply to full products and to explicit lists. Both are proved in Section C.3 (p. 64).

Lemma 6.11 (Surviving blocks in a tuple list)**.** *Suppose $d \ge 2$ and $\mathcal{P}(M; x, d)$. Let an interlace of M use a tuple list that is exactly 8-wise independent with the column weights of M as marginal. If*

an integer $a \geq 1$ satisfies $2^a x \leq 1$, and a restriction retains at least $2^{a-1}+1$ blocks, each of base row weight at least $2^a x$, together with a label set of total weight at least $9/25$, then the restricted interlace has cost at least $d + a$.

After a protocol prefix, the retained tuple labels may weigh less than the $9/25$ required by Lemma 6.11. In that case we condition on the event Y of retained labels and discard its exceptional coordinates (Lemma 6.10).

Theorem 6.12 (Surviving blocks after conditioning)**.** *Suppose $d \geq 2$ and $\mathcal{P}(M; x, d)$, and let an interlace of M use an exactly 8192-wise independent tuple list with weights ν and the column weights of M as marginal. Let $\nu(Y) \geq v > 0$, choose the exceptional set from Lemma 6.10 with any error bound $0 < \tau \leq 1$, and let an integer $a \geq 1$ satisfy $2^a x \leq 1$. If a restriction retains at least $2^{a-1}+1$ blocks indexed outside that exceptional set, each of base row weight at least $2^a x$, and a column set $K \subseteq Y$ with $(\nu \mid Y)(K) \geq 1/2$, then its cost is at least $d + a$.*

6.5 Where displaced inputs can go

We now prove a lower bound for restrictions of $U = \mathsf{I}_w(N)$ that keep almost all rows but few tuple inputs, following the copies of N through a hypothetical short protocol as in Section 2. The argument applies to both constructions. Each copy has row weight $1/w$ in U, but we measure row weight in units of one complete copy, and losses among a copy's Ω-rows in weights normalized within Ω. Recall that Ω has weight $\lambda = 255/256$ and $\mathcal{D}$ has weight $\mu = 1/256$.

Common hypotheses. In this section we use only the following hypotheses, where H_1, H_2, C are as in (1) and $H_1, H_2 \geq 2$.

(H1) Each copy satisfies the deletion bounds $\mathcal{P}(A; x, H_1)$, $\mathcal{P}(R_{\text{temp}}; x, H_2)$ and $\mathcal{P}(N_{\text{ref}}; x, C-1)$.

(H2) The matrix N has exact cost C, including on every restriction retaining $1 - \eta$ of each input set's total weight.

(H3) The weights on S give exact 8192-wise independence across its J-valued coordinates, and the bin tuple in Ω has the same independence order and the prescribed J_2 weights as marginal. The tuple weights ν of U have the original $\mathcal{K}$ weights at each coordinate and arise from three mutually independent label sequences, namely type/bin/item labels, base labels with the X weights, and template labels with the S weights. Each sequence has a value at every copy and is exactly 8192-wise independent across copies. At copy i, with template label z_i and base label u_i, the decoder outputs the template row (p, z_i) if the first sequence gives a template type with bin p, and the item row (a, u_i) if it gives item a. Thus, on any set of at most 8192 copies, the decoded labels have the product of the original $\mathcal{K}$ weights. The tuple coordinates $z_{i,\alpha}$, indexed by copy i and coordinate α, are also exactly 8192-wise independent: group at most 8192 coordinates by copy, then use independence between copies and among the tuple coordinates within each template input. By independence, fixing some of the three complete sequences leaves the remaining ones with their original weights. The proof of Lemma 6.14 fixes the first two sequences by averaging and then conditions on the retained template labels.

(H4) The widths and deletion thresholds satisfy the numerical bounds in (10): each quantity in the first column is below the corresponding bound in the last column. These bounds provide enough coordinates to encode the graph and ensure that, after deleting inputs and setting aside exceptional coordinates, enough inputs remain for the lower-bound arguments.

In both cases, (H1) follows from Lemma 6.6 and (H2) from Theorem 6.9. Lemma F.1 verifies (H4) for the parameter choices of Definition 6.1. For (H3) with full products, we take all complete sequences, including values ignored by the decoder, with full product weights. Decoding then induces the full product of the original weights on $\mathcal{K}^w$, so pulling back along the decoder preserves the weights of retained sets and the entries and depth of protocols. This representation is used only in the analysis and is never enumerated. In the explicit case the complete sequences are the actual indexed input labels (Definition 8.3 and Lemma 8.2).

By Lemma 6.10 with $v = \theta$ and $\tau = \eta/2$, conditioning on a tuple event of weight at least θ makes at most $e(\theta, \eta/2)$ copies exceptional. Write $E_{\mathrm{exc}} = 2e(\theta, \eta/2)$, so that at most $E_{\mathrm{exc}}/2$ copies are exceptional. The reduction need not identify them.

We stop each root-to-leaf path of the protocol just before Bob's first bit or just after Alice's $(\log w)$-th bit, whichever comes first. Up to these *stopping nodes* only Alice has spoken, so each row follows the same path for every remaining column. By following the Ω-rows, we assign distinct stopping nodes to all but at most E_{exc} copies (Lemma 6.13).

Lemma 6.13. *Under the common hypotheses above, suppose a restriction of U retains row weight at least $1 - \delta_U$ and a tuple event Y' of original weight at least θ, and has a protocol of depth at most $C + \log w - 1$. Apply Lemma 6.10 to Y' with $v = \theta$ and error bound $\eta/2$, and set aside the at most $E_{\mathrm{exc}}/2$ exceptional copies. At an Alice node of depth r, put $\kappa = w2^{-r}$. Follow each tracked set of Ω-rows of one remaining copy, of normalized weight at least $2\kappa x$, into every child in which it retains normalized weight at least κx. An unsplit row set contributes one. Once a row set first splits into two children meeting the threshold, each descendant contributes one-half. We add these contributions at each node. They count row sets, not input weight.*

Stop at the first node where Bob speaks or at depth $\log w$, whichever comes first. The latter stopping condition applies regardless of the next speaker. Paths that end earlier carry no row sets, and we extend them to depth $\log w$ by fixed Alice bits. At a node where Bob first speaks with $\kappa \geq 2$, the total contribution is at most $\kappa/2$. A node at depth $\log w$ has at most one row set. At least $w - E_{\mathrm{exc}}$ nodes at that depth are reached by unsplit row sets from distinct copies. Assign each such node to the copy of its row set. Over all other stopping nodes, the sum of κ is at most E_{exc}.

Proof. Since $\delta_U = \eta/(16w)$, initial deletion removes at most $w\delta_U = \eta/16$ in copy units, hence at most $\eta/16$ of any single copy, so every copy starts with normalized weight on Ω at least $1 - \eta/(16\lambda) > 2wx$, by (10). A row set always has a child retaining at least half its input weight. As with the marks in the tracking rule after Lemma 6.7, the first split of a row set preserves the total count, since $1 = \frac{1}{2} + \frac{1}{2}$, and further splits can only increase it. The descendants of one row set lie in disjoint subtrees, so each node holds at most one row set from each copy. Each of the at least $w - E_{\mathrm{exc}}/2$ copies outside the exceptional set starts with an unsplit row set, so the total count at the stopping nodes is at least $w - E_{\mathrm{exc}}/2$.

Until Bob speaks, the tuple set is the fixed Y'. For every non-exceptional copy i, the projection of Y' onto coordinate i has original $\mathcal{K}$ weight at least $1 - \eta/2 > 3/5$ (the display after Lemma 6.10 with $K = Y = Y'$), and together with the Ω-rows of copy i it gives a restriction of N_{ref} (Observation 4.5). Since a row set has row weight $\geq 2x$, the bound $\mathcal{P}(N_{\mathrm{ref}}; x, C-1)$ therefore requires at least $C - 1 > 0$ further bits, and no row set reaches a leaf. After we pad paths without row sets, every root-to-leaf path has one stopping node, so the sum of κ over these nodes is w.

At a node where Bob first speaks with $\kappa \geq 2$, suppose there were at least $\kappa/2 + 1$ row sets. They come from distinct copies whose coordinates lie outside the exceptional set for the same event Y'. The Bob child retaining more column weight has conditional weight at least one-half. We apply Theorem 6.12 with $M = N_{\mathrm{ref}}$, $d = C - 1$, $a = \log \kappa$, and K the column set of this child. Here a is an integer since w is a power of two, $\kappa x \leq wx \leq 1$, and each row set has weight at least $2\kappa x \geq \kappa x$.

The child requires $C - 1 + \log \kappa$ bits, but only $C + \log \kappa - 2$ remain. Consequently there are at most $\kappa/2$ row sets, and their total contribution is at most $\kappa/2$.

At a node at depth $\log w$, two row sets of weight at least $2x$ would require C bits by the same theorem with $a = 1$ and $K = Y'$, but only $C - 1$ bits remain. Thus a node at depth $\log w$ has at most one row set. If the set is unsplit, we assign the node to the copy supplying it. Otherwise, the set's contribution is at most one-half. Let c be the number of assigned nodes. The total contribution at every other stopping node is at most half its value of κ. Since each assigned node has $\kappa = 1$ and the values of κ sum to w,

$$w - E_{\text{exc}}/2 \le c + \tfrac{1}{2}(w - c), \qquad \text{hence} \qquad w - c \le E_{\text{exc}}.$$

An unsplit row set follows only one path, so the assigned nodes correspond to distinct copies. The sum of κ over all other stopping nodes is $w - c \le E_{\text{exc}}$. □

A surviving row of an assigned copy that misses its copy's assigned node reaches either an unassigned stopping node or a node assigned to another copy. Lemma 6.14 bounds the weight arriving at unassigned nodes, and Lemma 6.15 the weight arriving at nodes assigned to other copies. Their proofs are in Sections B.2 and B.3, respectively.

Lemma 6.14. *For the protocol and stopping nodes of Lemma 6.13, the total $\mathcal{D}$-weight, over all copies, reaching a stopping node of depth r where Bob first speaks before depth $\log w$ is less than $2s\kappa$, where $\kappa = w2^{-r}$. At a stopping node at depth $\log w$ this weight is less than $2s$. Weight is measured in units of one complete copy of N.*

Lemma 6.15. *In the same protocol, for each copy i with an assigned node, let ε_i be the normalized weight of copy i's Ω-rows missing at that node. For every other copy j with an assigned node, the $\mathcal{D}$-rows of copy j reaching the node assigned to copy i have total weight at most $\mu\delta_i$, where*

$$\delta_i = 2q_2x + 4\varepsilon_i.$$

The proof of Theorem 6.16 combines these bounds into the contradiction shown in Figure 3, in which P_i is the $\mathcal{D}$-weight of copy i surviving outside its assigned node and a_i is its initially deleted $\mathcal{D}$-weight.

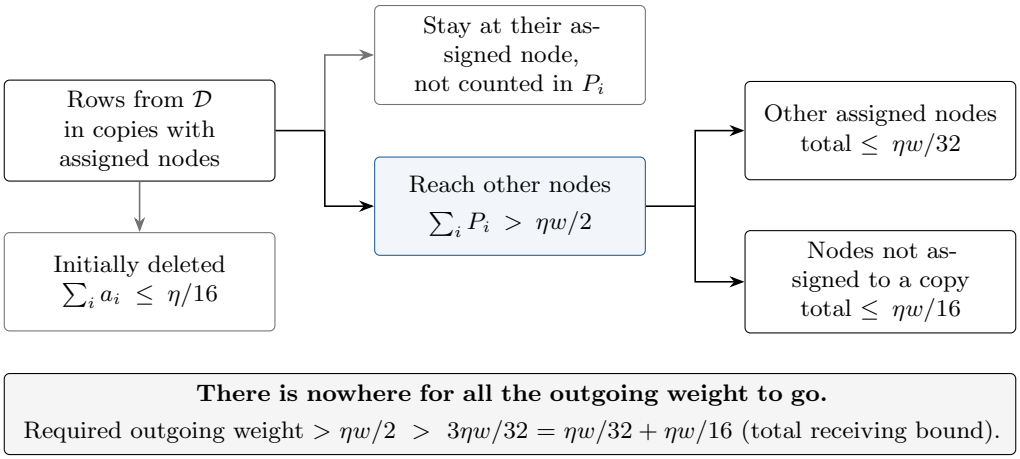


Figure 3: Destinations of the displaced $\mathcal{D}$-weight. Every surviving $\mathcal{D}$-row outside the node assigned to its copy must reach one of the two groups of stopping nodes on the right, whose combined bound on arriving weight is smaller than the outgoing weight. The lower group contains the nodes not assigned to copies, either where Bob first speaks before depth $\log w$ or at that depth. Rows deleted initially and rows staying at the node assigned to their copy do not count toward P_i. All weights are in complete-copy units, and arrow widths are not to scale.

The deletion bounds for W in the next theorem are one bit below its full cost, and the next subsection recovers that bit.

Theorem 6.16 (Lower bounds for W)**.** *Under the common hypotheses of Section 6.5,*

$$\mathrm{D}(W) = C + \log w, \qquad \mathcal{P}(W; \theta, C + \log w - 1).$$

Every restriction of W retaining row weight at least θ and column weight at least $1 - \delta_U$ has cost at least $C + \log w$.

Proof. Suppose a restriction of W with row weight $\ge \theta$ and column weight $\ge 1 - \delta_U$ has a protocol of depth $\le C + \log w - 1$. Transposing gives a restriction of U (recall $W = U^{\mathsf{T}}$) with a protocol of the same depth, retained row weight at least $1 - \delta_U$, and a retained tuple event Y' of weight at least θ. By Lemma 6.13 and $E_{\mathrm{exc}} < w/4$ from Lemma F.1, at least $w - E_{\mathrm{exc}} \ge 3w/4$ copies have assigned nodes, each with only $C - 1$ bits left. These copies lie outside the exceptional set, so conditioning on Y' with error bound $\eta/2$ gives each of them a column projection of original weight at least $1 - \eta/2$, which gives a restriction of N (Observation 4.5). By Theorem 6.9, the assigned node must miss more than η of the row weight in its assigned copy.

An unsplit set of Ω-rows discards less than $\sum_{r=0}^{\log w - 1} wx2^{-r} = 2(w-1)x$ on its path, since at depth r the discarded child retains less than $\kappa x = wx2^{-r}$. Initial deletion removes at most $\eta/16$ complete copies. Summing over copies with assigned nodes and using $w^2 x/\eta \le 1/(32\lambda)$ from (10), we obtain

$$\sum_i \varepsilon_i \le \frac{\eta}{16\lambda} + 2w(w-1)x \le \frac{\eta}{8\lambda}.$$

For these copies, the missing Ω-rows, initially deleted $\mathcal{D}$-rows, and $\mathcal{D}$-rows surviving elsewhere partition the part absent from their own assigned node. Thus $P_i + a_i + \lambda\varepsilon_i > \eta$, and

$$\sum_i P_i > \frac{3\eta w}{4} - \frac{\eta}{16} - \frac{\eta}{8} > \frac{\eta w}{2}.$$

We now show that the other stopping nodes cannot receive this much weight. The sum of κ over nodes not assigned to copies is at most E_{exc}, so Lemma 6.14 bounds their arriving weight by $2sE_{\mathrm{exc}} \le \eta w/16$. By Lemma 6.15, other assigned nodes receive at most

$$\mu w \sum_i (2q_2 x + 4\varepsilon_i) \le \eta w \left(\frac{1}{64} + \frac{1}{510}\right) < \frac{\eta w}{32},$$

using $4\mu w q_2 x \le \eta/32$ and the bound on losses within Ω above. Since only Alice has spoken, every surviving row reaches one stopping node. The total arriving weight is therefore less than $3\eta w/32$, contradicting $\sum_i P_i > \eta w/2$. This proves the restriction bound.

For the deletion bounds $\mathcal{P}(W; \theta, C + \log w - 1)$, we fix $j \in \{0, 1\}$ and transpose a restriction of W with row weight at least θ and column weight at least $\beta_j = (3/5)2^{-j}$. In U its tuple event Y has weight at least θ. We discard the $\mathcal{D}$-rows, leaving a restriction of the interlace of $N_{\mathrm{ref}} = N[\Omega, \mathcal{K}]$. Since the $\mathcal{D}$-rows have weight $1 - \lambda$, the retained weight on Ω, normalized within Ω and averaged over the copies, is at least $\beta'_j = (\beta_j - (1 - \lambda))/\lambda$. We put $\kappa = w2^{-j}$ and apply Lemma 6.10 to Y with $v = \theta$ and $\tau = 1$, which leaves at most $e(\theta)$ exceptional copies. Let m be the number of non-exceptional copies retaining normalized weight on Ω at least κx. The block count of Section 4 gives $w\beta'_j \le m + e(\theta) + w\kappa x$, and $wx \le 1/256$ and $4e(\theta)/w < 1/32$ (Lemma F.1) then give

$$\frac{m}{\kappa} \ge 2^j \beta'_j - wx - \frac{4e(\theta)}{w} \ge \frac{(3/5) - 4(1 - \lambda)}{\lambda} - \frac{1}{256} - \frac{1}{32} = \frac{44}{75} - \frac{1}{256} - \frac{1}{32} > \frac{1}{2}.$$

Since $\kappa/2$ is an integer, at least $\kappa/2 + 1$ blocks indexed outside the exceptional set remain. We apply Theorem 6.12 to N_{ref} with $a = \log \kappa$ and $K = Y$, which has conditional weight one. The parameter choices give $\kappa \ge 2$ and $\kappa x \le 1$. The lower bound is $C - 1 + \log \kappa = C + \log w - 1 - j$, which proves the deletion bounds for both values of j after transposition. Applied to the whole matrix, the restriction bound proved first gives $\mathrm{D}(W) \ge C + \log w$. Conversely, the player holding the copy index sends it in $\log w$ bits, and the players then run the cost-C protocol for N, giving equality. □

6.6 Restoring the bound for the next round

To use R^+ as the next base matrix, we need deletion bounds at its exact cost, whereas Theorem 6.16 gives them for W one bit lower. For full products, Theorem 6.17 recovers this bit from the full-cost bound of Theorem 6.16 for column weight at least $1 - \delta_U$. Its proof (p. 54) uses the induction of Section A.3.

Theorem 6.17 (Restoring the deletion bounds)**.** *Let W have $m \ge 1$ rows, $n \ge 1$ columns and $\mathrm{D}(W) \ge 3$. Let $\theta > 0$ and $0 < \delta \le 1/64$. Assume that every restriction retaining row weight at least θ and column weight at least $1 - \delta$ has cost at least $\mathrm{D}(W)$, and that $\mathcal{P}(W; \theta, \mathrm{D}(W) - 1)$ holds. Let $b \ge 1$ be an integer and $Q = 2^r$ a power of two with $Q\delta \ge 80b$ and $Q\theta \le 1/20$. Then $R^+ = (\mathsf{I}_Q(W))^{\mathsf{T}}$ satisfies*

$$\mathrm{D}(R^+) = \mathrm{D}(W) + r, \qquad \mathcal{P}(R^+; 2^{-b}, \mathrm{D}(W) + r).$$

Independently of the lower-bound hypotheses, $\mathrm{D}(W) \le y$ implies $\mathrm{D}(R^+) \le y + r$. The matrix R^+ has exactly n^Q rows and Qm columns.

The sample size in Theorem 9.4 grows with the rank of the matrix. The next lemma bounds the rank increase of one round in both cases.

Lemma 6.18. *In the construction of Definition 6.2, let* $h = 2\log s + \log q_2 + \log w + \log Q$ *and suppose* $|\mathcal{I}| \le 2sq_2$. *Then* $\operatorname{rank}(R^+) \le 2^h \operatorname{rank}(R)$. *The same bound holds if any product set of tuples is replaced by an indexed list of tuples from that set, or by a product of such lists.*

Proof. A width-t interlace stacks t row blocks, each obtained by selecting or repeating columns of its base matrix, so its rank is at most t times the base rank. The template rows form nested interlaces of widths q_1, q_2 and s, so the template part has rank at most $sq_1q_2 \operatorname{rank}(R)$. For each item i, the rows $\{i\} \times X$ have entries $R(u, \psi_i(Z))$, where the column $\psi_i(Z)$, namely $\xi_{\pi_\alpha(i)}$ or ξ_{s+1}, is independent of u, so they have rank at most $\operatorname{rank}(R)$. After the outer interlaces, rank subadditivity gives

$$\operatorname{rank}(R^+) \le wQ(sq_1q_2 + |\mathcal{I}|)\operatorname{rank}(R) \le wQsq_2(q_1+2)\operatorname{rank}(R) = 2^h \operatorname{rank}(R),$$

using $q_1 + 2 = s$. On indexed lists every block still selects or repeats base columns, and transposition preserves rank, so the count is unchanged. □

Proof of Theorem 5.3. As in Lemma 6.6, we suppress the subscript g in the explicit case. In both cases $h = 2\log s + \log q_2 + \log w + \log Q$. For n beyond the effective cutoff of Lemma F.2, that lemma gives $h = L + O_{\Delta,\gamma}(\log L)$ in the full-product case, and Lemma F.3 gives $h = 40\log s = 10\log \mathsf{N}$ in the explicit case.

We first prove (ii), so let the graph be NO, $\mathrm{D}(R) = d \ge 3$ and $\mathcal{P}(R; \rho, d)$. By Lemma 6.6, the deletion bounds for A, R_{temp} and N_{ref} and the collision bound hold. The construction has the exact marginals required by Theorem 6.9 (in the explicit case, those of S_{g} and $\mathcal{J}_{\mathrm{g}}$ in Definition 8.3), and Lemma F.1 gives the numerical bounds (10). Hence Theorem 6.9 applies, and N has cost C even after we delete at most η of each input set's weight. This gives (H1) and (H2) of Section 6.5, and Lemma F.1 gives (H4). For full products, (H3) is verified in Section 6.5. In the explicit case, Definition 8.3 provides the three independent sequences. Lemma 8.2 shows that decoding the nested template labels gives the exact 8192-wise independence of (H3), and that fixing the other two sequences leaves it unchanged. Thus, by (H1)–(H4), Theorem 6.16 gives $\mathrm{D}(W) = C + \log w$ and $\mathcal{P}(W; \theta, C + \log w - 1)$, and cost $\ge C + \log w$ on every restriction with row weight $\ge \theta$ and column weight $\ge 1 - \delta_U$. The two cases differ only in the restoration step.

- *Full-product case.* By Lemma F.2, $0 < \delta_U \le 1/64$, $Q\delta_U \ge 80L$, and $Q\theta = 1/512 < 1/20$, so Theorem 6.17 applies with $b = L$ and $\delta = \delta_U$, at threshold $2^{-L} = \rho$.
- *Explicit case.* Theorem 7.2 applies with $(\phi, \delta, Q) = (\rho, \delta_U, Q)$: its hypotheses on W are the bounds above, Lemma F.3 verifies its numerical hypotheses, and Definition 8.3 uses an exactly 8192-wise independent outer list with the column weights of W.

In both cases we restore the deletion bounds at threshold ρ and cost

$$C + \log w + \log Q = d + \underbrace{2\log s + \log q_2 + \log w + \log Q}_{h} + 1,$$

so $\mathrm{D}(R^+) = d + h + 1$ and $\mathcal{P}(R^+; \rho, d + h + 1)$.

For (i), on either promise, at each of the two outer interlaces the player holding the copy index sends it, in $\log w + \log Q$ bits, and the players then run the protocol of Proposition 6.4, of cost $\mathrm{D}(R) + 2\log s + \log q_2 + 1$. This gives $\mathrm{D}(R^+) \le \mathrm{D}(R) + h + 1$. On a YES graph Proposition 6.4 omits the final $+1$, giving $\mathrm{D}(R^+) \le \mathrm{D}(R) + h$. In the explicit case each input decodes locally to a full-product input with the same entry, and the protocol uses only decoded inputs, so it applies unchanged.

For (iii), the bound $q_2 \ge 4s|E|$ from (10) and $|E| \ge n/2$ imply $|\mathcal{I}| = 4s(n + s - 3) \le 2sq_2$, so Lemma 6.18 gives the rank bound in both cases. □

7 Interlacing on small input lists

7.1 Constructing exact marginals

The deterministic reduction replaces each full product Y^q by an explicit tuple list (Definition 4.4) whose induced weights on every set of at most 8192 coordinates equal the product weights.

We build the tuples one coordinate at a time. To each tuple we attach the zero–one vector indicating, for each set J of at most 8192 coordinates and assignment a on J, whether the tuple agrees with a. After each extension, we move weight along linear dependences among these vectors. Each move removes at least one tuple, keeps all weights nonnegative and leaves the joint weights on these coordinate sets unchanged. The moves continue until no more tuples remain than there are pairs (J, a). This is the sequential construction idea of Koller and Megiddo [15, Section 3]. The proof of the next lemma, including the implementation and bit bound, is in Section D. We store the rational weights explicitly and use them to measure retained input sets.

Lemma 7.1 (Small rational tuple lists)**.** *Fix an integer independence order $\mathsf{D} \geq 1$. Given normalized rational weights μ on an alphabet of size m and an integer $q \geq 1$, a deterministic algorithm constructs an exactly D-wise independent tuple list $\mathsf{Code}_{\mathsf{D}}(q, \mu)$ (Definition 4.4) with a label set of size at most $(1 + qm)^{\mathsf{D}}$. For fixed D, the construction time, total output length and rational-weight bit lengths are polynomial in q, m and the binary encoding length of μ.*

Below we always use Code_{8192}. Keeping the weights as binary-encoded rational numbers avoids expanding them into repeated labels. A denominator with polynomially many bits could require far too many copies to print in polynomial time.

7.2 Restoring the deletion bounds

The matrix W of Theorem 6.16 satisfies the deletion bounds one bit below its cost, and has full cost on restrictions retaining almost all column weight. One more interlace recovers the missing bit. The proof distinguishes whether Bob speaks in a protocol prefix. If not, the bound for almost all columns applies. Otherwise Lemma C.3 gives the missing bit.

Theorem 7.2. *Let $0 < \phi \leq 1$, $0 < \delta \leq 1/64$ and $\theta > 0$. Suppose W, with input weights μ_X, μ_Y, satisfies $\mathrm{D}(W) = C \geq 3$, $\mathcal{P}(W; \theta, C - 1)$, and*

$$\mathrm{D}(W[A, B]) \geq C \quad \text{whenever } \mu_X(A) \geq \theta \text{ and } \mu_Y(B) \geq 1 - \delta.$$

Let Q be a power of two satisfying

$$Q \geq 128\big(e(\phi, \delta/2) + 1\big), \qquad Q\theta \leq 1/20,$$

where $e(\cdot, \cdot)$, defined before Lemma 6.10, bounds the number of exceptional coordinates in that lemma. Form the Q-fold interlace of W on any exactly 8192-wise independent tuple list with marginal μ_Y, and transpose it to obtain R^+. Then

$$\mathrm{D}(R^+) = C + \log Q, \qquad \mathcal{P}(R^+; \phi, C + \log Q).$$

Independently of the lower-bound hypotheses, a protocol of depth U for W yields one of depth $U + \log Q$ for R^+.

Proof. To prove the lower bounds, we fix $j \in \{0,1\}$ and a restriction of R^+ retaining row weight $\geq \phi$ and column weight $\geq (3/5)2^{-j}$. After undoing the transpose, the retained columns form a label event Y of tuple-list weight $\nu(Y) \geq \phi$, and the within-block weights of the retained rows average at least $(3/5)2^{-j}$ over the Q blocks. We apply Lemma 6.10 to Y with $v = \phi$ and error bound $\delta/2$ and put

$$e = e(\phi, \delta/2), \qquad k = Q2^{-j}, \qquad a = \log k.$$

Since Q is a power of two with $Q \geq 128(e+1) > 2$, a is a positive integer. More than $k/2$ blocks are indexed outside the exceptional set and retain base-row weight at least $k\theta$. Otherwise the block-counting bound of Section 4 with e blocks set aside gives $(3/5)Q2^{-j} \leq k/2 + e + Qk\theta$, and hence

$$\frac{3}{5} \leq \frac{1}{2} + \frac{2^j e}{Q} + Q\theta \leq \frac{1}{2} + \frac{1}{64} + \frac{1}{20} < \frac{3}{5}.$$

We retain these blocks and all of Y.

Suppose the restriction to these blocks and Y has a protocol of depth at most $C+a-1$. We follow the protocol for a messages, starting from conditional column weight one, by the weight-halving walk from the proof of Theorem 6.12, which there runs for $a-1$ messages. At each Alice message, we assign every current block to a child retaining at least half its row weight and keep a child assigned at least half the blocks. At each Bob message, we keep a child retaining at least half the current column weight under $\nu \mid Y$. If the protocol ends earlier, we pad with dummy Alice messages. If α messages are Alice's and β are Bob's, then $\alpha + \beta = a$, and the surviving restriction has a column set $K' \subseteq Y$ and

$$\text{more than } 2^{\beta-1} \text{ blocks outside the exceptional set}, \qquad \text{row weight per block at least } 2^\beta\theta,$$
$$(\nu \mid Y)(K') \geq 2^{-\beta}, \qquad \text{at most } C-1 \text{ bits remaining.}$$

If $\beta = 0$, the columns are still all of Y. At least one block indexed outside the exceptional set remains, and its row weight is at least θ. The projection of Y onto its coordinate has base weight at least $1 - \varepsilon(\phi, \delta/2) \geq 1 - \delta$, by Lemma 6.10. Choosing retained representatives (Observation 4.5), we obtain a restriction of W that requires C bits by the hypothesis for row weight $\geq \theta$ and column weight $\geq 1 - \delta$.

If $\beta \geq 1$, we apply the short-continuation bound Lemma C.3 with $d = C-1$, $b = \beta - 1$, and $x = \theta$, under hypothesis (C) with $v = \phi$ and $\tau = \delta/2$. Its hypotheses hold. The bound $\mathcal{P}(W; \theta, C-1)$ is assumed, the blocks avoid the exceptional set, and the number of blocks is an integer exceeding $2^{\beta-1}$, so it is at least $2^b + 1$. Each block has row weight at least $2^{b+1}x$, and $(\nu \mid Y)(K') \geq 2^{-b-1}$. Also $d \geq 2$ and $2^{b+1}x = 2^\beta\theta \leq k\theta \leq 1$. The lemma therefore requires $d + 1 = C$ bits. In both cases this exceeds the $C-1$ remaining bits.

Thus the restriction to these blocks and Y requires at least $C+a = C+\log Q-j$ bits. Transposing back proves both deletion bounds and, for $j = 0$, the lower bound on the whole matrix. For the upper bound, Alice sends the block index i in $\log Q$ bits, Bob decodes coordinate i of his tuple, and they run the depth-U protocol for W. Exchanging the players gives the same cost for R^+. Taking $U = C$ proves equality. □

8 Deterministic gap tables

The proof of Theorem 5.3 covers explicit lists once their marginal weights and independence properties are verified. We now construct these lists and bound their encoding length, making each round a deterministic polynomial-time operation. Theorem 8.4 gives the resulting gap and encoding bound after k rounds.

The subscript g marks parameters and matrices of the explicit-list construction. We choose a power of two s_{g} above the fixed cutoff s_0 of Lemma F.3, with $\mathsf{N} := s_{\mathrm{g}}^4 \geq |V(G)|$ for the source graph G. Set $L_{\mathrm{g}} = \log \mathsf{N}$ and $\rho_{\mathrm{g}} = 1/\mathsf{N}$. The choices in Definition 6.1 then fix every other parameter as a function of s_{g}. In particular, $q_{2,\mathrm{g}} = s_{\mathrm{g}}^2/\rho_{\mathrm{g}} = s_{\mathrm{g}}^6$ since $\rho_{\mathrm{g}} = s_{\mathrm{g}}^{-4}$, $w_{\mathrm{g}} = s_{\mathrm{g}}^8$, $Q_{\mathrm{g}} = s_{\mathrm{g}}^{24}$, and the block threshold is $x_{\mathrm{g}} = \rho_{\mathrm{g}}^8$. These values are the same in every round, although the base matrix R changes from round to round.

Corollary 8.1 (Explicit case)**.** *For the construction $R \mapsto R_{\mathrm{g}}^+$ of Definition 8.3 with the parameters of Lemma F.3, parts (i) and (ii) of Theorem 5.3 hold with $L = L_{\mathrm{g}}$, so that $2^{-L} = \rho_{\mathrm{g}}$, and with*

$$h = h_{\mathrm{g}} := 2\log s_{\mathrm{g}} + \log q_{2,\mathrm{g}} + \log w_{\mathrm{g}} + \log Q_{\mathrm{g}} = 40 \log s_{\mathrm{g}} = 10 \log \mathsf{N}.$$

The construction is deterministic and runs in time polynomial in N and the explicit rational encoding length of R. Every row and column label of R_{g}^+ decodes to one of the full-product matrix R^+, with the same entry at each decoded pair.

After defining the lists below, we prove this corollary in Section 8.2.

8.1 Replacing products by explicit lists

A coordinate of a tuple in one small list can itself be a label of another small list. The next lemma shows that decoding such nested labels preserves the product weights on every set of at most 8192 coordinates. It also shows that fixing independently weighted labels of other lists does not change these weights. As in Section 4, inputs with the same decoded value keep separate labels and weights.

Lemma 8.2. *Let an inner indexed list (Definition 4.4) on $Z^{q_{\mathrm{in}}}$ be exactly* 8192*-wise independent with marginal μ_Z. Let an outer list on q_{out}-tuples of inner labels be exactly* 8192*-wise independent with the inner label weights as marginal. Then the decoded coordinates, indexed by pairs in $[q_{\mathrm{out}}] \times [q_{\mathrm{in}}]$, are exactly* 8192*-wise independent with marginal μ_Z. If the outer-list label has product weights with the tuple of labels from finitely many other lists, fixing those other labels at any joint value of positive weight leaves the outer-label weights, and hence the conclusion above, unchanged.*

Proof. Given at most 8192 prescribed decoded coordinates, we group them by outer position. Outer independence factors their joint weight over these groups, and inner independence factors each group into its prescribed coordinate marginals. Each factorization involves at most 8192 coordinates, as both independence hypotheses require. By the product form, fixing the other labels leaves the outer-label weights, which determine the decoded weights, unchanged. □

Definition 8.3. Let $R : X \times Y \to \{0,1\}$ be a rationally weighted base matrix. With the parameter choices of Lemma F.3, construct the items and the indicators $v_i(\alpha)$ as in Definition 6.2. Replace the product alphabets Γ, S, and J_2^s of that definition by the indexed lists Γ_{g}, S_{g}, and $\mathcal{J}_{\mathrm{g}}$ of Lemma 7.1, defined by

$$\Gamma_{\mathrm{g}} = \mathsf{Code}_{8192}(2s_{\mathrm{g}} - 1, \mu_Y), \qquad J_{\mathrm{g}} = [q_{1,\mathrm{g}}] \times X, \qquad S_{\mathrm{g}} = \mathsf{Code}_{8192}(q_{2,\mathrm{g}}, \mu_{J_{\mathrm{g}}}),$$

$$J_{2,\mathrm{g}} = [q_{2,\mathrm{g}}] \times \Gamma_{\mathrm{g}}, \qquad \mathcal{J}_{\mathrm{g}} = \mathsf{Code}_{8192}(s_{\mathrm{g}}, \mu_{J_{2,\mathrm{g}}}).$$

Integer factors are uniform, and all other factors have their base or indexed weights. On decoded labels set

$$A_{\mathrm{g}}(\gamma, (r,u)) = R(u, \gamma_r), \qquad R_{\mathrm{temp,g}}(z, (\alpha, \gamma)) = A_{\mathrm{g}}(\gamma, z_\alpha).$$

Use the raw entry rules, family tags, and type weights of Definition 6.2. For the family corresponding to Ω, the column weights are the product of the weights of a $\mathcal{J}_{\mathrm{g}}$ label and the weights induced on

the tail by Γ_{g}. For the family corresponding to $\mathcal{D}$, the weights are induced from $J_{2,\mathrm{g}}$ under the map $j = (\alpha, \gamma) \mapsto (\mathrm{tail}(\gamma), j, \ldots, j)$ defining $\mathcal{D}$. These rules define $B_{\mathrm{raw,g}}$, its transpose N_{g}, and its normalized restriction $N_{\mathrm{ref,g}}$ to inputs with the Ω family tag.

The lower-bound proof fixes the type/bin/item sequence and the base-label sequence before restricting the template labels, which must then still have their original weights, as required by (H3) of Section 6.5. To interlace w_{g} copies of N_{g}, we therefore combine three tuple lists whose labels have product weights. Each list has tuples of length w_{g} and is exactly 8192-wise independent, with the following marginals. Type/bin/item labels have the weights induced by the raw-row weights, base labels have weights μ_X, and template labels have the S_{g} weights. We enumerate the Cartesian product of these three label sets. At each position, the type label in the first sequence determines whether the decoded row is a template row (p, z), from the bin and template labels, or a row (i, u) indexed by an item, from the item and base labels. Since the three sequences are independent and each is exactly 8192-wise independent, their coordinatewise combination gives an exactly 8192-wise independent tuple list with the raw-row marginal. Interlacing N_{g} on this combined list defines U_{g}. Set $W_{\mathrm{g}} = U_{\mathrm{g}}^{\mathsf{T}}$. We interlace W_{g} at width Q_{g} on Code_{8192} with the column weights of W_{g} and then transpose to obtain R_{g}^{+}. We store all input labels and rational weights explicitly, keeping inputs with equal decoded values separate.

As in Observation 4.5, we obtain projected restrictions from chosen representatives. The per-copy deletion bounds and the collision bound are the explicit case of Lemma 6.6. The proof of the deletion bounds uses Lemma 6.10 and Theorem 6.12 to handle retained events of weight as small as x_{g}, and the collision bound follows from Lemma 6.11. With these bounds and the exact independence of Lemma 8.2, the collision and stopping-node arguments of Theorem 6.9 and Section 6.5 apply to the indexed construction.

8.2 Iteration, encoding size and deterministic running time

To prove Corollary 8.1, we need only add the encoding bound to the explicit case of Theorem 5.3.

Proof of Corollary 8.1. We suppress the subscript g. Each width q_1, q_2, w, Q is a fixed polynomial in N (Lemma F.3). Therefore the fixed number of calls to Lemma 7.1, all of order 8192, and enumeration of the Cartesian products of their label sets take time bounded by a fixed polynomial in N and the encoding length of R. The indexed labels and rational weights needed by the next round have encoding length within the same bound.

Parts (i) and (ii) follow from the proof of Theorem 5.3, with Theorem 7.2 used in place of Theorem 6.17. We apply Theorem 7.2 with $(\phi, \delta, Q) = (\rho, \delta_U, Q)$, and Lemma F.3 verifies its numerical hypotheses. The upper-bound protocol uses only the decoded inputs and not their weights, so it also works for the indexed construction. □

Theorem 8.4. *There are effective absolute constants $\mathsf{A} \geq 2$ and $C_0 > 0$ and a deterministic algorithm with the following guarantee. Given a YES or NO source graph G from Theorem 5.2, an integer $k \geq 1$, and a power of two s_{g} such that $\mathsf{N} = s_{\mathrm{g}}^4 \geq |V(G)|$ meets the parameter conditions of Lemma F.3, it produces an explicit total Boolean matrix $F_{G,k}$ and a public integer B_k^{g} such that*

$$\textit{YES: } \mathrm{D}(F_{G,k}) \leq B_k^{\mathrm{g}}, \qquad \textit{NO: } \mathrm{D}(F_{G,k}) = B_k^{\mathrm{g}} + k.$$

Its running time and full binary output length are at most $\mathsf{N}^{C_0 \mathsf{A}^k}$. The constants depend only on the fixed source reduction and construction. The output is square with power-of-two side length at least N. The public threshold is

$$B_k^{\mathrm{g}} = d_{0,\mathrm{g}} + k h_{\mathrm{g}},$$

where $d_{0,\mathrm{g}}$ is the cost of the initial matrix in Lemma 5.4 with $L = L_{\mathrm{g}}$:

$$d_{0,\mathrm{g}} = \log \mathrm{pow}_2(64(L_{\mathrm{g}} + 1)) + 1, \qquad L_{\mathrm{g}} = \log \mathsf{N}.$$

Proof. We use the parameters of Lemma F.3 for the given s_{g} and let R_0 be the initial matrix of Lemma 5.4 with $L = L_{\mathrm{g}} = \log \mathsf{N}$. We apply Definition 8.3 k times, using each output as the next base matrix and storing each indexed table with its rational input weights. Finally, we forget the weights and duplicate inputs to make a square table with power-of-two side at least N.

By Lemma 5.4, the initial matrix has exact cost $d_{0,\mathrm{g}} \geq 3$ and satisfies the deletion bounds at threshold $\rho_{\mathrm{g}} = 2^{-L_{\mathrm{g}}}$. Proposition 5.5, in the explicit case (Corollary 8.1), therefore gives

$$\mathrm{D}(R_k) \leq d_{0,\mathrm{g}} + kh_{\mathrm{g}} \quad \text{on YES}, \qquad \mathrm{D}(R_k) = d_{0,\mathrm{g}} + k(h_{\mathrm{g}} + 1) \quad \text{on NO}$$

for the matrix R_k after k rounds. Forgetting weights and duplicating rows and columns to make the table square change neither cost (Section 4). This proves the gap with $B_k^{\mathrm{g}} = d_{0,\mathrm{g}} + kh_{\mathrm{g}}$.

For size and time, let Z_i be $\mathsf{N}+2$ plus the complete encoding length after round i, including labels and normalized rational weights. The polynomial bound for the initial matrix and Corollary 8.1 give effective constants C_1 and a power of two $\mathsf{A} \geq 2$, independent of i and k, such that, for all $i \geq 0$,

$$Z_0 \leq \mathsf{N}^{C_1}, \qquad Z_{i+1} \leq Z_i^{\mathsf{A}}, \qquad \text{time for round } i+1 \leq Z_i^{\mathsf{A}},$$

and hence $Z_i \leq \mathsf{N}^{C_1 \mathsf{A}^i}$. The padded side is below $2Z_k$ (Section 4). After enlarging the effective constant C_0 if necessary, we bound the sum of the k round times and the final square padding by $\mathsf{N}^{C_0 \mathsf{A}^k}$. Computing the public threshold and the parameters takes time within the same bound. □

We now derive Theorem 1.1 by taking a fixed number k of rounds, and Corollary 1.5 by taking $k = \Theta(\log \log N)$ rounds for a sparse formula on N variables.

Proof of Theorem 1.1. For fixed k, we apply the graph reduction of Theorem 5.2 and then Theorem 8.4. The least value $\mathsf{N} = s_{\mathrm{g}}^4$ meeting the parameter conditions obeys $\mathsf{N} \leq 16 \max\{|V(G)|, \mathsf{N}_0\}$ for the effective constant $\mathsf{N}_0 = s_0^4$, by Lemma F.3. Since the graph reduction runs in polynomial time, $|V(G)|$ is polynomial in the formula length, and since k is fixed, so is $\mathsf{N}^{C_0 \mathsf{A}^k}$. Thus $B = B_k^{\mathrm{g}}$ separates cost $\leq B$ from cost exactly $B + k$, on a square table with power-of-two side. □

Proof of Corollary 1.5. We choose the number k of rounds so that the gap k exceeds twice the allowed error $c_* \log \log T$. At the same time, k must be small enough that the construction and the estimator run in quasipolynomial time in the size of the sparse formula. Let A be the effective power of two from Theorem 8.4, and set $c_* = 1/(8 \log \mathsf{A})$. Suppose such an estimator exists. For a sparse formula on N variables and the graph G of Theorem 5.2, the least power of two $s_{\mathrm{g}} \geq \max\{N^{1/4}, |V(G)|^{1/4}, s_0\}$ meets the conditions of Lemma F.3, and $\max\{N, |V(G)|\} \leq \mathsf{N} = s_{\mathrm{g}}^4 < 16 \max\{N, |V(G)|, \mathsf{N}_0\}$. Since $|V(G)| \leq N(\log(N+2))^p$ for the fixed p of Theorem 5.2, we have $\mathsf{N} = N^{O(1)}$. For sufficiently large N, we choose

$$k = \left\lfloor \frac{\log \log N}{\log \mathsf{A}} \right\rfloor.$$

Then $k = \Theta(\log \log N)$ and $\mathsf{A}^k \leq \log N$. The square output of Theorem 8.4 has threshold B_k^{g} and T entries, where

$$N^2 \leq T \leq \mathsf{N}^{C_0 \mathsf{A}^k}, \qquad \log \log T \leq k \log \mathsf{A} + O(\log \log N).$$

For sufficiently large N, the second bound is at most $3k \log \mathsf{A}$: indeed it is at most $2 \log \log N + O(1)$, whereas $k \log \mathsf{A} = \log \log N + O(1)$. Consequently

$$2c_* \log \log T \leq 3k/4 < k.$$

The estimator's guarantee holds on all sufficiently large tables, and $T \geq N^2$. Hence, for sufficiently large N, comparing the estimate with the public midpoint $B_k^{\mathrm{g}} + k/2$ decides satisfiability. The construction time is at most $\mathsf{N}^{C_0 \mathsf{A}^k}$, and the estimator's running time is a fixed polynomial in $T \leq \mathsf{N}^{C_0 \mathsf{A}^k}$. Since $\mathsf{N} = N^{O(1)}$ and $\mathsf{A}^k \leq \log N$, both are bounded by

$$2^{O((\log N)^2)}.$$

For the finitely many remaining values of N, including those with $k < 1$, exhaustive search over assignments decides satisfiability. This proves the claimed deterministic running-time implication. □

9 Final sampling and growing error

In the full-product case of Proposition 5.5, the iterated matrix R_k has cost at most B_k on YES graphs and exactly $H_k = B_k + k$ on NO graphs, but it is too large to print, so we replace it by a sampled submatrix. A single sample must be hard for every protocol of depth below the NO-case bound. A small sample suffices because the rank bound $\operatorname{rank}(R_k) \leq r_k$ of Proposition 5.5 limits the number of column sets on which such protocols can be correct (Theorem 9.3).

9.1 Counting column sets on which short protocols are correct

For the union bound in Section 9.2, we count the sets of columns on which a short protocol can be correct on a retained row set. One bound covers all row sets at once. We fix the tree, Alice's message functions and the leaf outputs, and let Bob's bits depend on the column.

Definition 9.1 (Columns on which Bob can complete the protocol)**.** Let a depth-t tree P specify a complete binary tree of depth t, the speaker at each internal node, all Alice maps (from her row to her bit at each of her nodes), and all leaf outputs. Bob's maps are unspecified. For a Boolean matrix $F : X \times Y \to \{0, 1\}$ and $U \subseteq X$, let $G_P(U)$ be the set of columns for which some assignment of bits to Bob's nodes makes the tree correct on every row in U. The assignment may depend on the column, but not on Alice's row.

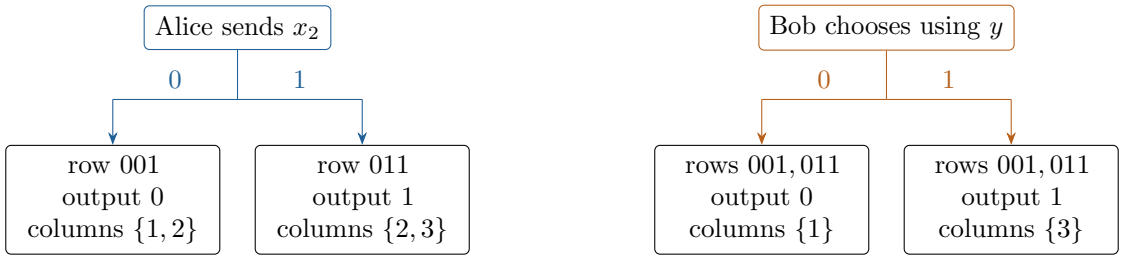
Observation 9.2 (Lemma 3.2 in [17])**.** *For retained sets $U \subseteq X$, $V \subseteq Y$ and an integer $t \geq 0$, with correctness on an empty restriction understood vacuously,*

$$\mathrm{D}(F[U, V]) \leq t \quad \Longleftrightarrow \quad V \subseteq G_P(U) \text{ for some depth-}t \text{ protocol tree } P \text{ of this kind.} \tag{2}$$

Figure 4 uses two one-bit trees to show why leaf requirements combine by intersection at an Alice node and by union at a Bob node. We label each leaf by the columns on which its output is correct for the rows reaching it.

The same matrix in both trees: $F(x, y) = x_y$, the yth bit of x.

	Bob's column y		
Alice's row x	1	2	3
001	0	0	1
011	0	1	1



(a) **Both required: intersection.** $\{1, 2\} \cap \{2, 3\} = \{2\}$.

(b) **Either works: union.** $\{1\} \cup \{3\} = \{1, 3\}$.

Figure 4: Alice's branch depends on the row, so a column on which the tree is correct must satisfy both leaf conditions. Bob's branch depends on the column, so either leaf may work, but his chosen leaf must answer both rows correctly.

Mackenzie proves the count in Theorem 9.3 with the row-sum argument of Oum, Sæther and Vatshelle [20, Theorem 4.2]: the columns on which selected Boolean rows all vanish are exactly the zero coordinates of the integer sum of those rows. Counting these sums instead of the row subsets gives a bound that depends only on the rank and not on the number of rows, which keeps the final sample small.

Theorem 9.3 (Lemma 4.3 in [17]). *Let $r \geq 1$ and $t \geq 0$ be integers. If a Boolean matrix F has real rank at most r, then, over all depth-at-most-t trees P specified as in Definition 9.1 and all retained row subsets U, there are at most*

$$2^{(r+1)\lceil \log(r+2) \rceil\, 2^{t+1}} \tag{3}$$

distinct sets $G_P(U)$. Repeated row and column labels are allowed.

Theorem 9.3 follows from Mackenzie's bound $2^{B(r)2^t - 1}$ for depth exactly t, where $B(r) = (r+1)\lceil \log(r+1) \rceil + 2$. Since $B(r) \leq 2(r+1)\lceil \log(r+2) \rceil$ for $r \geq 1$, and padding turns a tree of depth $t' < t$ into a depth-t tree with the same sets $G_P(U)$, this gives (3). Repeated labels are handled by passing to distinct values [17, Section 2].

9.2 Sampling columns and then rows

We sample the same number m of rows and columns, and we take m to be a power of two, so the output has the square power-of-two form of Corollary 1.4. In the proof, we first show that, with high probability, the restriction to the sampled columns still has cost at least H on every row set of original weight $\geq 3/5$. We then analyze the row sample, which is independent of the columns.

Theorem 9.4 (Sampling rows and columns). *Let $F : X \times Y \to \{0, 1\}$ have real rank at most r, where $r \geq 1$ is an integer, and let μ_X, μ_Y be input distributions on nonempty finite indexed sets. Suppose an integer $H \geq 1$ satisfies, for all $U \subseteq X$ and $V \subseteq Y$,*

$$\mu_X(U) \geq 3/5, \quad \mu_Y(V) \geq 3/5 \quad \Longrightarrow \quad \mathrm{D}(F[U, V]) \geq H.$$

Put

$$m = \mathrm{pow}_2\big(16(r+1)\,\lceil \log(r+2)\rceil\; 2^H\big),$$

where $\mathrm{pow}_2(u)$ *is the least power of two at least* u*. Draw* m *rows from* μ_X *and* m *columns from* μ_Y*, all independently, retaining separate indices for repeated draws. The resulting* $m \times m$ *matrix has cost at least* H *except with probability less than* $1/100$*. Every protocol for* F *also computes the sampled matrix, so the sampled matrix has cost at most* $\mathrm{D}(F)$ *on every sample outcome.*

Proof. Separate indices for repeated draws do not change costs. Each player can decode an index to its label, and conversely a protocol on the indexed matrix restricts to one occurrence of each label.

For every U with $\mu_X(U) \geq 3/5$ and every depth-$(H-1)$ tree P specified as in Definition 9.1, we have $\mu_Y(G_P(U)) < 3/5$: otherwise Observation 9.2 with $V = G_P(U)$ gives $\mathrm{D}(F[U,V]) \leq H-1$, contradicting the restriction hypothesis. As U and P vary, these sets form a family of size at most $2^{m/16}$ by Theorem 9.3, since $m \geq 16(r+1)\,\lceil \log(r+2)\rceil\; 2^H$. By a union bound, the probability that all sampled columns lie in one of them is at most $2^{m/16}(3/5)^m$. Outside this event, the sampled-column matrix $F_Y(x,j) = F(x,y_j)$ has cost $\geq H$ on every original row set of weight $\geq 3/5$.

We condition on such a column outcome. The transpose F_Y^{T} still has rank at most r. Consider the sets $G_P(U)$ for this transpose, where $U = [m]$ is its full row set and P has depth $H-1$. Each has $\mu_X(G_P(U)) < 3/5$: otherwise Observation 9.2 and transposition give $\mathrm{D}(F_Y[G_P(U),[m]]) \leq H-1$, contradicting the preceding guarantee. There are again at most $2^{m/16}$ such sets. The row draws remain independent with distribution μ_X, so the probability that they all lie in one of these sets is at most $2^{m/16}(3/5)^m$. By Observation 9.2 for the transpose of the final matrix, a protocol of depth at most $H-1$ on the final matrix would put all sampled rows in one of these sets. Since $r, H \geq 1$ give $m \geq 32$, the total failure probability is at most

$$2 \cdot 2^{m/16}(3/5)^m \leq 2^{1-7m/16} \leq 2^{-13} < 1/100.$$

The universal upper bound follows by decoding sampled indices and running a protocol for F. □

We sample only once, after all amplification rounds on full products. Theorem 9.4 bounds the failure probability but gives no test certifying that a particular sample succeeds.

9.3 Printing the sample and comparing errors

We generate the sampled row and column labels and evaluate R_k only on the pairs they form.

We take a source graph on n vertices above the fixed effective cutoff in Lemma F.2, use the parameter choices of that lemma at every round, and write

$$L = \lceil \log n \rceil, \qquad B_k = d_0 + kh, \qquad H_k = B_k + k, \qquad r_k = 2^{B_k - 1}$$

as in Proposition 5.5, which gives $\mathrm{rank}(R_k) \leq r_k$. The reduction uses only this bound and never computes a rank. A *random tape* is the sequence of random bits used by the algorithm. The next lemma bounds the running time on every tape, not only on likely ones.

Lemma 9.5. *Let* R_k *be the matrix of Proposition 5.5 in the full-product case, and let* m *be the sample size of Theorem 9.4 with* $r = r_k$ *and* $H = H_k$*. An algorithm generates* m *row and* m *column labels of* R_k *and evaluates* R_k *on all their pairs, using* $n^{O(k+1)}$ *time and random bits on every tape, with an absolute implied constant independent of* k*. Every output label is an actual input of* R_k*, and the sample counts and tape length do not depend on the tape. The joint distribution of the generated labels is within total variation less than* $1/100$ *of the ideal independent sample of Theorem 9.4.*

Here total variation is the maximum, over events, of the absolute difference between the probabilities that two distributions assign to the event. The proof in Section E gives the recursive construction with a fixed number of random bits. The label distribution is only close to the ideal one, but every printed entry is an exact entry of R_k. Combining the lemma with Theorem 9.4, we obtain the following reduction.

Theorem 9.6 (Randomized reduction)**.** *There is a randomized algorithm that, given a source graph on n vertices above the fixed cutoff in Lemma F.2 and an integer $k \geq 1$, outputs a total Boolean matrix M and an integer B. The table is square with power-of-two side length. On every tape the algorithm uses $n^{O(k+1)}$ time and random bits, with an absolute implied constant independent of k. For a YES source, $\mathrm{D}(M) \leq B$ on every tape. For a NO source,*

$$\Pr[\mathrm{D}(M) = B + k] \geq 49/50.$$

Here $B = B_k$ and the tape-independent entry count is

$$m = \mathrm{pow}_2\big(16(r_k + 1)\,\lceil \log(r_k + 2) \rceil\; 2^{H_k}\big), \qquad T = m^2.$$

Proof. The algorithm computes the parameters of Lemma F.2 and the initial matrix of Lemma 5.4, then B_k, H_k, r_k, m. It generates m columns and m rows with Lemma 9.5, evaluates all m^2 row–column pairs (keeping repeated labels as distinct indices), and outputs (M, B_k). The time and random-bit bounds follow from Lemma 9.5, and the output is $m \times m$ with m a power of two. Every generated label is an input of R_k. On a YES source, decoding the labels and running a protocol for R_k of cost at most B_k (Proposition 5.5) shows that every sample has cost at most B_k.

On a NO source, Proposition 5.5 gives

$$\mathrm{D}(R_k) = H_k, \qquad \mathcal{P}(R_k; 2^{-L}, H_k), \qquad \mathrm{rank}(R_k) \leq r_k.$$

Because $2^{-L} \leq 1/2 < 3/5$, the $j = 0$ lower bound applies whenever both retained weights are at least $3/5$. Thus, by Theorem 9.4, an ideal sample has cost below H_k with probability less than $1/100$. By Lemma 9.5, using the implemented sampler instead of the ideal independent sample adds less than $1/100$ to this failure probability. Therefore, except with probability less than $1/50$, the output has cost $\geq H_k$. Decoding the inputs and using $\mathrm{D}(R_k) \leq H_k$ (Proposition 5.5) gives cost $\leq H_k$ on every tape, so the cost is then exactly $B_k + k$. □

9.4 From an approximation algorithm to a 3-SAT algorithm

We now combine the two constructions with an estimator, proving the general implication stated in Theorem 9.7. Each construction produces a known threshold B and a gap k. The size bound of the construction determines which approximation errors are small enough to distinguish the two cases.

Theorem 9.7 (From approximation algorithms to 3-SAT algorithms)**.** *There is an effective absolute integer $\mathsf{A} \geq 2$ and, for each fixed clause-density bound c_s, an effective integer $C \geq 1$ with the following property. Suppose a deterministic or bounded-error randomized polynomial-time estimator approximates $\mathrm{D}(M)$ with additive error at most a nondecreasing function $a(T)$ on every sufficiently large explicit total Boolean table with T entries. Let the gap $k = k(N) \geq 1$ be an integer function computable in time polynomial in $N + k(N)$. Choose either construction and its size bound:*

Construction	$S(N,k)$
Randomized	$(N+2)^{C(k+1)}$
Deterministic	$(N+2)^{C\mathsf{A}^k}$

If, for all sufficiently large N,

$$2a\big(S(N,k(N))\big) < k(N),$$

then 3-SAT on N variables and at most c_sN clauses has an algorithm running in time $S(N,k(N))^{O(1)}$. It is deterministic when both the construction and the estimator are deterministic, and bounded-error randomized otherwise. The exponent may depend on the estimator and the algorithm computing k. It suffices for the estimator to work on square tables whose side length is a power of two.

Proof of Theorem 9.7. Apply Theorem 5.2, then Theorem 9.6 in the randomized case or Theorem 8.4 in the deterministic case. Since the source graph and the size parameter N of Theorem 8.4 are $N^{O(1)}$, a fixed C bounds both the table size T and construction time by the chosen $S(N,k)$. The constant A is the absolute constant from Theorem 8.4. The table sizes T tend to infinity with N, so the estimator's guarantee applies for all sufficiently large N.

Because a is nondecreasing, $a(T) \le a(S(N,k)) < k/2$. An accurate estimate is therefore below the midpoint $B + k/2$ on a YES output and above it on a successful NO output. If the estimator is randomized, repeat this comparison independently a constant number of times and take the majority answer, reducing its error below $1/100$. The randomized construction adds failure probability at most $1/50$, and the deterministic construction adds none. Thus the combined algorithm has error less than $1/3$ whenever randomness is used, and is exact otherwise.

Construction, computing k, estimation, and comparison with the binary rational midpoint all take time $S(N,k)^{O(1)}$. In particular, the assumed polynomial cost of computing k is $S(N,k)^{O(1)}$ because $S(N,k) \ge N + k$. Exhaustive search handles the finitely many smaller sizes. □

To read this in terms of a proposed 3-SAT running time $R(N) = 2^{g(N)}$, suppose $g(N)/\log(N+2)$ grows without bound and the chosen number of rounds can be computed within the stated time. A running time $2^{O(g(N))}$ allows

$$k = \Theta\left(\frac{g(N)}{\log(N+2)}\right) \qquad \text{in the randomized construction,}$$
$$k = \Theta\left(\log\frac{g(N)}{\log(N+2)}\right) \quad \text{in the deterministic construction.}$$

The constant in the second choice is chosen so that $\mathsf{A}^k \log(N+2) = O(g(N))$. Additive error below $k/2$ on tables of size $S(N,k)$ therefore yields a 3-SAT algorithm running in time $2^{O(g(N))}$. If such a 3-SAT algorithm is excluded by an explicit hypothesis, the same statement gives inapproximability. To quantify a further improvement over a known algorithm, choose a smaller target running time R. The theorem does not assume that the best known running time is optimal.

For Corollary 1.4, we use the exact sample size rather than the coarser bound $S(N,k)$. This gives the coefficient $1/8$ in the randomized error bound, and the resulting algorithm is the same.

Proof of Corollary 1.4. Suppose a bounded-error randomized polynomial-time estimator achieves additive error $a(T) = c\log T/\log\log T$ for some fixed $0 \le c < 1/8$ on every sufficiently large square table of power-of-two side length. Fix $8c < \alpha < 1$. We obtain a bounded-error $2^{O(N^\alpha)}$ algorithm for sparse 3-SAT.

For a formula with N variables, Theorem 5.2 produces a source graph with n vertices, where

$$2N/\Delta \le n \le N(\log(N+2))^p$$

for the fixed integer p of that theorem. The lower bound follows from $N \le |E| \le \Delta n/2$, also given by that theorem. Thus the graph is above the fixed cutoff for every sufficiently large N, and

$$L = \lceil \log n \rceil = \log N + O(\log\log N).$$

Choose a fixed rational β with $8c < \beta < \alpha$, and put $k = \lfloor N^\beta / \lceil \log N \rceil \rfloor$. This integer is computable in polynomial time, using an integer root for the fixed rational power. By Lemmas 5.4 and F.2, $h = L + O(\log L)$ and $d_0 = O(\log L)$, so

$$B_k = d_0 + kh = (1 + o(1))k \log N.$$

Substituting $r_k = 2^{B_k - 1}$ and $H_k = B_k + k$ into Theorem 9.6 gives $\log m = 2B_k + k + O(\log(B_k + 1))$. Hence the printed size satisfies

$$\log T = 2 \log m = 4B_k + 2k + O(\log(B_k + 1)) = (4 + o(1))k \log N, \tag{4}$$

$$\log \log T = (\beta + o(1)) \log N, \qquad k = \left(\frac{\beta}{4} + o(1) \right) \frac{\log T}{\log \log T}. \tag{5}$$

Consequently $k > 2a(T)$ for all sufficiently large N. The table already has the square power-of-two form required by the assumed estimator.

For each estimator run, we test whether its output is below the rational midpoint $B_k + k/2$. An accurate estimate answers this test correctly on every YES output and every successful NO output. We repeat the run independently a constant number of times and take the majority answer, which reduces the decision error on each such table below $1/100$. This uses only comparison with a binary rational, which the estimator model supports (Definition 4.1). The reduction fails with probability at most $1/50$, so the decision error is less than $3/100 < 1/3$.

Both construction time and random-bit usage are $n^{O(k+1)} = 2^{O(N^\beta)}$ on every tape. The estimator also takes $2^{O(N^\beta)}$ time because $\log T = O(N^\beta)$. For all sufficiently large N, we have $k \geq 1$, the graph is above the fixed cutoff, and T meets the estimator's size threshold. The finitely many smaller values of N are decided by exhaustive search within a fixed bound. Since $\beta < \alpha$, the total running time is $2^{O(N^\alpha)}$, as claimed. □

10 Strictly alternating protocols

We now prove the result that settles Yao's original question. If we simulate the multi-bit messages of the unrestricted construction by inserting the missing turns, the YES-case upper bound can double while the known NO-case lower bound stays the same, so the additive gap is lost. We therefore use a separate matrix designed for strict alternation. The end of Section 2 sketches this construction. In the YES case, a balanced coloring of the source hypergraph gives a short alternating protocol (Proposition 10.4). In the NO case, a short protocol colors each vertex by Alice's most common response rules for her early turns. A monochromatic edge then gives many array positions on which her remaining messages produce too few answer patterns (Lemma 10.5). With an explicit array list (Corollary 10.6), the table has polynomial size for each fixed gap, and no repeated amplification is needed.

10.1 Starter, output and matrix inputs

Changing the starting player or the output convention changes the cost by at most two bits, so the gap transfers between the conventions of Definition 4.2.

Lemma 10.1. *On nonempty input sets, each cost $C \in \{D_A^{\mathrm{alt}}, D_B^{\mathrm{alt}}, D_{\mathrm{free}}^{\mathrm{alt}}, Y_{\mathrm{free}}\}$ satisfies*

$$|C(M) - D_A^{\mathrm{alt}}(M)| \leq 2.$$

Proof. Prepending one fixed bit sent by the new starter changes the starter, so $|D_A^{\mathrm{alt}} - D_B^{\mathrm{alt}}| \le 1$ and

$$D_A^{\mathrm{alt}} - 1 \le D_{\mathrm{free}}^{\mathrm{alt}} \le D_A^{\mathrm{alt}}.$$

Public output also gives receiver output. Conversely, at a nonzero-depth terminal transcript the receiver is the next scheduled speaker and can announce the answer in one bit. Public stopping tells both players when to make this announcement. At depth zero, choose the designated output player as the starter and let that player announce the answer. Thus

$$Y_{\mathrm{free}} \le D_{\mathrm{free}}^{\mathrm{alt}} \le Y_{\mathrm{free}} + 1,$$

and the claimed two-bit bound follows. Each of these conversions preserves public stopping and uses an input-independent starter. □

A hypergraph coloring is proper if no edge is monochromatic. In the YES case the source hypergraph has a balanced coloring with few colors. In the NO case it has no proper coloring even with the larger number of colors that a short protocol would induce. We use the balanced hypergraph coloring reduction of Guruswami and Lee [8], modified so that every edge has distinct vertices.

Proposition 10.2 (Hypergraph coloring promise)**.** *Fix an integer $t \ge 4$, and set*

$$s_{\mathrm{a}} = 2^t, \qquad K_{\mathrm{a}} = s_{\mathrm{a}}^2, \qquad S_{\mathrm{a}} = 2^{s_{\mathrm{a}}}.$$

There is a deterministic polynomial-time reduction from 3-SAT to K_{a}-uniform hypergraphs with nonempty vertex sets and distinct vertices in every edge such that:

1. *a satisfiable formula has an s_{a}-coloring with at least $s_{\mathrm{a}} - 1$ vertices of every color in every edge;*
2. *an unsatisfiable formula has no proper S_{a}-coloring.*

The polynomial may depend on t.

The proof in Section G.2 derives this promise from the cited result and ensures that every edge has distinct vertices.

Alice holds a vertex and a tuple of addresses, each indexing the first coordinate of Bob's array. Bob holds the array, a coordinate selecting one of Alice's addresses, and an edge index. If the vertex is outside the selected edge, or the edge index is unused, the answer is zero.

Definition 10.3. Let $\mathcal{H} = (V, E)$ be a K_{a}-uniform source hypergraph as in Proposition 10.2, with a fixed ordering of the vertices in each edge. Set

$$q_{\mathrm{a}} = 2^{t-1} = s_{\mathrm{a}}/2, \qquad H_{\mathrm{a}} = 2^b = \mathrm{pow}_2(\max\{1, |E|\}).$$

Index the edges by distinct $\alpha \in [H_{\mathrm{a}}]$, and let $\iota_\alpha(v)$ be the position of v in that edge, or zero if v is absent or the index is unused. Alice's address tuples are all of

$$\mathcal{C} = [H_{\mathrm{a}}]^{q_{\mathrm{a}}},$$

each used once. For any labeled Boolean array list $\mathcal{W} = (w^\ell)$ on

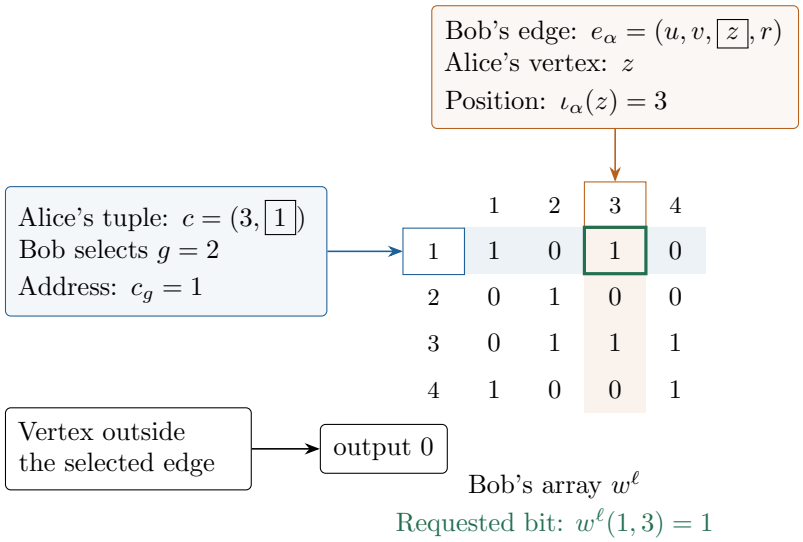
$$\Omega_{\mathrm{a}} = [H_{\mathrm{a}}] \times [K_{\mathrm{a}}],$$

define

$$M_{\mathrm{a}}((v,c),(g,\alpha,\ell)) := \begin{cases} w^{\ell}(c_g, \iota_\alpha(v)), & \iota_\alpha(v) > 0, \\ 0, & \iota_\alpha(v) = 0. \end{cases}$$

Rows are indexed by $(v,c) \in V \times \mathcal{C}$. Column labels lie in $[q_{\mathrm{a}}] \times [H_{\mathrm{a}}] \times [|\mathcal{W}|]$. Equal arrays retain their separate labels.

In Figure 5, g selects an address from Alice's tuple, which gives the row of Bob's array, and the position of her vertex in Bob's edge gives the column.



	1	2	3	4
1	1	0	1	0
2	0	1	0	0
3	0	1	1	1
4	1	0	0	1

Figure 5: The address selects a row and the vertex's edge position selects a column, giving $w^{\ell}(1,3) = 1$. A vertex outside the edge gives zero without reading the array. The indices are illustrative, and the arrows show how the entry is selected, not a protocol's messages.

10.2 The protocol given by a coloring

Once Bob has sent the edge, a balanced coloring lets Alice name her position among the at most $2s_{\mathrm{a}} - 1$ positions of her color in that edge. The schedule in the proof of Proposition 10.4 exchanges color and coordinate, then edge and address, then this position, and ends with Bob's answer bit. It costs

$$B_{\mathrm{a}} := (2t-1) + 2b + 2(t+1) + 1 = 2b + 4t + 2.$$

Proposition 10.4. *For the table of Definition 10.3, an s_{a}-coloring with at least $s_{\mathrm{a}} - 1$ vertices of every color in every edge gives*

$$D_A^{\mathrm{alt}}(M_{\mathrm{a}}) \le B_{\mathrm{a}}.$$

Proof. Fix a row (v,c), a column (g,α,ℓ) and a balanced coloring χ. Every edge contains at least $s_{\mathrm{a}} - 1$ vertices of each of the $s_{\mathrm{a}} - 1$ colors other than $\chi(v)$. Consequently the vertices whose color is $\chi(v)$ occupy at most

$$K_{\mathrm{a}} - (s_{\mathrm{a}} - 1)^2 = 2s_{\mathrm{a}} - 1$$

positions. Alice's message value is the rank of her position among these, or one extra value for "outside the edge." There are therefore at most $2s_{\mathrm{a}} = 2^{t+1}$ message values, encoded in $t+1$ bits. The ranking and the encoding are fixed publicly for every color and edge index. An unused edge index has only the value for "outside the edge." Alice determines her message value from v, α and $\chi(v)$.

First Alice sends the t bits of $\chi(v)$, while Bob sends the $t-1$ bits of g between consecutive Alice bits. This uses $2t-1$ bits and leaves Alice knowing g. Next Bob sends the b bits of α, while Alice sends the b bits of c_g, alternating. This uses $2b$ bits, after which Alice knows α and Bob knows c_g. Since Bob speaks next, he sends $t+1$ dummy bits while Alice sends the $t+1$ bits encoding her message value, after which Bob sends one answer bit. He answers zero for "outside the edge." Otherwise he decodes, from $\chi(v)$, α and the rank, the position $\iota_\alpha(v)$ and answers $w^\ell(c_g, \iota_\alpha(v))$. The three stages use B_{a} bits in total. □

10.3 Few response patterns cannot cover many array positions

Bob's early messages may depend on the whole array, so the transcript on one column does not determine how Alice behaves on other columns. To analyze an arbitrary protocol, we therefore also consider Bob histories that do not occur on a given column. For a fixed row, we record Alice's rth bit after every possible history of $r-1$ Bob bits, including hypothetical histories. Her own preceding bits are then determined. For this count, we pad branches that stop early with fixed bits in the alternating order. For example, her first two turns require three response bits: the initial bit and one response to each possible first Bob bit. In general, her first t turns require $\sum_{r=1}^{t} 2^{r-1} = 2^t - 1$ response bits, giving at most 2^{2^t-1} possible response rules for these turns.

Among all Boolean arrays, the fraction agreeing with prescribed bits on k distinct positions is 2^{-k}. We need only an upper bound of $2^{-k}+\varepsilon$ on this fraction for a small $\varepsilon > 0$, and a short explicit list already satisfies this weaker bound (Corollary 10.6). Here we count all array labels equally, keeping repeated arrays as separate labels. For a list $\mathcal{W} = (w^\ell)_{\ell\in\mathcal{L}}$ on a finite position set Ω, with $\mathcal{L}$ finite and nonempty, we require

$$\frac{|\{\ell \in \mathcal{L} : w^\ell|_I = z\}|}{|\mathcal{L}|} \le 2^{-|I|} + \varepsilon \qquad (I \subseteq \Omega,\ z \in \{0,1\}^I). \tag{6}$$

Adding these bounds over P patterns on the same set I, at most a $P(2^{-|I|}+\varepsilon)$ fraction of all labels matches one of them. Every fraction below is measured against the original list size $|\mathcal{L}|$, even after a transcript has removed some columns.

Lemma 10.5 rules out a short protocol when many requested positions share Alice's early response rules. Her early messages cannot distinguish those positions, so many array labels remain after them. The remaining messages produce only a few answer patterns, and by Equation (6) these match fewer labels than survive. In the lemma, m is the total number of Alice turns, and t is the number of initial Alice turns during which her response rules agree across the rows for every possible Bob history.

Lemma 10.5 (Lower bound for alternating protocols)**.** *Let m, t be integers with $m \ge 2t \ge 2$. Suppose the nonempty labeled array list $\mathcal{W}$ satisfies Equation* (6) *with $\varepsilon \le 2^{-4m}$. If $J \subseteq \Omega$ has at least 2^{m-t+1} positions, then the matrix*

$$L_J(j,\ell) := w^\ell(j)$$

has no Alice-first alternating protocol of depth at most $2m-1$ in which all rows $j \in J$ give Alice the same response rules for her first t turns on every possible Bob history.

Proof. Suppose such a protocol exists. We pad every shorter branch to depth $2m-1$ with fixed bits in the alternating order. At the final Alice node, her row and the transcript determine the old output leaf, so we replace her last bit by that output. This uses no extra bit and, since $m > t$, preserves her first t responses.

We stop with about $\log m$ turns per player remaining, so that Bob's remaining response rules require fewer bits to describe than there are surviving array positions. Put $h = \lceil \log m \rceil \le m - t$ and $Q = 2^h$. The inequality follows from $m \le 2(m-t) \le 2^{m-t}$. We follow the first $2(m-h)-1$ messages from the root. After each Bob message we keep the child with more column labels, and after each Alice message the child with more rows. Alice's first t turns discard no rows because these response rules agree. Each of her remaining $m-h-t$ turns retains at least half the rows, and each of the $m-h-1$ Bob turns retains at least half the columns. The resulting rectangle $X \times Y$, with $Y \subseteq \mathcal{L}$, therefore satisfies

$$|X| \ge 2^{h+1} = 2Q, \qquad \frac{|Y|}{|\mathcal{L}|} \ge 2^{-(m-h-1)}.$$

Exactly h Bob–Alice pairs remain. We record Bob's next bit after every hypothetical history of earlier Alice bits, which gives $1 + 2 + \cdots + 2^{h-1} = Q - 1$ bits. Fixing these bits and Alice's response rules determines her final output on every retained row. On any $2Q$ of those rows, there are therefore at most 2^{Q-1} output patterns (vectors of outputs on these rows).

Since $m \le Q < 2m$, we have $\varepsilon \le 2^{-2Q}$. Every column in Y must match one of these patterns. Applying Equation (6) to each pattern and adding the bounds gives

$$\frac{|Y|}{|\mathcal{L}|} \le 2^{Q-1}(2^{-2Q} + \varepsilon) \le 2^{-Q}.$$

This contradicts the lower bound on the surviving fraction because $Q \ge m > m - h - 1$. □

Since Equation (6) holds for every set of positions, one array list works against every protocol. We use the accuracy $\epsilon_{\mathrm{a}} = 2^{-4(b+3t)} = 2^{-12t} H_{\mathrm{a}}^{-4}$, which is inverse polynomial in H_{a} for each fixed t. At this accuracy, the finite-field construction of Section H gives such a list with polynomially many arrays in H_{a}, and we use it in place of the set of all Boolean arrays in M_{a}.

Corollary 10.6. *For every fixed $t \ge 4$, a nonempty labeled array list on Ω_{a} satisfying Equation* (6) *with*

$$\epsilon_{\mathrm{a}} := 2^{-4(b+3t)}$$

can be generated in deterministic time polynomial in H_{a}.

The construction and proof are in Section H.

10.4 From a hypergraph to an explicit gap

Given a short protocol, we color each vertex v by Alice's most common response rules for her first t turns among the rows (v, c) with $c \in \mathcal{C}$. A monochromatic edge then gives the many distinct array positions that Lemma 10.5 requires.

Proposition 10.7. *Use the complete address list $\mathcal{C}$ and the explicit array list of Corollary 10.6 in M_{a}. A balanced coloring of $\mathcal{H}$ as in Proposition 10.4 gives*

$$D_A^{\mathrm{alt}}(M_{\mathrm{a}}) \le B_{\mathrm{a}},$$

whereas the absence of a proper S_{a}-coloring gives

$$D_A^{\mathrm{alt}}(M_{\mathrm{a}}) \ge B_{\mathrm{a}} + 2t - 8.$$

For each fixed t, the full table can be printed in deterministic time polynomial in $|V| + H_{\mathrm{a}}$.

Proof. The upper bound is Proposition 10.4. For the lower bound, we put $m = b + 3t - 3$, so $2m - 1 = B_{\mathrm{a}} + 2t - 9$, and suppose that an Alice-first protocol has depth at most $2m - 1$. We pad it to that depth and, for each vertex v, partition its address tuples by Alice's response rules on her first t turns. By the count of response bits above, there are at most $2^{2^t-1} \le S_{\mathrm{a}}$ choices of these rules. A largest class C_v therefore has

$$|C_v| \ge H_{\mathrm{a}}^{q_{\mathrm{a}}}/S_{\mathrm{a}}.$$

We color v by this class's response rules. Since the hypergraph has no proper S_{a}-coloring, some edge e_α is monochromatic, that is, all its vertices receive the same response rules.

For $v \in e_\alpha$ and $g \in [q_{\mathrm{a}}]$, let $J_{v,g} = \{c_g : c \in C_v\}$ be the set of addresses occurring in coordinate g. Since C_v is contained in the product of these projections,

$$\prod_{g=1}^{q_{\mathrm{a}}} |J_{v,g}| \ge |C_v| \ge H_{\mathrm{a}}^{q_{\mathrm{a}}}/S_{\mathrm{a}}.$$

A column (g, α, ℓ) reads only coordinate g of Alice's tuple, so we need one coordinate g that is good for all vertices of the edge. Multiplying the last display over the K_{a} vertices of e_α and choosing g to maximize $\prod_{v\in e_\alpha} |J_{v,g}|$ gives $\prod_{v\in e_\alpha} |J_{v,g}| \ge (H_{\mathrm{a}} S_{\mathrm{a}}^{-1/q_{\mathrm{a}}})^{K_{\mathrm{a}}}$. By AM–GM, this coordinate satisfies

$$\begin{aligned} N := \sum_{v\in e_\alpha} |J_{v,g}| &\ge K_{\mathrm{a}} H_{\mathrm{a}} S_{\mathrm{a}}^{-1/q_{\mathrm{a}}} \\ &= \frac{K_{\mathrm{a}} H_{\mathrm{a}}}{4} = 2^{m-t+1}, \end{aligned}$$

where $S_{\mathrm{a}} = 4^{q_{\mathrm{a}}}$. For each $v \in e_\alpha$ and $u \in J_{v,g}$, we choose a tuple $c(v,u) \in C_v$ with $c(v,u)_g = u$. We retain the corresponding rows and all columns (g, α, ℓ). Their entries are

$$M_{\mathrm{a}}((v, c(v,u)), (g, \alpha, \ell)) = w^\ell(u, \iota_\alpha(v)).$$

The positions $(u, \iota_\alpha(v))$ are N distinct elements of Ω_{a}, because different vertices occupy different positions in the edge and different addresses give different first coordinates. All retained rows give Alice the same response rules for her first t turns. Since $m = b + 3t - 3 \ge 2t$ for $t \ge 4$ and $\epsilon_{\mathrm{a}} \le 2^{-4m}$, the restriction to these rows and columns contradicts Lemma 10.5. Integer depths give $D_A^{\mathrm{alt}}(M_{\mathrm{a}}) \ge 2m = B_{\mathrm{a}} + 2t - 8$.

There are $|V| H_{\mathrm{a}}^{q_{\mathrm{a}}}$ rows and $q_{\mathrm{a}} H_{\mathrm{a}} |\mathcal{W}|$ columns. For fixed t, the number q_{a} is constant and Corollary 10.6 constructs the arrays in polynomial time. Enumerating all tuples and evaluating the displayed entry rule therefore prints the entire table in polynomial time. □

To finish, we choose t so that the gap $2t - 8$ of Proposition 10.7 is at least the requested gap. We then transfer the resulting hardness to the other starter and output conventions with Lemma 10.1.

Proof of Theorem 1.2 and Corollary 1.3. Given a fixed requested gap $k \ge 1$, we choose $t = \lceil k/2 \rceil + 4$. Then $2t - 8 = 2\lceil k/2 \rceil \ge k$. Composing the reduction of Proposition 10.2 with Proposition 10.7, we obtain, in deterministic polynomial time, a table whose cost D_A^{alt} is at most B_{a} in the YES case and at least $B_{\mathrm{a}} + k$ in the NO case, proving Theorem 1.2.

An additive-u estimate of any cost in Lemma 10.1 is, by the triangle inequality, an additive-$(u+2)$ estimate of D_A^{alt} with the same success probability. We choose any fixed integer $k > 2(u + 2)$ and compare the estimate with $B_{\mathrm{a}} + k/2$. This gives $\mathsf{P} = \mathsf{NP}$ for a deterministic estimator and $\mathsf{NP} \subseteq \mathsf{BPP}$ for a bounded-error randomized estimator, proving Corollary 1.3. □

The address length $q_a = 2^{t-1}$ and the degree of the polynomial bounding the running time of the source reduction may both grow with t. Hence the construction is polynomial only for each fixed gap, and this argument gives no growing-error result for strict alternation.

11 Remaining approximation questions

We obtained fixed-error NP-hardness in both models and algorithmic implications for growing errors in the unrestricted model. The quantitative consequences depend on how fast the explicit table grows with the requested gap. Smaller tables would give faster SAT algorithms from the same approximation guarantee, and stronger inapproximability under assumptions excluding those SAT algorithms.

For unrestricted protocols, our deterministic construction outputs a table of size at most $N^{2^{O(k)}}$ for an N-variable sparse formula and gap k (Theorem 8.4). Can smaller lists preserve the restriction bounds needed for the next round, or can another construction improve the $\log \log T$ error scale?

For strict alternation, a construction whose size and running time depend more mildly on the gap could let the gap grow with the table size, giving analogous approximation-to-SAT implications in Yao's model.

A Interlacing estimates and the initial matrix

We first prove the projection inequality and the small-block bounds used in Section C.3. The surviving-block bound proved there (Lemma 6.11) is then used in the separation theorem and the column-loss induction below. We use these estimates to prove Lemma 6.5 and Theorem 6.17, and finish with the initial matrix and template-layer bounds (Lemmas 5.4 and 6.6).

A.1 Projecting column tuples and retained weight

After a protocol prefix, a restriction of an interlace may retain many blocks but only a small fraction of its column weight. We select a few blocks whose projected column tuples still have enough weight for a lower bound. We measure the loss by the logarithm Δ_p below, which increases by one each time the retained column weight is halved.

Fix a Boolean matrix $M : X \times Y \to \{0, 1\}$ with positive normalized row weights μ and column weights ν. Every weight below uses these original weights or their products. For a nonempty $K \subseteq Y^p$, define $\Delta_p(K) = \log(1/\nu^p(K))$, where ν^p is the product weight on Y^p. The projection onto no coordinates has loss $\Delta_0 = 0$, since Y^0 is a single point of weight one.

We show that projection cannot increase the loss, and that projecting onto a suitable set of l of the p coordinates leaves at most an l/p fraction of it.

Lemma A.1. *Let $K \subseteq Y^p$ be nonempty, and write $\pi_J K$ for its projection onto the coordinates in $J \subseteq [p]$. Then*

$$\Delta_{|J|}(\pi_J K) \le \Delta_p(K),$$
$$\Delta_{|J_0|}(\pi_{J_0} K) + \Delta_{|J_1|}(\pi_{J_1} K) \le \Delta_p(K) \qquad (J_0 \sqcup J_1 = [p]).$$

For every integer $1 \le l \le p$, some l-element set J satisfies

$$\Delta_l(\pi_J K) \le \frac{l}{p} \Delta_p(K).$$

Viewing K as a column set of $\mathsf{I}_p(M)$, the interlace on the blocks in J, restricted to the columns $\pi_J K$ and to any retained rows of these blocks, is a submatrix of the restriction of $\mathsf{I}_p(M)$ to these rows and the columns K, and each block keeps its row weights (Observation 4.5).

Proof. The first two inequalities follow from $K \subseteq \pi_J^{-1}(\pi_J K)$ and $K \subseteq \pi_{J_0} K \times \pi_{J_1} K$, using the product weights.

For the averaged bound, we regard ν^p as a distribution, let Z have distribution $Q = \nu^p(\cdot \mid K)$, and write Q_J for the distribution of Z_J. For a distribution P and a positive distribution P' on a finite set, their relative entropy is $D_{\mathrm{KL}}(P\|P') = \sum_z P(z)\log(P(z)/P'(z))$, and entropy is $\mathsf{H}(P) = -\sum_z P(z)\log P(z)$, where zero terms are omitted. Since $Q(z)/\nu^p(z) = 1/\nu^p(K)$ for $z \in K$,

$$D_{\mathrm{KL}}(Q\|\nu^p) = \Delta_p(K).$$

Since Q_J is supported on $\pi_J K$, nonnegativity of relative entropy gives

$$\Delta_{|J|}(\pi_J K) \le D_{\mathrm{KL}}(Q_J\|\nu^J).$$

For a uniformly chosen l-subset J, the chain rule and the fact that conditioning does not increase entropy give

$$\mathbb{E}_J \mathsf{H}(Z_J) \ge \mathbb{E}_J \sum_{i\in J} \mathsf{H}(Z_i \mid Z_1,\dots,Z_{i-1}) = \frac{l}{p}\mathsf{H}(Z).$$

Also, since ν^J is a product, each coordinate contributes to $-\mathbb{E}\log\nu^J(Z_J)$ exactly when it lies in J, which happens with probability l/p. Since $D_{\mathrm{KL}}(Q_J\|\nu^J) = -\mathsf{H}(Z_J) - \mathbb{E}\log\nu^J(Z_J)$, subtracting the entropy inequality yields

$$\mathbb{E}_J \Delta_l(\pi_J K) \le \mathbb{E}_J D_{\mathrm{KL}}(Q_J\|\nu^J) \le \frac{l}{p} D_{\mathrm{KL}}(Q\|\nu^p) = \frac{l}{p}\Delta_p(K).$$

Hence some J satisfies the desired bound. The last assertion is Observation 4.5. □

A.2 Two or three surviving blocks

Later arguments, such as the proof of Lemma 6.11, reduce a large collection of surviving blocks to two or three. We therefore first prove lower bounds for these small collections. Each proof splits into cases according to whether Alice or Bob speaks first and uses both column-weight thresholds of Definition 5.1.

Lemma A.2 (Small block bounds)**.** *Suppose $\mathcal{P}(M; x, H)$ with $H \ge 2$. A restriction of a full-product interlace of M retaining the listed number of blocks, each of at least the listed base row weight, and a column set of at least the listed product weight has at least the listed cost whenever the listed row weight is at most one.*

blocks	*row weight per block*	*product column weight*	*cost*
2	$2x$	$9/50$	H
2	$2x$	$9/25$	$H+1$
3	$4x$	$9/50$	$H+1$

Proof. For the first row, Lemma A.1 gives a singleton column projection of weight $\ge \sqrt{9/50} > 3/10$. Applying $\mathcal{P}(M; x, H)$ with $j = 1$ to that block gives cost $\ge H - 1 \ge 1$, so the restriction is nonconstant. If Alice speaks first, we follow that block into a child retaining base row weight $\ge x$.

The retained columns are unchanged, so the case $j = 1$ of $\mathcal{P}(M; x, H)$ again gives suffix cost at least $H - 1$. If Bob speaks first, his denser child has column weight $\geq 9/100$ and hence a singleton projection of weight $\geq 3/10$. The suffix at that child again costs at least $H - 1$. Including the first bit gives cost at least H.

For the second row, the first row gives cost at least $H \geq 2$, hence nonconstancy. If Alice speaks, some singleton projection has column weight $\geq \sqrt{9/25} = 3/5$, and by $j = 0$ in $\mathcal{P}(M; x, H)$, a child retaining half the row weight of that block costs at least H. If Bob speaks, his denser child has column weight $\geq 9/50$, so the first row gives suffix cost at least H. Including the first bit gives cost at least $H + 1$.

For the third row, the projection onto any two blocks has product column weight $\geq 9/50$ and row weights $\geq 2x$, so the first row shows that the restriction is not constant. If Alice speaks, we assign each block to a child retaining at least half its row weight. One child receives two blocks of weight $\geq 2x$, and the first row gives suffix cost H. If Bob speaks, his denser child retains column weight $\geq 9/100$. The averaged projection inequality gives two of its three coordinates with product weight at least

$$(9/100)^{2/3} > 9/50.$$

Their row weights are still $\geq 4x \geq 2x$, so the first row again gives suffix cost H. Including the first bit proves the claim. By Observation 4.5, each projection is a restriction of the original matrix. □

The next theorem shows that a short protocol for a width-u interlace must begin with $\log u$ Alice bits, and that each node reached by these bits keeps most of the row weight of one block and little of every other block. Its proof uses Lemma 6.11, whose proof in Section C does not use the theorem.

Theorem A.3. *Suppose $\mathcal{P}(M; \rho, d)$ with $d \geq 2$. Let $u \geq 2$ be dyadic and $1 \geq \xi > 2u\rho$. Interlace M at width u using an exactly 8-wise independent tuple list with the base-column marginal, as in Definition 4.4. Every restriction retaining base row weight at least ξ in every block and tuple-label weight at least $18/25$ has the following property: every protocol of depth at most $d + \log u$ for it begins with $\log u$ Alice bits along every path. The u nodes at depth $\log u$ are in bijection with the blocks: each such node retains more than $\xi - 2(u-1)\rho$ of its assigned block and less than 2ρ of every other block. Full products are a special case, and fixed bits count as transmitted bits.*

Proof. We maintain the following invariant at each protocol-tree node of depth $r < \log u$: there are $\kappa = u2^{-r}$ assigned blocks, each retaining base row weight $\geq \xi 2^{-r} > 2\kappa\rho$, and the retained column set is the original one. The root has this property. By Observation 4.5, the projection onto the coordinate of one assigned block is a restriction of M with column weight $\geq 18/25 > 3/5$, so $\mathcal{P}(M; \rho, d)$ gives cost at least $d \geq 2$ and the node is not a leaf. Since $1 \geq \xi > 2u\rho$, every row threshold $2^a\rho$ with $1 \leq a \leq \log u$ used below is less than one.

If Bob speaks, one child retains column weight $\geq 9/25$ and all κ assigned blocks. By Lemma 6.11 with $a = \log \kappa$, it needs $d + \log \kappa$ bits, but has only $d + \log u - r - 1 = d + \log \kappa - 1$ left. Thus Alice speaks. We assign each block to a child retaining at least half its weight. If a child received more than $\kappa/2$ blocks, it would contain at least $\kappa/2 + 1$ blocks of weight $\xi 2^{-r-1} > \kappa\rho$, and Lemma 6.11 with $a = \log \kappa$ would again require $d + \log \kappa$ bits, exceeding the $d + \log \kappa - 1$ left. Both children therefore receive exactly $\kappa/2$ blocks. This proves the invariant, and hence that Alice sends the first $\log u$ bits. In particular, a fixed Alice bit is impossible, since it would assign every block to one child.

At depth $\log u$, each node has one assigned block of retained base row weight at least $\xi/u > 2\rho$. If such a node retained two blocks of weight at least 2ρ, they would require $d + 1$ bits by Lemma 6.11

with $a = 1$, whereas only d remain. Every other block thus has weight less than 2ρ. Each original block is assigned to exactly one node, and only Alice has spoken, so its weight at the other $u - 1$ nodes totals less than $2(u-1)\rho$, proving the final bound. □

A.3 Induction on column loss

For a restriction of $\mathsf{I}_p(M)$ with surviving columns K, Bob's messages shrink K and may increase the loss $\Delta_p(K)$. The following induction gives lower bounds that tolerate this loss, and it uses only lower bounds on the base matrix.

Lemma A.4. *Let $C \geq 3$, $0 < \theta \leq 1/4$, and $0 < \delta \leq 1/64$. Suppose $\mathcal{P}(M; \theta, C-1)$ (Definition 5.1) and every restriction of M retaining row weight at least θ and column weight at least $1 - \delta$ has cost at least C. Let $e \in \{0, 1\}$, let $a \geq 2e$ be an integer with $2^a\theta \leq 1$, and let $p > 2^{a-1}$ be an integer. If a restriction G of $\mathsf{I}_p(M)$ retains base row weight at least $2^a\theta$ in every block and its nonempty column set K satisfies*

$$\Delta_p(K) \leq (p - 2^{a-1})\delta + \frac{4e}{3},$$

then $\mathrm{D}(G) \geq C + a - e$. Here C is only a lower bound, and we do not assume $\mathrm{D}(M) = C$.

Proof. Say that G is in state (a, e), or (a, e, p), when it satisfies the hypotheses for these values. We prove the lower bound by induction on

$$\mathfrak{m}(a, e) := 2a + 1 - e.$$

An Alice bit halves the row thresholds and lowers a. A Bob bit increases the loss by at most one. We cover this increase either by raising e to one, which adds $4/3$ to the allowed loss, or by projecting onto about half the blocks with Lemma A.1. Each such step lowers $\mathfrak{m}$.

Since $C - 1 \geq 2$, $4\theta \leq 1$ and $\mathcal{P}(M; \theta, C-1)$ hold, we may apply Lemma 6.11 with $d = C - 1$ and $x = \theta$ to full-product interlaces of M. At column weight $9/25$, its cases $a = 1, 2$ give cost at least C for two blocks of row weight 2θ, and at least $C + 1$ for three blocks of row weight 4θ. We call these the two-copy and three-copy bounds. Every projection below with loss $\leq 3\delta + 4/3$ has column weight $> 9/25$, because

$$3\delta + \frac{4}{3} \leq \frac{265}{192} < \frac{7}{5} < -\log(9/25).$$

The last inequality follows from $2^7 \cdot 9^5 < 25^5$.

We first verify the base states, where an Alice bit would leave the range $a \geq 2e$. For $(a, e) = (0, 0)$, the averaged projection bound in Lemma A.1 gives a singleton $J \subseteq [p]$ with

$$\Delta_1(\pi_J K) \leq \frac{p - 1/2}{p}\delta < \delta.$$

Since $2^{-v} \geq 1 - v$ for $0 \leq v \leq 1$, applied with $v = \Delta_1(\pi_J K)$, the projected column weight is greater than $1 - \delta$. The block in J retains row weight $\geq \theta$, so the hypothesis gives cost at least C.

For $(a, e) = (2, 1)$, the loss bound and Lemma A.1 give a three-coordinate projection with loss at most

$$\frac{3}{p}\left((p-2)\delta + \frac{4}{3}\right) \leq 3\delta + \frac{4}{3} < -\log(9/25).$$

Its column weight exceeds $9/25$, and each block retains weight $\geq 4\theta$. The three-copy bound above gives cost at least $C + 1 = C + 2 - 1$. There is one further base case, $(a, e, p) = (3, 1, 5)$, since the

Bob step below needs $p \geq 6$. Here $\Delta_5(K) \leq \delta + 4/3 < -\log(9/25)$, and every block retains weight $\geq 8\theta$. Its column weight therefore exceeds $9/25$, and Lemma 6.11 with $a = 3$ applies to these five blocks, giving $C + 2 = C + 3 - 1$ bits.

We also handle the Bob-first case of $(a, e) = (1, 0)$ now, since $(1, 1)$ is not a valid state. The general step below covers its Alice-first case. Some singleton projection has loss $\leq (p-1)\delta/p < \delta$, so its column weight exceeds $1 - \delta$. Its block retains row weight $\geq 2\theta \geq \theta$, so the hypothesis gives it cost at least $C \geq 3$, and the restriction is nonconstant. After a Bob bit, the denser child has a two-coordinate projection of loss at most

$$\frac{2}{p}\left((p-1)\delta + 1\right) \leq 2\delta + 1 < 3\delta + \frac{4}{3},$$

so the two-copy bound gives cost at least C after that bit. Every other valid state has $a \geq 2$, $p \geq 3$, and base row weight $\geq 4\theta$. By Lemma A.1, some three-coordinate projection has loss $\leq 3\delta + 4/3$, so by the three-copy bound the restriction is not constant.

Now consider any case not yet treated, and the first bit of a protocol for its restriction. If Alice speaks, we assign every block to a child retaining at least half its row weight. For child $i \in \{0, 1\}$, let p_i be the number of assigned blocks and let L_i be the loss after projecting the retained column set onto those blocks. Then $p_0 + p_1 = p$, and the partition inequality in Lemma A.1 gives $L_0 + L_1 \leq \Delta_p(K)$. Put

$$s_i = \max\{p_i - 2^{a-2}, 0\}.$$

Since $s_0 + s_1 \geq p - 2^{a-1} > 0$, at least one s_i is positive. The loss bounds required of the children with $s_i > 0$ satisfy

$$\begin{aligned}\sum_{i:s_i>0}\left(s_i\delta + \frac{4e}{3}\right) &\geq (p - 2^{a-1})\delta + \frac{4e}{3} \\ &\geq \sum_{i:s_i>0} L_i.\end{aligned}$$

At least one eligible child contributes the constant $4e/3$, and the final inequality uses the assumed loss bound on K and $L_0 + L_1 \leq \Delta_p(K)$. Consequently some child with $s_i > 0$ has $L_i \leq s_i\delta + 4e/3$. Each of its assigned blocks retains weight $\geq 2^{a-1}\theta$, and there are $p_i > 2^{a-2}$ such blocks. Since $(a, e) \neq (2, 1)$, we have $a - 1 \geq 2e$, so the projection of this child onto its assigned blocks, a submatrix of the child by Lemma A.1, is in state $(a - 1, e, p_i)$. The induction hypothesis gives cost at least $C + a - 1 - e$ in that child, and Alice's bit gives $C + a - e$. For $(a, e) = (1, 0)$, this is the step to state $(0, 0)$.

If Bob speaks, one child has at least half the current column weight, so its loss is $\leq \Delta_p(K) + 1$. When $e = 0$ and $a \geq 2$, this child satisfies state $(a, 1)$, because

$$\Delta_p(K) + 1 \leq (p - 2^{a-1})\delta + \frac{4}{3}.$$

Its induction cost is at least $C + a - 1$, and adding Bob's bit proves the claim.

It remains to treat a Bob bit when $e = 1$ and $a \geq 3$. The extra term $4/3$ has then already been used, so we instead project the denser child onto about half the blocks and show that the projection is in state $(a - 1, 1)$. The five-block base case already handles $p = 5$, so $p \geq 6$. We keep $p' = \lceil p/2 \rceil$ blocks, chosen below by Lemma A.1. Then

$$\frac{1}{2} \leq \frac{p'}{p} \leq \frac{4}{7}, \qquad \frac{p'}{p}(p - 2^{a-1}) \leq p' - 2^{a-2}, \qquad \frac{4}{7}\left(\frac{4}{3} + 1\right) = \frac{4}{3}.$$

By the averaged projection bound in Lemma A.1, some p' coordinates of the denser child therefore have loss at most

$$\frac{p'}{p}\left((p-2^{a-1})\delta+\frac{4}{3}+1\right)\le(p'-2^{a-2})\delta+\frac{4}{3}.$$

The projected child retains base row weight $\ge 2^{a-1}\theta$ in each of its $p'>2^{a-2}$ blocks, so it is state $(a-1,1)$. Its induction cost is at least $C+a-2$, and adding Bob's bit gives $C+a-1$, as required. □

We prove Lemma 6.5 by applying Lemma A.4 to $p=9u/16$ blocks. The induction needs more than $u/2$ blocks, and the extra $u/16$ blocks give slack $(9u/16-u/2)\delta=u/1024$ at $\delta=1/64$ in the loss bound. The retained column weight may therefore be as small as $2^{-u/1024}$.

Proof of Lemma 6.5. Let $K\subseteq Y^q$ be the retained column set, where q is the width of the interlace. Since $u\ge 16$ is dyadic, $9u/16$ is an integer. We select a set J of exactly $p=9u/16$ blocks retaining base row weight $\ge ux$ and project K onto them. The hypothesis $\nu^q(K)\ge 2^{-u/1024}$ and Lemma A.1 give $\Delta_p(\pi_J K)\le\Delta_q(K)\le u/1024$, and the projected interlace is a restriction of the original one. Apply Lemma A.4 with

$$a=\log u,\quad \theta=x,\quad \delta=1/64,\quad C=H,\quad e=0.$$

Here $x\le 1/16$, $p>2^{a-1}=u/2$, and $2^a x=ux\le 1$. The hypothesis $\mathcal{P}(M;x,H)$ implies the weaker bounds $\mathcal{P}(M;x,H-1)$ and, since $63/64\ge 3/5$, cost $\ge H$ at column weight $63/64$. These are the two lower-bound hypotheses of Lemma A.4. Finally,

$$\Delta_p(\pi_J K)\le u/1024=(9u/16-u/2)/64=(p-2^{a-1})\delta.$$

The induction gives cost at least $C+a=H+\log u$. □

We use Lemma A.4 to restore both deletion bounds at full cost for $R^+=(\mathsf{I}_Q(W))^{\mathsf{T}}$.

Proof of Theorem 6.17. Put $C=\mathrm{D}(W)$ and $f=\lceil b/\delta\rceil$. The assumptions $b\ge 1$, $\delta\le 1/64$ and $Q\delta\ge 80b$ imply

$$f\delta\ge b,\qquad (f-1)/Q<1/80,\qquad f<Q/40,\qquad Q\ge 5120.$$

Indeed, $b/\delta\ge 64$, $Q\delta\ge 80b$, and $f\le b/\delta+1<2b/\delta$. Fix $j\in\{0,1\}$ and put

$$a_j=r-j,\qquad p_j=2^{a_j-1}+f,\qquad t_j=2^{a_j}\theta.$$

These choices give $a_j\ge 0$, $2^{a_j-1}<p_j<Q$, and $t_j<1$.

Consider a restriction of R^+ to which $\mathcal{P}(R^+;2^{-b},C+r)$ applies with this j, and transpose it to a restriction of $\mathsf{I}_Q(W)$. By Definition 5.1, the latter retains tuple weight $\ge 2^{-b}$ and total block-row weight $\ge(3/5)2^{-j}$. If fewer than p_j blocks had base row weight at least t_j, then, as each block has weight $1/Q$, this total would be less than

$$\frac{p_j-1}{Q}+t_j<2^{-j-1}+\frac{1}{80}+\frac{1}{20\cdot 2^j}\le\frac{3}{5}\,2^{-j},$$

a contradiction, where the last inequality uses $j\le 1$. We select a set J of p_j blocks meeting the row-weight threshold. For the retained tuple set K, Lemma A.1 gives

$$\Delta_{p_j}(\pi_J K)\le\Delta_Q(K)\le b\le f\delta=(p_j-2^{a_j-1})\delta.$$

The projected restriction therefore satisfies the hypotheses of Lemma A.4 with $e = 0$, $a = a_j$, $p = p_j$, and row threshold $2^{a_j}\theta = t_j$. Its two lower-bound hypotheses are those assumed for W, and $\theta < 1/4$ follows from $Q\theta \le 1/20$. The lemma gives cost at least $C + a_j = C + r - j$, so Observation 4.5 and transposition prove $\mathcal{P}(R^+; 2^{-b}, C + r)$.

Sending the copy index and then running a protocol for W costs $r + \mathrm{D}(W)$ bits. This gives $\mathrm{D}(R^+) = \mathrm{D}(W) + r$, and it turns any upper bound y on $\mathrm{D}(W)$ into the bound $y + r$ on $\mathrm{D}(R^+)$. By definition the transpose has n^Q rows and Qm columns. □

A.4 The initial matrix and template layers

For the initial matrix $R_0(z, j) = z_j$ on $\{0,1\}^{p_0} \times [p_0]$, we use real rank. Every restriction tested by the deletion bounds still has large rank, and large rank forces large depth.

Proof of Lemma 5.4. Since p_0 is a power of two, $a = \log p_0$ is an integer. Fix $j \in \{0,1\}$, a retained row set $V \subseteq \{0,1\}^{p_0}$ with $|V| \ge 2^{p_0 - L}$, and a retained column set $J \subseteq [p_0]$ with $|J| \ge \lceil (3/5)2^{-j}p_0 \rceil$. Under uniform weights, these are the restrictions tested by $\mathcal{P}(R_0; 2^{-L}, d_0)$. Let v be the dimension of the real span of V. The matrix with row set V has rank v (with all p_0 columns), so by Mackenzie [17, Lemma 4.1(1)] it has at most 2^v distinct rows, that is, $|V| \le 2^v$. The lower bound on $|V|$ therefore implies $v \ge p_0 - L$.

Deleting the $p_0 - |J|$ coordinates outside J can lower the row rank by at most $p_0 - |J|$. Thus the restricted initial matrix has rank r satisfying

$$r \ge |J| - L \ge \frac{3}{5}\,2^{-j}p_0 - L > 2^{a-j-1}.$$

The strict inequality follows from

$$\frac{3}{5}\,2^{-j}p_0 - 2^{a-j-1} = \frac{p_0}{10 \cdot 2^j} \ge \frac{p_0}{40} > L,$$

where $p_0/40 - L > 0$ follows from $p_0 \ge 64(L+1)$.

By Mackenzie [17, Lemma 4.1(2)], a matrix of depth at most $a - j \ge 1$ has rank at most 2^{a-j-1}. Hence the restriction has depth at least $a + 1 - j = d_0 - j$, proving $\mathcal{P}(R_0; 2^{-L}, d_0)$. At $j = 0$, this gives $\mathrm{D}(R_0) \ge a + 1$. Conversely, Bob sends the column index in a bits and Alice sends the corresponding bit, so $\mathrm{D}(R_0) \le a + 1$, and hence $\mathrm{D}(R_0) = a + 1 = d_0$. The unit-vector rows form the $p_0 \times p_0$ identity submatrix, so $\mathrm{rank}(R_0) = p_0 = 2^{d_0 - 1}$.

By minimality of $p_0 = \mathrm{pow}_2(64(L+1))$, $p_0 < 128(L+1)$. If $L = \lceil \log n \rceil$ and $n \ge 4$, then $L + 1 \le \log n + 2 \le 2 \log n$, so $p_0 < 256 \log n$ and $2^{p_0} < n^{256}$. Enumerating the 2^{p_0} bit strings and the p_0 column indices therefore prints the table deterministically in polynomial time. □

Proof of Lemma 6.6. We treat the full-product and explicit cases together, except for the step that turns many blocks above the threshold into a lower bound. Fix $j \in \{0,1\}$, put $\beta_j = (3/5)2^{-j}$, and let $\lambda = 255/256$ be the total weight of the template rows (Definition 6.2). After transposition, each of the three deletion bounds concerns a restriction of an interlace with the following data. Here q is the interlace width, u is the block parameter to which we apply Lemma 6.5 or Theorem 6.12, and the last column is a lower bound on the average block retention.

target	base	q	u	average block retention
A	R	$q_1 = s - 2$	$s2^{-j}$	β_j
R_{temp}	A	q_2	$q_2 2^{-j}$	β_j
N_{ref}	R_{temp}	s	$s2^{-j}$	$(\beta_j - (1-\lambda))/\lambda$

In each case the retained tuple-label event Y has weight $\geq x$. For N_{ref}, we discard the non-template rows. They have total weight $1-\lambda$, which gives the displayed average retention. We then map each retained Ω label to its bin part, the tuple $(j_1,\dots,j_s)$ in the full-product case and its $\mathcal{J}_{\mathrm{g}}$ label in the explicit case, forgetting the tail ξ, which template entries ignore. This map induces the original weights on its target labels, so its image still has weight $\geq x$, and choosing one retained label for each image label gives the template entries (Observation 4.5).

For the three rows, let m be the number of blocks retaining base row weight at least the threshold $t = u\rho$, ux, and ux, respectively. The block-counting bound $q\beta \leq m + qt$ of Section 4, with β the tabulated retention, gives the following lower bounds for m/u:

$$\frac{3}{5}(1-2/s) - q_1\rho, \qquad \frac{3}{5} - q_2 x, \qquad \frac{2^j(\beta_j - (1-\lambda))}{\lambda} - sx.$$

All three exceed $9/16$ by the numerical bounds of Lemma F.1, which both parameter schedules satisfy. They also give the row thresholds $u\rho \leq s\rho \leq 1$ and $ux \leq \max\{q_2 x, sx\} \leq 1$, and $u \geq s/2 \geq 16$ since $s \geq 256$.

Full-product case. Here $x = 2^{-8L} \geq 2^{-u/1024}$ by Lemma F.2, so Lemma 6.5 applies to the more than $9u/16$ blocks above the threshold.

Explicit case. We apply Lemma 6.10 to Y with $v = x$ and $\tau = 1$. Outside at most $e(x)$ exceptional coordinates, the marginals of $\nu \mid Y$ on at most eight coordinates stay close to the product weights, and $e(x) \leq s/128 \leq u/32$ by Lemma F.3. Discarding the blocks at these coordinates leaves more than $9u/16 - u/32 > u/2$ blocks above the threshold. Since u is dyadic, at least $u/2+1$ of them remain. These blocks and $K = Y$, for which $(\nu \mid Y)(K) = 1$, meet the hypotheses of Theorem 6.12 with $a = \log u$, which gives the costs computed below.

In both cases the base bound is $\mathcal{P}(R;\rho,d)$ for the first row and, for the other two rows, the deletion bounds just established for A and R_{temp}. The required costs are therefore

$$d + \log s - j = H_1 - j, \qquad H_1 + \log q_2 - j = H_2 - j, \qquad H_2 + \log s - j = C - 1 - j,$$

which are the three asserted pairs of deletion bounds.

It remains to prove the collision bound for the interlace on Γ. The first q_1 positions and three distinct tail positions give $q_1 + 3 = s + 1 = (2s)/2 + 1$ blocks. By hypothesis, each retains base row weight $\geq x_{\mathrm{cap}} = 2s\rho$, and the retained list labels have weight greater than $1/2 > 9/25$. In either case the list Γ is exactly 8-wise independent with the base-column marginal, so Lemma 6.11 with $a = \log(2s)$ gives $d + \log(2s) = H_1 + 1$ bits. □

B Bounds on transferred inputs

We prove Lemmas 6.14 and 6.15, which bound the $\mathcal{D}$-weight arriving at a stopping node (defined in Lemma 6.13) and the weight sent to the stopping node assigned to another copy. Throughout, we assume the common hypotheses (H1)–(H4) of Section 6.5 and drop the subscript g of the explicit construction. In each source copy, the $\mathcal{D}$-inputs have total weight $\mu = 1/256$, and the template rows and the Ω-inputs each have total weight $\lambda = 255/256$.

Transferred weights below are in units of one complete copy, whereas the fraction ε_i of Lemma 6.15 is normalized within copy i's Ω-inputs. Lemma 6.10 controls conditional marginals on at most eight nonexceptional coordinates, and the projection inequality following it converts this control into bounds in the original product weights. By Observation 4.5, choosing one retained representative for each projected label turns these projected tables into submatrices of the restriction.

B.1 Two template inputs at one stopping node

The next lemma uses the orientation of B_{raw}, transposed relative to N. Its rows are template-label pairs, and its two column families come from one copy's Ω-inputs and another copy's $\mathcal{D}$-inputs.

At the node assigned to copy i, an Ω-input of copy i reads the template label at coordinate i, and a $\mathcal{D}$-input of a copy $j \neq i$ reads the template label at coordinate j. By (H3), these two labels are independent under the original weights. The next lemma shows that if enough of both families is retained, the protocol needs one bit more than the cost H_2 given by $\mathcal{P}(R_{\mathrm{temp}}; x, H_2)$, even if the two families have the same column labels in J_2.

Lemma B.1 (Two template inputs)**.** *Let $A : \Gamma \times J \to \{0,1\}$ satisfy $\mathcal{P}(A; x, H_1)$ with $H_1 \geq 2$, and let q_2 be dyadic with $2q_2x \leq 1/8$. Take two weighted lists S_o, S_f of words in J^{q_2}, each exactly 8-wise independent with the column weights of A as marginal, and use product weights on $S_o \times S_f$. Put $H_2 = H_1 + \log q_2$ and give $J_2 = [q_2] \times \Gamma$ the product of uniform index weights and the row weights of A on Γ. Rows are pairs of template labels $(z_o, z_f) \in S_o \times S_f$. The columns form two separately tagged copies of J_2, with entries*

$$A(\gamma, (z_o)_\alpha) \quad \textit{and} \quad A(\gamma, (z_f)_\alpha),$$

respectively. The tags remain distinct even when the inner indices coincide. Retain a set of pairs of original product weight at least $49/64$, columns reading z_o of normalized weight at least $1 - \varepsilon$, and columns reading z_f of normalized weight at least δ, each tagged copy of J_2 having total weight one, where

$$\varepsilon \geq 0, \qquad 0 < \delta \leq 1, \qquad \delta > 2q_2x + \varepsilon.$$

Then the retained matrix requires at least $H_2 + 1$ bits.

Proof. Put $t = 2q_2x$. We group columns by their tag and index α. These $2q_2$ groups have total retained Γ-weight at least

$$q_2(1 - \varepsilon + \delta) > q_2(1 + t).$$

If at most q_2 groups had weight at least t, this sum would be at most $q_2 + q_2t$. Thus at least $q_2 + 1$ groups retain weight $\geq t$.

We transpose the matrix. These groups are now distinct row blocks of an interlace of A. Concatenating S_o and S_f with product weights gives an exactly 8-wise independent tuple list on $2q_2$ tagged coordinates, since any set of at most eight coordinates splits between the two independent lists, and each list has the required marginals on its part. We keep the surviving pair labels as tuple labels, whose original weight is $\geq 49/64 > 9/25$, and apply Lemma 6.11 with the integer $a = \log(2q_2)$ and row threshold $2q_2x \leq 1/8$. This gives $H_1 + \log(2q_2) = H_2 + 1$ bits, as required. □

In the application (the proof of Lemma 6.15), we fix a bin at each of two copies i and j, and the pairs (z_o, z_f) occur only in the bin fixed at copy i, which we call the *distinguished bin*. In the other bins of copy i, we know nothing about the entries in the columns of copy j, so the next claim extends Lemma B.1 to allow arbitrary entries there.

Claim B.2. *Use the source weights of Definition 6.2 and the common hypotheses (H1)–(H4) of Section 6.5, and put $\lambda_0 = 2q_2x$, $\tau_0 = 2sx$, $C = H_2 + \log s + 1$, and*

$$a_{\min} = \frac{99/100 + \lambda^2 - 1}{\lambda^2}.$$

Retain Ω-inputs of normalized weight at least $1 - \varepsilon$, and separately tagged $\mathcal{D}$-inputs of another copy whose weight-preserving map $(\mathrm{tail}(\gamma), j, \ldots, j) \mapsto j$ to J_2 has retained image weight at least

δ. Number the bins so that the distinguished bin is bin one. This bin retains a set $P_* \subseteq S^2$ of pairs (z_o, z_f) of product weight at least $a_{\min}$, and every other bin retains a set of labels of S-weight at least τ_0. As in Definition 6.2, write a column in Ω as $Z = (\xi, j_1, \dots, j_s)$, where $\xi \in Y^{s+1}$ and $j_p = (\alpha_p, \gamma_p) \in J_2$. Require entries

	$Z \in \Omega$	*other-copy image* (α, γ)
$(z_o, z_f) \in P_*$	$A(\gamma_1, (z_o)_{\alpha_1})$	$A(\gamma, (z_f)_\alpha)$
$(p, z),\ p \neq 1$	$A(\gamma_p, z_{\alpha_p})$	*arbitrary.*

If $\varepsilon \geq 0$, $0 < \delta \leq 1$, and $\delta > \lambda_0 + 4\varepsilon$, then every such matrix requires at least C bits.

Proof. Suppose this matrix has a protocol of depth at most $C-1$. In each bin, we follow the retained set of template labels through Alice's bits until Bob first speaks, using the tracking rule stated before Lemma 6.8 with $t = \tau_0$. The only marked set is the set P_* of pairs in the distinguished bin. For the thresholds, we measure the weight of P_* and of its unsplit descendants by their projections onto z_o, and a split replaces the two parts by their projections onto z_o, which are sets of labels from S. A pair set weighs no more than its projection onto z_o, so by the tracking rule an unsplit pair set loses less than τ_0 of its original weight. The numerical bounds of Lemma F.1 give $\tau_0 \leq 10^{-3}$ and $a_{\min} - \tau_0 > 49/64$. An unsplit set of pairs therefore retains weight $> 49/64$, and every retained set at depth r has original weight $\geq \tau_0 2^{-r}$.

The hypothesis on δ implies $\varepsilon < 1/4$, so the columns from Ω retain normalized weight $> 3/4$, and so does their projection onto any bin. Since $H_2 > 0$, the bounds $\mathcal{P}(R_{\text{temp}}; x, H_2)$ rule out a leaf containing a retained set of weight $\geq x$. A retained set also cannot reach depth $\log s + 1$: it would have weight at least x, but only $H_2 - 1$ bits would remain. Thus every retained set reaches a Bob node by depth $\log s$.

At each stopping node v, let h_v be the number of tracked sets and z_v the number of unsplit pair sets among them. By the tracking rule, $\sum_v (h_v + z_v)$ equals $s+1$, the s initial sets plus the one initial mark. Thus the stopping nodes satisfy the hypotheses of Lemma 6.7 with $q = s$. Let v, at depth r, be the node given by that lemma, and write $\kappa = \kappa_v = s2^{-r}$, $h = h_v$, $z = z_v$. The sets at v have distinct bin coordinates, since descendants in different branches never rejoin, and each has original weight $\geq \tau_0 2^{-r} = 2\kappa x$. By Lemma 6.7, $h \geq \kappa/2 + 1$ if $\kappa \geq 2$, while for $\kappa = 1$, either $h \geq 2$ or $h = z = 1$.

We first treat the case $\kappa = 1$, $h = z = 1$, before restricting columns or projecting the pair. Since only Alice has spoken, both tagged column families are still unchanged. We project columns from Ω onto j_1 and other-copy columns onto their J_2 images, choosing one surviving column per image label. These maps preserve the original weights, so the projected weights are at least $1 - \varepsilon$ and δ. The set of pairs has weight $> 49/64$. Thus Lemma B.1 requires $H_2 + 1$ bits, while the stopping node has only $C - 1 - r = H_2$ bits left. The original weights on pairs of template labels are the product of their weights on S, as required by that lemma.

In all other cases, at Bob's bit we choose the child with more weight from Ω, which retains weight $> 3/8 > 9/25$. We replace any set of pairs by its projection onto z_o, which has at least its weight. The template entries at the chosen child now form a restriction of an interlace of R_{temp} with one block for each tracked bin. The original column list of this interlace is the bin part $(j_1, \dots, j_s)$ of Ω, which by (H3) has exactly 8192-wise independent bin coordinates under the original weights in either construction. We forget the tail ξ, which template entries ignore, but keep the whole tuple $(j_1, \dots, j_s)$ in the full-product case and its whole list index in the explicit case. The retained image has original weight $> 3/8$, and choosing one surviving row or column for each projected label gives the required entries (Observation 4.5). Thus Lemma 6.11 applies in both constructions. For $\kappa \geq 2$ we take $a = \log \kappa$. The $\kappa/2 + 1$ blocks have row weight $\geq 2\kappa x \geq \kappa x$, so the lemma requires

$H_2 + \log \kappa$ bits, but the child has only $H_2 + \log \kappa - 1$ bits left. In the remaining case $\kappa = 1$, $h \geq 2$, there are at least two blocks of weight $\geq 2x$, and $a = 1$ gives the stronger bound $H_2 + 1$. The conditions $2^a x \leq 1$ hold since $2^a \leq s$ and $sx \leq 10^{-4}$ by (10). This contradiction does not use the other-copy entries outside the distinguished bin, which proves the claim. □

B.2 Total weight arriving at a stopping node

We first bound the total $\mathcal{D}$-weight arriving at one stopping node, including weight from copies without an assigned node. Too much incoming weight leaves many blocks of substantial row weight, and Theorem 6.12, applied after conditioning, then requires more communication than remains.

Proof of Lemma 6.14. At a node at depth $\log w$ we put $\kappa = 1$, so that both bounds read $2s\kappa$. Fix any collection of copies, and suppose the weight of their inputs from $\mathcal{D}$ arriving at one stopping node is at least $2s\kappa$. Let $d_i \in [0, 1]$ be the fraction of copy i's inputs from $\mathcal{D}$ reaching this node, and set $d_i = 0$ outside this collection. Since each such family has weight $\mu = 1/256$,

$$D_0 := \sum_i d_i \geq 512 s\kappa.$$

These fractions are fixed by the Alice-only transcript. At a node where Bob first speaks before depth $\log w$, we choose a child whose column set Y^* has original weight $\geq \theta/2$, which is possible since the tuple event Y' of Lemma 6.13 has weight $\geq \theta$. At a node at depth $\log w$ we put $Y^* = Y'$.

To apply the template lower bounds, we need retained column tuples in which at least half of the incoming weight D_0 comes from copies i at which the tuple's label is a template row. Let T_i indicate that the label at coordinate i of a column tuple is a template row. When the type, bin, and item labels are drawn according to their fixed weights, the centered variables $T_i - \lambda$ are 256-wise independent. The weighted form of Lemma C.1 with $r = 128$ and coefficients d_i, followed by Markov's inequality, gives

$$\Pr\left\{\sum_i d_i T_i < D_0/2\right\} \leq C_{128} D_0^{-128} \leq C_{128} s^{-128} \leq \theta/4.$$

Here $C_{128} = (2^{14}/(\lambda - 1/2)^2)^{128}$ is the fixed constant of (8), and the last inequality is a row of (10). Thus Y^* intersects the event $\sum_i d_i T_i \geq D_0/2$ in weight $\geq \theta/4$. This event depends only on the type/bin/item sequence. Averaging over that sequence and the base-label sequence, we fix a configuration whose retained set Z of template-label sequences has original weight $\geq \theta/4$ and whose template positions have total weight $\sum_i d_i T_i \geq D_0/2$. By (H3), fixing those two sequences leaves the template-label sequence with its original weights and exactly 8192-wise independent J coordinates $z_{i,\alpha}$. We apply Lemma 6.10 to Z with $v = \theta/4$ and $\tau = 1$, which gives an exceptional set of at most $e(\theta/4) \leq sq_2$ coordinates (i, α) by (10).

By Definition 6.2, each retained input (α, γ) from $\mathcal{D}$ at a copy i with $T_i = 1$ reads $A(\gamma, z_{i,\alpha})$, regardless of its fixed bin. We group these inputs by (i, α) and measure each group within its fiber, that is, by the fraction of the inputs from $\mathcal{D}$ of copy i with index α that it contains. Since the index α is uniform, the q_2 groups of copy i then have total weight $q_2 d_i$, so the groups with $T_i = 1$ have total base-row weight $\geq q_2 D_0/2$. Set $K_0 = 2s\kappa q_2$ and $t = K_0 x$. From $w^2 q_2 x \leq 1$ and $512 s\kappa \leq D_0 \leq w$,

$$t \leq \frac{2s\kappa}{w^2} \leq \frac{1}{256w} < 1.$$

There are at most wq_2 groups, each of weight at most one. If h groups have weight at least t, the total weight of all groups is at most $h + wq_2t$, so

$$h \geq \frac{q_2 D_0}{2} - \frac{K_0}{w} \geq K_0(128 - 1/w) > K_0 \geq K_0/2 + sq_2.$$

After deleting the exceptional coordinates, more than $K_0/2$ blocks outside that exceptional set remain, each of base row weight $\geq K_0 x$. We retain $K_0/2 + 1$ of them and all labels in Z. Distinct triples (i, α, γ) label distinct incoming rows, so these rows and Z form a restriction with the required entries. Row weights use the original Γ weights, and the retained column set is all of Z, so it has weight one under the conditional weights on Z.

Since K_0 is dyadic, we can apply Theorem 6.12 with base bounds $\mathcal{P}(A; x, H_1)$ and integer $a = \log K_0$. It requires $H_1 + \log K_0 = C + \log \kappa$ bits. Since the protocol has depth at most $C + \log w - 1$, if Bob first speaks at a depth $r < \log w$, the chosen child has only $C + \log \kappa - 2$ bits left, and a node at depth $\log w$ has only $C - 1$. This proves the bound for every collection of copies. □

B.3 Transfer to another copy's assigned node

We next bound the $\mathcal{D}$-weight that one copy can send to the node assigned to another copy i, which already retains all but the normalized fraction ε_i of copy i's Ω-inputs. A large transfer would yield a restriction of the form in Claim B.2, which needs more bits than remain at that node.

Proof of Lemma 6.15. Let ε_i be the normalized Ω-weight missing at the node assigned to copy i, and put $\delta_i = 2q_2 x + 4\varepsilon_i$. If $\delta_i \geq 1$, the bound $\mu\delta_i$ is automatic, since the inputs from $\mathcal{D}$ in one copy have weight only $\mu \leq \mu\delta_i$. Otherwise, suppose the transfer exceeds $\mu\delta_i$. Then the node receives a set of inputs labeled by $\mathcal{D}$ from some copy $j \neq i$ of actual weight $\mu\delta$ for some $\delta > \delta_i$.

Both copies have assigned nodes, so their indices lie outside the exceptional set for Y', and only Alice has spoken before reaching the assigned node. By Lemma 6.10 with error bound $\eta/2$ and the display after it, the projection of Y' onto coordinates i and j has original product weight $\geq 1 - \eta/2 > 99/100$. We restrict to the event that both coordinates are template rows. This event has weight λ^2, so the retained relative weight is at least

$$a_{\min} = \frac{99/100 + \lambda^2 - 1}{\lambda^2}.$$

We average over the two bin labels and fix a pair of bins attaining this relative weight. The retained pairs of template labels in these two bins form a set $P_* \subseteq S^2$ of original product weight $\geq a_{\min}$. The bin fixed at copy i, which we call p_i, is the distinguished bin. The projection at coordinate i, which lies outside the exceptional set, loses at most $\eta/2$ of its row weight by the same error bound, while each bin has weight λ/s, so every other bin retains weight at least

$$1 - \frac{\eta s}{2\lambda} > 2sx.$$

We choose one surviving completion in Y' for every pair in P_* and for every retained template label in the other bins. The bin label at i distinguishes the representatives of the other bins from those of the pairs, and the projected labels distinguish the representatives within each bin. Thus the chosen completions define a restriction of the original table with the claimed entries.

After transposition, the columns labeled by Ω from copy i have normalized weight $1 - \varepsilon_i$. The columns labeled by $\mathcal{D}$ from j have J_2-image weight $\delta > \delta_i$, since that image map preserves their

original weights. A pair (z_o, z_f) reads $A(\gamma_{p_i}, (z_o)_{\alpha_{p_i}})$ on columns labeled by Ω and $A(\gamma, (z_f)_\alpha)$ on columns labeled by $\mathcal{D}$. Other bins have the same template entries on the first family, and their entries on the second family are arbitrary. This is the restriction in Claim B.2, after renaming the distinguished bin. The claim requires C bits because $\delta > 2q_2x + 4\varepsilon_i$, whereas the assigned node has only $C-1$ bits left. This contradiction proves that the transfer is at most $\mu\delta_i$, also when $\varepsilon_i = 0$. □

C Conditioning and surviving blocks

We prove Lemma 6.10 by bounding the number of exceptional coordinates, and then show that a short protocol must leave two or three blocks to which Lemma A.2 still applies. For these estimates only, we view a label drawn according to its fixed weights as a random variable. Expectations then express weighted sums over the fixed tuple list. The 1024th-moment bounds use the original weights ν, under which independence is exact, whereas the fourth-moment bound under $\nu \mid Y$ in Lemma C.2 uses only the approximate marginals outside the exceptional coordinates.

C.1 A moment count and the exceptional coordinates

We use the following moment bound twice: to bound the number of exceptional coordinates, and, in the proof of Lemma 6.14, to bound the probability that too few template positions survive. In the second use the coefficients are the fixed fractions d_i of incoming weight, so the bound applies directly to the incoming template weight $\sum_i d_i T_i$.

Lemma C.1 (Moment bound)**.** *Let $r \geq 1$ be an integer. Suppose $X_1, \dots, X_m$ have mean zero, $|X_i| \leq 1$, and every subfamily of at most $2r$ variables is independent. For fixed coefficients $a_i \in [0,1]$ with $W = \sum_i a_i \geq 1$,*

$$\mathbb{E}\left|\sum_i a_i X_i\right|^{2r} \leq (r^2 W)^r.$$

Unit coefficients give the same bound with $W = m$, and the empty sum is zero. In particular, centered indicators with 1024-wise independence satisfy

$$\mathbb{E}\left|\sum_i X_i\right|^{1024} \leq (2^{18}m)^{512}, \qquad \mathbb{E}\left|\sum_i X_i\right|^{256} \leq (2^{14}m)^{128}.$$

Proof. We expand $(\sum_i a_i X_i)^{2r}$ and group the terms by the pattern of equal indices among the $2r$ factors. A term in which some index appears once has expectation zero. Every remaining pattern has $k \leq r$ index classes, each of size at least two. Labeling the classes in order of first appearance shows that there are at most r^{2r} such patterns. If their sizes are $n_1, \dots, n_k$, then, since $|X_i| \leq 1$, the absolute contribution from a fixed pattern is at most

$$\sum_{\substack{i_1,\dots,i_k \\ \text{distinct}}} a_{i_1}^{n_1} \cdots a_{i_k}^{n_k} \leq W^k \leq W^r,$$

because $a_i^{n_j} \leq a_i$. Summing over patterns proves the bound. Taking $r = 512$ and $r = 128$ gives the displayed special cases. □

We take the exceptional coordinates to be the union of a maximal disjoint family of sets of at most eight coordinates with large marginal error. Lemma C.1 bounds the size of this family.

Proof of Lemma 6.10. Put $\varepsilon = \varepsilon(v,\tau)$. We choose a maximal disjoint family $J_1, \dots, J_m$ of sets of at most eight coordinates with $d_{\mathrm{TV}}((\nu \mid Y)_{J_i}, \mu^{J_i}) > \varepsilon$. If $m = 0$, we delete nothing. Otherwise, since ν has marginal μ^{J_i} on J_i, the definition of total variation distance gives an event indicator f_i on each J_i with $\mathbb{E}_{\nu|Y} f_i - \mathbb{E}_\nu f_i > \varepsilon$ (take its complement if necessary). Under the original weights ν, any subfamily of at most 1024 such indicators is independent, because its disjoint sets use at most 8192 coordinates.

Put $S = \sum_i (f_i - \mathbb{E}_\nu f_i)$, so $\mathbb{E}_{\nu|Y} S > m\varepsilon$. Jensen's inequality and Lemma C.1 give

$$v(m\varepsilon)^{1024} < \nu(Y)\mathbb{E}_{\nu|Y}|S|^{1024} \le \mathbb{E}_\nu |S|^{1024} \le (2^{18}m)^{512}.$$

Thus $m < 2^{18}\varepsilon^{-2}v^{-1/512}$. We delete $E = \bigcup_i J_i$, so $|E| \le 8m \le e(v,\tau)$. Every remaining set of at most eight coordinates has marginal within ε of the product, since otherwise it could be added to the disjoint family, contradicting maximality. □

C.2 Some coordinate pair has a large projection

When many blocks survive, we apply the two-block case of Lemma A.2 to a pair of coordinates on which the retained labels project with product weight at least $9/25$. Such a pair exists because, if every one of many disjoint pairs had a smaller projection, the retained set itself would be too small.

Lemma C.2. *Fix disjoint coordinate pairs $J_1, \dots, J_m$ of a tuple list with marginal μ and label weights ν. Suppose a label set K satisfies $\mu^2(\pi_{J_i}K) < 9/25$ for every $i \in [m]$. For an exactly 8-wise independent list,*

$$\nu(K) < \frac{2}{m^2} < \frac{9}{50m} \qquad (m \ge 16).$$

For an exactly 8192-wise independent list, let $K \subseteq Y$ with $\nu(Y) \ge v > 0$, and choose the pairs outside the exceptional set from Lemma 6.10, with any error bound $0 < \tau \le 1$. Then

$$(\nu \mid Y)(K) < \frac{1}{4m} \qquad (m \ge 16).$$

Thus a retained set of conditional weight at least $1/(4m)$ forces one of these pairs to have product projection weight at least $9/25$.

Proof. Write $P_i = \pi_{J_i}K$ and put

$$p_i = \mu^2(P_i) < 9/25, \qquad X_i = \mathbf{1}[z_{J_i} \in P_i] - p_i, \qquad S = \sum_{i=1}^m X_i.$$

On K every indicator equals one, so $S > (16/25)m$. Under the full product weights, the centered indicators are independent, and $\mathbb{E}X_i^2, \mathbb{E}X_i^4 \le 1/4$. Only terms with no singleton index survive the fourth-power expansion. Hence

$$\mathbb{E}_{\mathrm{prod}} S^4 = \sum_i \mathbb{E}X_i^4 + 6\sum_{i<j} \mathbb{E}X_i^2 \mathbb{E}X_j^2 \le \frac{m}{4} + \frac{6}{16}\binom{m}{2} \le \frac{m^2}{4}.$$

Exact 8-wise independence gives the same fourth moment under ν, since each term uses at most four pairs. Markov's inequality for S^4 gives $\nu(K) < 2/m^2$, because $1/[4(16/25)^4] < 2$. For $m \ge 16$, $2/m^2 \le 1/(8m) < 9/(50m)$.

Now suppose the pairs lie outside the exceptional set for the fixed event Y. Every term in S^4 lies in $[-1,1]$ and depends on a set J of at most eight coordinates outside that set. Since ν has

marginal μ^J, Lemma 6.10 shows that conditioning on Y changes the term's expectation by at most 2ε, where $\varepsilon = \varepsilon(v,\tau)$. Hence $\mathbb{E}_{\nu|Y} S^4 \le m^2/4 + 2\varepsilon m^4$, and Markov's inequality gives

$$p := (\nu \mid Y)(K) \le \frac{2}{m^2} + \frac{2\varepsilon}{(16/25)^4} \le \frac{2}{m^2} + 16\varepsilon.$$

The error term 16ε does not decrease with the number of pairs, so for large m we also use the original weights ν. Any subfamily of at most 1024 pair indicators involves at most 2048 coordinates and is therefore independent under ν. Since $K \subseteq Y$ and $\nu(Y) \ge v$, Lemma C.1, Markov's inequality and $(25/16)^2 < 4$ give

$$p = \frac{\nu(K)}{\nu(Y)} \le \frac{(2^{18}m)^{512}}{v((16/25)m)^{1024}} \le \frac{(2^{20}/m)^{512}}{v}.$$

Combining the bounds and using $\min\{u,w\} \le u^{511/512} w^{1/512}$ yields

$$\begin{aligned} mp &\le \frac{2}{m} + \min\{16\varepsilon m, 2^{20\cdot 512} v^{-1} m^{-511}\} \\ &\le \frac{2}{m} + (16\varepsilon)^{511/512} 2^{20} v^{-1/512}. \end{aligned}$$

The last term is at most

$$2^{20-26\cdot 511/512} v^{255/131072} < \frac{1}{32},$$

by $\varepsilon \le 2^{-30} v^{1/256}$ and $0 < v \le 1$. For $m \ge 16$, this gives $mp < 1/8 + 1/32 < 1/4$, as required. □

C.3 Small projections and short continuations

Under $\mathcal{P}(M; x, d)$ with $d \ge 2$, Lemma A.2 with $H = d$ shows that $d+1$ bits are needed for two blocks of row weight $2x$ with product column weight $9/25$, and for three blocks of row weight $4x$ with product column weight $9/50$, whenever this row threshold, $2x$ or $4x$, is at most one. We apply Lemma A.2 to the projections of Observation 4.5, whose column weight $\mu^J(\pi_J K)$ is a product weight even when the labels have conditional weights. The proofs of Lemma 6.11 and Theorem 6.12 follow a protocol for $a-1$ messages and then apply the next lemma, which shows that enough surviving blocks and column weight still force $d+1$ bits.

Lemma C.3. *Suppose $d \ge 2$ and $\mathcal{P}(M; x, d)$, and let $b \ge 0$ be an integer with $2^{b+1}x \le 1$. Consider an interlace of M on a tuple list whose marginal is the column weights of M, and a restriction of it that retains at least $2^b + 1$ blocks, each of base row weight at least $2^{b+1}x$, and a label set K. Assume one of the following.*

(E) The list is exactly 8-wise independent with weights ν, and $\nu(K) \ge (9/25)2^{-b}$.

(C) The list is exactly 8192-wise independent with weights ν, $\nu(Y) \ge v > 0$, the exceptional set is chosen from Lemma 6.10 with error bound $0 < \tau \le 1$, the retained blocks have indices outside that set, $K \subseteq Y$, and $(\nu \mid Y)(K) \ge 2^{-b-1}$.

Then the restriction requires at least $d+1$ bits.

Proof. Let J be any set of at most eight retained block indices. Under (E), exact marginals give $\mu^J(\pi_J K) \ge \nu(K) \ge (9/25)2^{-b}$ (Observation 4.5). Under (C), the indices in J lie outside the exceptional set. With $\varepsilon = \varepsilon(v,\tau) \le 2^{-30}$, Lemma 6.10 gives

$$\mu^J(\pi_J K) \ge (\nu \mid Y)(K) - \varepsilon \ge 2^{-b-1} - \varepsilon.$$

For $b \le 4$ this is again $\ge (9/25)2^{-b}$, since the smallest margin is $(7/50)2^{-4} = 7/800 > 2^{-30}$.

For $b \le 4$ we select the blocks directly. If $b = 0$, any two blocks meet the $9/25$ threshold above, and if $b = 1$, any three meet the $9/50$ threshold. If $b = 2$, the projection onto five retained blocks has product weight $\ge 9/100$, so Lemma A.1 with $l = 2$ gives a pair of weight $\ge (9/100)^{2/5} > 9/25$. If $b = 3$ or $b = 4$, the projection onto eight retained blocks has weight $\ge 9/400$, so Lemma A.1 gives a pair of weight $\ge (9/400)^{1/4} > 9/25$.

For $b \ge 5$, we choose $m = 2^{b-1} \ge 16$ disjoint pairs from the at least $2m + 1$ retained blocks. If every pair projection had product weight less than $9/25$, Lemma C.2 would give $\nu(K) < 9/(50m) = (9/25)2^{-b}$ under (E), and $(\nu \mid Y)(K) < 1/(4m) = 2^{-b-1}$ under (C), a contradiction.

Each selected block has row weight $\ge 2^{b+1}x$, hence $\ge 2x$, and $\ge 4x$ in the triple at $b = 1$. These thresholds are at most one by $2^{b+1}x \le 1$. Thus Lemma A.2 gives $d + 1$ bits in every case. □

Proof of Lemma 6.11 and Theorem 6.12. In Lemma 6.11 we measure column weight under ν, and in Theorem 6.12 under the fixed conditional weights $\nu \mid Y$, where the event Y and its exceptional set stay fixed. Suppose a protocol for the restriction has depth at most $d + a - 1$. We follow it for $a - 1$ messages by a weight-halving walk. At each Alice message, we assign every current block to a child retaining at least half its row weight and keep a child assigned at least half the blocks. At each Bob message, we keep a child retaining at least half the current column weight. If the protocol ends earlier, we pad the leaf with dummy Alice messages, which leave the restriction unchanged. If A messages are Alice's and B are Bob's, then $A + B = a - 1$, at most d bits remain, and the surviving restriction, with label set K', has

$$\begin{aligned} \text{number of blocks} &\ge \lceil (2^{a-1} + 1)/2^A \rceil \ge 2^B + 1, \\ \text{row weight per block} &\ge 2^{a-A}x = 2^{B+1}x, \\ \text{column weight of } K' &\ge 2^{-B} \cdot \text{initial column weight}. \end{aligned}$$

In Theorem 6.12 the surviving blocks still have indices outside the exceptional set. The initial column weight is $\nu(K) \ge 9/25$ in Lemma 6.11 and $(\nu \mid Y)(K) \ge 1/2$ in Theorem 6.12, so K' satisfies hypothesis (E), respectively (C), of Lemma C.3 with $b = B$. Its row threshold is at most one because $B + 1 \le a$ and $2^a x \le 1$. Hence it requires $d + 1$ further bits, while only d remain. □

D Constructing small rational tuple lists

The deterministic reduction prints its tuple lists together with their rational weights, so Lemma 7.1 must bound the encoding length and not only the number of tuples. We use the sequential construction of Koller and Megiddo [15, Section 3]: extend each word by every coordinate value with its product weight, then reduce the list while preserving the joint weights on all sets of at most D coordinates. We recall the procedure and bound the bit length of the weights.

Proof of Lemma 7.1. Let Y be the alphabet, $|Y| = m$, and put $P = (1 + qm)^{\mathsf{D}}$. We start with the empty word of weight one. At step $i \in [q]$, we append every value $y \in Y$ to every retained word and multiply its weight by $\mu(y)$. On the resulting candidate set $\mathcal{U} \subseteq Y^i$, the weights satisfy, by induction on i,

$$\sum_{u \in \mathcal{U}} w_u \mathbf{1}[u_J = a] = \prod_{j \in J} \mu(a_j) \qquad (J \subseteq [i],\ |J| \le \mathsf{D},\ a \in Y^J). \tag{*}$$

There are at most P equalities in (*), including the normalization equation for $J = \varnothing$, since $\sum_{k \le \mathsf{D}} \binom{i}{k} m^k \le (1 + qm)^{\mathsf{D}}$. The constraint column $a(u)$ of a word u has entry $\mathbf{1}[u_J = a]$ in row

(J, a). The reduction below removes linear dependences among these columns until at most P words remain, and it preserves the equalities and the nonnegativity of the weights. Each stage therefore extends at most P words to at most mP candidates. After q stages the equalities give the required D-wise independence.

Words of weight zero are deleted. A reduction step takes a nonzero integer dependence $\sum_u c_u a(u) = 0$, whose row $J = \varnothing$ gives $\sum_u c_u = 0$, so some $c_u > 0$. It sets $w'_u = w_u - tc_u$, where $t = \min_{c_u>0} w_u/c_u > 0$, which keeps the weights nonnegative and (*) valid, and removes at least one word. With more than P words, we take $r \le P$ columns forming a basis of their span, one further column, and r rows on which the basis columns are independent. The signed $r \times r$ minors of this $r \times (r+1)$ zero–one matrix give a nonzero integer relation among these columns, valid on these rows. This relation holds in every row, since the basis spans the further column with coefficients determined by these rows. By Hadamard's bound, each coefficient has absolute value at most $\max\{1, r^{r/2}\}$, hence $O(P \log(P+1))$ bits. Exact linear algebra finds this relation in polynomial time, and fixed orderings of words and rows fix every choice.

Let L_μ be the total binary encoding length of the rationals $\mu(y)$. An extension multiplies the common denominator of the weights by at most the product of the denominators in μ, adding at most L_μ bits. A reduction step multiplies it by the coefficient $c_u > 0$ attaining the minimum ratio, adding $O(P \log(P+1))$ bits. A stage has at most mP reduction steps, each deleting a word. Across q extensions and at most qmP reduction steps, every denominator therefore has bit length

$$O\big(qL_\mu + qmP^2 \log(P+1)\big).$$

Numerators obey the same bound because the weights lie in $[0, 1]$. Hence all arithmetic has polynomial cost for fixed D, and so does the output of at most P weighted words of length q. □

E Printing a sample of the implicit table

The randomized reduction uses the implicit table only through sampled rows, columns and their entries, which we generate without printing the table. We first bound the length and generation cost of complete row and column labels, then implement the elementary draws with fixed blocks of fair bits. Evaluating one entry reads only one base entry per recursion level, but the sampler outputs complete labels, including the coordinates that the evaluation never reads.

Proof of Lemma 9.5. We fix a source graph on n vertices above the size cutoff in Lemma F.2, put $L = \lceil \log n \rceil$, and use the parameter choices of that lemma at every level. Write $R : X \times Y \to \{0, 1\}$ for the current base matrix.

Labels and entries. The two interlaces $U = \mathsf{I}_w(N)$, $W = U^{\mathsf{T}}$ and $R^+ = (\mathsf{I}_Q(W))^{\mathsf{T}}$ give the row and column labels of R^+ (the successor labels)

$$\mathbf{a} = ((i_t, Z_t))_{t\in[Q]}, \qquad \mathbf{b} = (j, (k_i)_{i\in[w]}),$$

where $i_t \in [w]$ and $j \in [Q]$ are uniform, Z_t are tagged columns of B_{raw}, and $k_i \in \mathcal{K}$ are rows of B_{raw}. We draw all these coordinates independently, each according to its weight in Definitions 4.3 and 6.2. The entry rule is

$$R^+(\mathbf{a}, \mathbf{b}) = B_{\mathrm{raw}}(k_{i_j}, Z_j).$$

The entry rules in Definition 6.2 select exactly one entry of R. To generate a row k_i or column Z_t of B_{raw}, we sample according to the row-type and column-family weights from that definition and draw every listed coordinate, including those not visited by the eventual entry evaluation. A

$\mathcal{D}$-input needs one draw of $j = (\alpha, \gamma) \in J_2$, since its other positions copy j or tail(γ). We generate only actual input coordinates, not the auxiliary sequences used only in the lower-bound proof. The numbers of rows and columns of R drawn for each label are

label	base-row draws	base-column draws
template row	q_2	0
row for an item	1	0
Ω column	0	$(s+1) + s(2s-1) = 2s^2 + 1$
$\mathcal{D}$ column	0	$2s - 1$
successor row	0	$\le Q(2s^2+1)$
successor column	$\le wq_2$	0

Apart from these recursive draws, generating indices and scanning activity lists cost polynomial time in $n + s + q_2 + w + Q + L + 1$, copying a label costs time linear in its length, and no coloring is required. Every width s, q_1, q_2, w, Q is $2^{O(L)}$ by Lemma F.2, and so is every draw count in the table. Hence we can choose common computable integer bounds J_i for the label length, generation time, elementary draw count and entry-evaluation time at level i such that

$$J_0 = O(L), \qquad J_{i+1} \le 2^{c_0 L} J_i, \qquad J_k = 2^{O((k+1)L)},$$

where c_0 is fixed and effective. Here each elementary draw counts as one operation, and we implement it with fair bits below. At the initial matrix R_0 of Lemma 5.4, a row consists of $p_0 = O(L)$ independent uniform bits, a column is a uniform index in $[p_0]$, and an entry can be read in $O(L)$ time. This gives $J_0 = O(L)$.

Since $h = O(L)$ and $d_0 = O(L)$ by Lemma F.2, we have $B_k = O((k+1)L)$. The direct sample size in Theorem 9.4 then gives $m = 2^{O((k+1)L)}$. We generate and store the m row labels and m column labels, then evaluate all m^2 row–column entries. Including the stored tuples and printed entries, this costs

$$O((2m + m^2)J_k) = 2^{O((k+1)L)} = n^{O(k+1)}.$$

All constants are independent of k.

Fixed-bit implementation. Let $N_{\mathrm{draw}} \ge 1$ be a computable bound on the elementary draws for all $2m$ labels, so $N_{\mathrm{draw}} = 2^{O((k+1)L)}$. On every tape, each draw has at most

$$A = \max\{512, 4s(n+s), q_2, w, Q, p_0\} = 2^{O(L)}$$

outcomes, and the weights of each draw have a common denominator $\le A$. This holds for type draws (whose weights in Definition 6.2 have denominators 512 and 256), item indices, bins, template positions, activity indices, interlace-copy indices and draws from the initial matrix. Rows and columns of base matrices above R_0 are generated recursively, not drawn as single outcomes. We set $b = \lceil \log(200 A N_{\mathrm{draw}}) \rceil$ and reserve N_{draw} blocks of b fair bits. A uniform draw from $[a]$, $a \le A$, reads the next block as a uniform integer $U \in \{0, \dots, 2^b - 1\}$ and returns

$$\mathrm{Draw}(a) = 1 + \left\lfloor aU/2^b \right\rfloor. \tag{7}$$

A draw whose weights have common denominator $a \le A$ maps consecutive intervals of values of Draw(a), whose lengths are the numerators of its weights, to its types. Unused blocks are ignored.

Each value in $[a]$ has $\lfloor 2^b/a \rfloor$ or $\lceil 2^b/a \rceil$ preimages, so a draw is within total variation $a/2^b \le 1/(200 N_{\mathrm{draw}})$ of the ideal one, and mapping values to types cannot increase this distance. Because every draw reads a fresh block, this bound holds conditionally on all earlier outcomes, even when

the alphabet depends on them. We couple the implemented and ideal draws one at a time. By a union bound over the at most N_{draw} draws, the implemented and ideal draws differ somewhere with probability at most $N_{\mathrm{draw}}/(200N_{\mathrm{draw}}) = 1/200 < 1/100$.

Every tape thus produces exactly m legal row labels and m legal column labels, and all m^2 entries are evaluated exactly. The tape has fixed length bN_{draw}, where $b = O((k+1)L)$. Arithmetic on $O(b+\log A)$-bit integers adds only polynomial overhead. Hence the running time and the number of random bits are $n^{O(k+1)}$ on every tape. □

F One parameter recipe for both constructions

We first verify the numerical requirements of the recipe in Definition 6.1 that both constructions share. For full products and for explicit lists, we then check the source-matrix deletion bounds and the final interlace.

Recall $n = |V|$, $\lambda = 255/256$, $\mu = 1/256$, and $\zeta = \gamma/24$. The source promise gives, with δ_v as in Lemma 6.3,

$$n/2 \le |E| \le \Delta n/2, \qquad \delta_v = \frac{\gamma|E|}{2\Delta n} \ge \frac{\gamma}{4\Delta}.$$

From Lemma 6.13 we use the bound E_{exc} on the sum of κ over stopping nodes with no assigned copy, and from the proof of Lemma 6.14 the weighted moment constant C_{128}:

$$E_{\mathrm{exc}} = 2e(\theta, \eta/2), \qquad C_{128} = \left(\frac{2^{14}}{(\lambda - 1/2)^2}\right)^{128}. \tag{8}$$

The definition of $e(v,\tau)$ before Lemma 6.10 gives the estimate

$$\begin{aligned} e(v,\tau) &\le 8 + 2^{23}\tau^{-2}v^{-1/512} + 2^{83}v^{-5/512} \\ &\le 2^{84}\tau^{-2}v^{-1/64} \qquad (0 < v, \tau \le 1). \end{aligned} \tag{9}$$

Lemma F.1 (Common numerical bounds)**.** *There is an effective dyadic constant s_0, depending only on the fixed graph promise, such that the recipe in Definition 6.1 with $n \le \rho^{-1}$, $\rho \le s^{-4}$, and $s \ge s_0$ satisfies all source and internal-interlace bounds listed in* (10) *below. The choice of s_0 is independent of the base matrix and the number of rounds.*

Proof. The middle column of the table below lists upper bounds on the quantities in the first column, obtained by substituting the recipe. Each of these bounds tends to zero as s increases. The last column lists the bound that the proof requires for each quantity.

quantity	upper bound	required allowance
$s\eta$	s^{-1}	$2^{-16}\min\{1, \delta_v, \mu\zeta\}$
$4s\lvert E\rvert/q_2$	$2\Delta/s$	1
$s^2\rho$	s^{-2}	$1/4$
sq_1/q_2	s^{-4}	$\zeta/128$
$\max\{q_2x, sx, wx\}$	s^{-24}	10^{-4}
w^2x/η	s^{-14}	$1/(32\lambda)$
w^2q_2x	s^{-10}	1
wq_2x/η	s^{-16}	$1/(128\mu)$
$e(\theta/4)/(sq_2)$	$2^{85}s^{-53/8}$	1
$C_{128}s^{-128}/\theta$	$512C_{128}s^{-104}$	$1/4$
$2sE_{\mathrm{exc}}/(\eta w)$	$2^{89}s^{-5/8}$	$1/16$

(10)

For the two rows involving e, where $e(v) = e(v,1)$, (9) with $\theta = 2^{-9}s^{-24}$ and $\eta = s^{-2}$ gives $e(\theta/4) \le 2^{85}s^{3/8}$ and $E_{\mathrm{exc}} \le 2^{88}s^{35/8}$, while $sq_2 = s^3/\rho \ge s^7$. The other rows follow by substituting the widths of Definition 6.1 and using $|E| \le \Delta n/2$, $n \le \rho^{-1}$, and $\rho \le s^{-4}$. We choose a dyadic $s_0 \ge 256$ so that every middle entry is below its allowance, using $\gamma/(4\Delta)$ in place of δ_v. Clearing the fixed rational exponents makes this a finite set of integer-power comparisons. Each middle entry decreases with s, so the same choice works for every $s \ge s_0$. None of these comparisons involves the base matrix or the number of rounds.

The table's second row gives $q_2 \ge 4s|E|$. Since q_2 is a power of two, $q_2 \ge q_0 = \mathrm{pow}_2(4s|E|)$, as Definition 6.2 requires. In particular, $|E| \ge n/2$ gives $q_2 \ge 2sn \ge 2(n+s-3)$, so all item labels fit: $|\mathcal{I}| = 4s(n+s-3) \le 2sq_2$, as Lemma 6.18 requires. The table's last bound also implies the two bounds used to count assigned nodes:

$$E_{\mathrm{exc}} \le \frac{\eta w}{32s} < \frac{w}{4}, \qquad e(\theta) \le e(\theta,\eta/2) = \frac{E_{\mathrm{exc}}}{2} \le \frac{\eta w}{64s} < \frac{w}{128}.$$

Here $e(\theta) = e(\theta,1) \le e(\theta,\eta/2)$ because $e(v,\tau)$ is nonincreasing in the error bound τ.

The remaining checks of fixed fractions appear in the proofs that use them. For each such proof, we list the table rows it uses and derive here only the bounds that it does not display.

- *Block counts* (proof of Lemma 6.6). Since $s \ge 256$ and $s\rho < 10^{-3}$, the three lower bounds for m/u satisfy

  $$\frac{3}{5}(1-2/s) - q_1\rho > \frac{9}{16}, \qquad \frac{3}{5} - q_2x > \frac{9}{16}, \qquad \frac{44}{75} - sx > \frac{9}{16},$$

  the last since $2^j(\beta_j - (1-\lambda))/\lambda \ge ((3/5) - 4(1-\lambda))/\lambda = 44/75$ for $j \in \{0,1\}$. The bound on m/κ in the proof of Theorem 6.16 uses $wx \le 1/256$ and $4e(\theta)/w < 1/32$ from the display above.

- *Item deletion and collisions* (proof of Theorem 6.9). The row $s\eta$ gives $4096\eta sn \le \delta_v n/16$ and $4096\eta s(s-3) \le (s-3)/16 < s/2$. The rows $s\eta$ and sq_1/q_2 give $8sq_1 + 4s\eta q_2/\mu \le (1/16 + 2^{-14})\zeta q_2 < \zeta q_2/4$.

- *Row and template thresholds* (proofs of Theorem 6.9, Lemma 6.13, and Claim B.2). The rows $s^2\rho$, $s\eta$ and $\max\{q_2x, sx, wx\}$ give $2s\rho \le 1/8$, $2wx + \eta/(16\lambda) < 1$, $2sx \le 10^{-3}$, and hence $999/1000 - 2(s-1)x > 998/1000$ in the proof of Theorem 6.9 and $(99/100 + \lambda^2 - 1)/\lambda^2 - 10^{-3} > 49/64$ in Claim B.2.

- *Conservation* (proof of Theorem 6.16). That proof uses the rows w^2x/η, wq_2x/η and $2sE_{\mathrm{exc}}/(\eta w)$ and displays the resulting bounds.

- *Moment bound* (proof of Lemma 6.14). That proof bounds the tail estimate $C_{128}D_0^{-128}$ obtained from Lemma C.1 by $\theta/4$, using $D_0 \ge 512s\kappa \ge 512s$ and the row $C_{128}s^{-128}/\theta$.

□

F.1 Full products

In the randomized reduction, the graph-encoding width $q_2 = s^2 2^L$ is the only width that contributes a term linear in L to the per-round overhead h. The widths s, $w = s^8$, and $Q = s^{24}$ are polynomial in $\log n$. A smaller overhead allows more amplification rounds before the final sample becomes too large.

Lemma F.2 (Parameters for full products)**.** *Put $L = \lceil \log n \rceil$, $\rho = 2^{-L}$, and $s = \mathrm{pow}_2(2^{16}(L+1))$, and use the recipe of Definition 6.1. For every sufficiently large source graph, these choices satisfy the hypotheses of Lemma F.1, the numerical conditions for the full-product deletion bounds for the source matrices, and the numerical hypotheses of Theorem 6.17 with $b = L$ and $\delta = \delta_U$. The per-round overhead is*

$$h = 2\log s + \log q_2 + \log w + \log Q = L + 36\log s = L + O(\log L).$$

There is an effective fixed cutoff beyond which also $L \geq 16$, $d_0 \leq L$, and $h \leq 7L$, where d_0 is the cost of the initial matrix in Lemma 5.4. Computing the widths has polynomial bit complexity and uses neither the base matrix nor the number of rounds.

Proof. Since $s = \mathrm{pow}_2(2^{16}(L+1))$, we have $2^{16}(L+1) \leq s < 2^{17}(L+1)$. Hence $s \geq s_0$ and $2^{-L} \leq s^{-4}$ beyond an effective fixed cutoff. Since also $n \leq 2^L = \rho^{-1}$, the common lemma applies. In each application of the deletion bounds for the source matrices, the block parameter satisfies $u \geq s/2$. Since $\log(1/x) = 8L$,

$$u \geq 2^{15}(L+1) > 4096\log(1/x),$$

which gives $x > 2^{-u/4096}$ and hence the block threshold $x \geq 2^{-u/1024}$ of Lemma 6.5. Lemma F.1 gives the block counts and row-weight bounds required in the proof of Lemma 6.6. The remaining hypotheses of Theorem 6.17 for the final interlace are

$$0 < \delta_U = \frac{1}{16s^{10}} < \frac{1}{64}, \qquad Q\theta = \frac{1}{512} < \frac{1}{20}, \qquad Q\delta_U = \frac{s^{14}}{16} \geq 80L.$$

Thus Theorem 6.17 applies with $b = L$.

Substituting the widths gives the displayed value of h. Writing $\ell = \lceil \log(L+1) \rceil$, the initial matrix of Lemma 5.4 has $d_0 = \log p_0 + 1 = \ell + 7$ and $\log s = 16 + \ell$, so $d_0 \leq L$ and $h \leq 7L$ eventually. Since $s = \Theta(L+1)$ with explicit constants, the ratios $s^4 2^{-L}$, d_0/L and $(h-L)/L$ are at most fixed multiples of $(L+1)^4 2^{-L}$ or $(1+\log(L+1))/L$. These bounds decrease to zero beyond a computable point. Hence the cutoff is effective and valid at every larger size. All widths except $q_1 = s - 2$ are powers of two with exponents $O(L)$; forming those powers and subtracting two from s gives the stated bit complexity. □

F.2 Explicit lists

For the deterministic reduction, we choose the least dyadic $s \geq s_0$ with $s^4 \geq n$ and set $\rho = s^{-4}$. Then every width is polynomial in s, including the final-interlace width $Q = s^{24}$ required by Theorem 7.2.

Lemma F.3 (Parameters for explicit lists)**.** *After increasing the effective dyadic constant s_0 if necessary, every dyadic $s \geq \max\{n^{1/4}, s_0\}$ satisfies the bounds below. Choose the least such s and put $\mathsf{N} = s^4$, $L_{\mathrm{g}} = \log \mathsf{N}$, and $\rho = \mathsf{N}^{-1}$. With these choices, the common recipe satisfies all numerical hypotheses of the explicit-list construction and of the final interlace in Theorem 7.2. Its overhead is*

$$h_{\mathrm{g}} = 2\log s + \log q_2 + \log w + \log Q = 40\log s = 10L_{\mathrm{g}}.$$

With $\mathsf{N}_0 = s_0^4$, the padding satisfies $n \leq \mathsf{N} < 16\max\{n, \mathsf{N}_0\}$. Every width is polynomial in N. One round has a fixed polynomial cost in N and the complete encoding of the base matrix, with exponent independent of the round number.

Proof. Here $\rho = s^{-4}$, so the common lemma applies. Only the deletion bounds for the source matrices at tuple weight x and the final interlace need additional checks. By (9),

$$e(x) \le 2^{84} s^{1/2} \le s/128, \qquad e(\rho, \delta_U/2) \le 2^{94} s^{20+1/16}.$$

We increase s_0 so that $2^{84} s^{1/2} \le s/128$ and $128(2^{94} s^{20+1/16} + 1) \le s^{24} = Q$. Dividing each condition by its right side leaves fixed multiples of negative powers of s, so this enlargement is effective and valid for every larger s.

For the deletion bounds for the source matrices, $u \ge s/2$, and by the three block-count bounds in the proof of Lemma 6.6, more than $9u/16$ blocks meet the required row-weight threshold. Discarding the blocks at the at most $e(x) \le s/128 \le u/32$ exceptional coordinates leaves more than $9u/16 - u/32 > u/2$ blocks. The row thresholds obey $u\rho \le s\rho \le 1$ for the first matrix and $ux \le \max\{q_2 x, sx\} \le 1$ in the other two. For the final interlace, $0 < \rho \le 1$, $0 < \delta_U < 1/64$, $Q\theta = 1/512 < 1/20$, and the bound $Q \ge 128(e(\rho, \delta_U/2) + 1)$ above are the numerical hypotheses of Theorem 7.2 with $\phi = \rho$ and $\delta = \delta_U$.

The least dyadic choice satisfies $s < 2\max\{n^{1/4}, s_0\}$, proving the factor-16 size bound. This s can be found by comparing integer fourth powers, without approximating a root. The widths are $q_2 = s^6$, $w = s^8$, and $Q = s^{24}$, giving the stated overhead. The independence order stays fixed at 8192, so by Lemma 7.1 each of the fixed number of tuple-list constructions has cost polynomial in these widths and in the complete rational encoding of the base matrix. Neither the parameter choices nor the exponent of the per-round polynomial cost depends on the number of rounds. □

G Coloring promises from SAT

The interlacing reduction uses bounded-degree graph coloring (Theorem 5.2), and the strict-alternation reduction uses balanced hypergraph coloring (Proposition 10.2). This appendix derives both hardness results from the literature and checks the input conventions that the communication constructions need.

G.1 Bounded-degree graph coloring of almost-linear size

The graph source of Theorem 5.2 needs a constant coloring gap and must keep the almost-linear size of sparse SAT. Composing the reductions of Dinur and of Bogdanov, Obata and Trevisan gives both. We record this composition and our graph conventions because the size bound determines the final additive-error bound.

Proof of Theorem 5.2. For a SAT instance of encoding length L, Dinur's almost-linear PCP construction gives a constant-alphabet binary constraint graph of size $L(\log L)^{O(1)}$, with perfect completeness and a constant unsatisfied-constraint fraction in the NO case [5, Lemma 8.3 and proof of Theorem 8.1]. Dinur's degree reduction makes the degree constant, preserves perfect completeness, and changes the size and the NO-case unsatisfied fraction by only constant factors [5, Lemma 1.7]. A constant-size Boolean encoding of the labels and constraints, followed by the reduction of Bogdanov, Obata, and Trevisan [3, Lemma 11], gives bounded-occurrence 3-SAT with perfect completeness, a constant NO-case gap and almost-linear size. The gap-preserving reduction of Bogdanov, Obata, and Trevisan [3, Theorem 14 and its proof, pp. 98–99] then gives a bounded-degree graph of linear size.

It remains to match the conventions of Theorem 5.2. We remove isolated vertices and add a disjoint edge, so that the graph is nonempty. This preserves completeness and, when the graph already has an edge, decreases the NO-case gap by at most a factor of two. For trivial inputs

we use fixed promise graphs. For a sparse formula on N variables, $L = O(N \log(N+2))$, so the resulting graph H has $N(\log(N+2))^{O(1)}$ vertices and edges. Taking $\max\{1, \lceil N/|E(H)| \rceil\}$ disjoint copies gives $N \le |E(G)| \le N + |E(H)|$ without changing the degree or gap. The remaining bounds $n/2 \le |E(G)| \le \Delta n/2$ follow from the absence of isolated vertices and the maximum degree Δ. Parallel edges are retained, and edge occurrences are counted with multiplicity. We choose a rational $\gamma > 0$ below the resulting gap and enlarge the fixed exponent p to cover the constant factors. The cited reductions and these adjustments are deterministic and polynomial-time. □

G.2 Balanced hypergraph coloring with distinct vertices

The hypergraph promise of Proposition 10.2 comes from the hardness result of Guruswami and Lee [8], whose edges may contain repeated vertices. We replace repeated occurrences by distinct labeled clones and check that the YES and NO conditions of Proposition 10.2 are preserved.

Proof of Proposition 10.2. Recall the fixed parameters of Proposition 10.2

$$s_{\mathrm{a}} = 2^t, \qquad K_{\mathrm{a}} = s_{\mathrm{a}}^2, \qquad S_{\mathrm{a}} = 2^{s_{\mathrm{a}}}.$$

By Theorem 1.1 of Guruswami and Lee [8], it is NP-hard to distinguish K_{a}-uniform occurrence hypergraphs admitting an s_{a}-coloring with at least $s_{\mathrm{a}} - 1$ occurrences of each color in every edge from those with no independent set containing at least a $1/(2S_{\mathrm{a}})$ fraction of the vertices. Here, as in their reduction, an edge of an occurrence hypergraph is a family of K_{a} vertex occurrences indexed by $[s_{\mathrm{a}}] \times [s_{\mathrm{a}}]$, listed in a fixed order. A vertex may occur more than once in an edge, and colors in an edge are counted with multiplicity, as in their completeness proof. Hypergraphs of the second kind admit no proper S_{a}-coloring, since some color class would contain at least a $1/S_{\mathrm{a}}$ fraction of the vertices and so would not be independent.

We replace each vertex u by $K_{\mathrm{a}}S_{\mathrm{a}}$ labeled clones. For each ordered occurrence edge $(u_1, \dots, u_{K_{\mathrm{a}}})$, we include every tuple

$$(u_1^{(j_1)}, \dots, u_{K_{\mathrm{a}}}^{(j_{K_{\mathrm{a}}})})$$

of pairwise distinct clones, in the displayed order, as an edge. If an occurrence edge appears several times, we treat its copies as separately labeled edges. Each original edge produces at least one such tuple, since each vertex has $K_{\mathrm{a}}S_{\mathrm{a}} \ge K_{\mathrm{a}}$ clones. Giving clones their source vertex's color preserves the color multiset of every edge, and hence the YES promise.

Conversely, suppose the cloned hypergraph has a proper S_{a}-coloring. We give each source vertex a color appearing on at least K_{a} of its clones. Suppose some occurrence edge were monochromatic in this induced coloring. Each source vertex occurs at most K_{a} times in that edge, so we can choose different clones of that color for its positions. This gives a monochromatic edge of distinct clones, contradicting properness. Hence the induced coloring is proper, contradicting the NO promise.

Each vertex produces $K_{\mathrm{a}}S_{\mathrm{a}}$ clones and each edge produces at most $(K_{\mathrm{a}}S_{\mathrm{a}})^{K_{\mathrm{a}}}$ cloned edges. These are constants for fixed t, so enumeration is deterministic and polynomial-time, although the polynomial may depend on t. If the vertex set is empty, we add one isolated labeled vertex, since Proposition 10.2 requires a nonempty vertex set. Because each ordered edge has distinct vertices, the position $\iota_\alpha(v)$ in Definition 10.3 is well defined. □

H Finite-field arrays

The alternating lower bound needs a short, explicitly generated list of arrays on which no prescribed bit pattern occurs much more often than it would for independent fair bits (Equation (6)). We construct that list here, completing the algorithmic step promised in Corollary 10.6.

We use the polynomial-root construction of Alon, Goldreich, Håstad and Peralta [1, Construction 3 and Proposition 3]. The conversion from parity bounds to pattern frequencies is also recorded in that paper, with credit to Vazirani [1, Lemma 1 and Corollary 1]. We use the following direct consequence and keep its seed labels explicit so that the reduction can print the full array list.

Lemma H.1 (Construction 3, Proposition 3 and Corollary 1 in [1])**.** *For integers $n, r \geq 1$, a deterministic algorithm prints a nonempty labeled list $(w^\ell)_{\ell\in\mathcal{L}}$ of n-bit arrays satisfying*

$$\frac{|\{\ell \in \mathcal{L} : w^\ell|_I = z\}|}{|\mathcal{L}|} \leq 2^{-|I|} + 2^{-r} \qquad (I \subseteq \{0,\dots,n-1\},\ z \in \{0,1\}^I).$$

Its running time and complete output length are polynomial in $n2^r$.

To instantiate the construction, let R be the least power of two with $R \geq n2^r$. The construction lists $R^2 < 4n^2 4^r$ arrays, indexed by $\mathbb{F}_2^{\log R} \times \mathbb{F}_R$, whose parities over nonempty coordinate sets have bias at most $(n-1)/R \leq 2^{-r}$, and the bias-to-pattern conversion gives the displayed bound. A field representation of $\mathbb{F}_R$ can be found by exhaustive search in time polynomial in R, and we keep repeated arrays as distinct labels.

Finally, we apply Lemma H.1 on the positions Ω_{a} of Definition 10.3 with error ϵ_{a}, and check that the list has size polynomial in H_{a} for each fixed t.

Proof of Corollary 10.6. We number Ω_{a} by $0,\dots,n-1$ and apply Lemma H.1 with

$$n = K_{\mathrm{a}} H_{\mathrm{a}}, \qquad r = 4(b+3t).$$

Here $2^{-r} = \epsilon_{\mathrm{a}}$ and $n2^r = K_{\mathrm{a}} 2^{12t} H_{\mathrm{a}}^5$, which is polynomial in H_{a} for fixed t. The integers n and r have $O_t(b+1)$ bits, and the polynomial bound also covers $H_{\mathrm{a}} = 1$. □

AI methodology

The original NP-hardness result and the subsequent ETH-based inapproximability result were obtained manually by the authors, who developed the underlying constructions and mathematical insights. The authors also initiated the effort to sparsify interlacing while retaining its communication lower bounds. AI assistance contributed to the later development: completing proof details, finding and checking further arguments, extending the results to strict alternation, simplifying the proofs, and formalizing them in Lean.

We used OpenAI GPT models through ChatGPT Pro and Codex, together with Anthropic's Claude models, to develop and check arguments, search for counterexamples, and compare proposed proofs with the literature. GPT and Claude also helped refine the writing and simplify the proofs by removing redundant cases, combining shared arguments, and reducing parameter bookkeeping. The inputs to this work included the authors' manuscripts, specific mathematical questions, and proposed arguments.

The final sparsification approach was developed with AI assistance, including the connection to rank-sensitive counting: bounding the number of different column sets on which a short protocol can be correct in terms of the matrix rank. The counting ingredient used here builds on the row-sum argument of Oum, Sæther and Vatshelle [20, Theorem 4.2], in its protocol-specific form given by Mackenzie [17, Lemma 4.3]; related counting and covering methods for hard submatrices appear in Hrubeš [11]. AI helped connect these existing methods to the authors' sparsification objective and develop their application to the reduction. The contribution here is their use to retain the required communication lower bound in a much smaller table.

AI assistance also contributed to amplification followed by a final sampling step, to the bounded-independence and conditioning arguments that keep the deterministic construction small, and to the separate construction for strictly alternating protocols. The latter uses the classical small-bias construction of Alon, Goldreich, Håstad and Peralta [1]. ChatGPT Pro supplied substantive mathematical proposals; Codex helped complete their proofs and check the parameter choices and running-time bounds; Claude contributed mathematical and editorial review.

The mathematical and editorial work was followed by AI-assisted formalization in Lean. The companion repository [13] provides the machine-checked statements and proofs, together with build instructions. The interactive proof inspector [14] connects the manuscript to those proofs and lets readers examine their definitions, hypotheses, and dependencies. Cited source reductions remain explicit hypotheses of the formal statements. The authors retain responsibility for the mathematical claims, the use and attribution of prior work, and the final text.

References

[1] Noga Alon, Oded Goldreich, Johan Håstad, and René Peralta. Simple constructions of almost k-wise independent random variables. *Random Structures & Algorithms*, 3(3):289–304, 1992. doi: 10.1002/rsa.3240030308. URL `https://web.math.princeton.edu/~nalon/PDFS/aghp4.pdf`. Our theorem numbering refers to the author-hosted version dated 22 February 2002: Lemma 1, Corollary 1, Construction 3 and Proposition 3.

[2] Beyond the Library. Formal proof inspector for Refuting the Direct Sum Conjecture. FormaliaXiv, 2026. URL `https://beyondthelibrary.github.io/formal_archive/p/Refuting/`. Accessed 25 October 2026.

[3] Andrej Bogdanov, Kenji Obata, and Luca Trevisan. A lower bound for testing 3-colorability in bounded-degree graphs. In *Proceedings of the 43rd Annual IEEE Symposium on Foundations of Computer Science*, pages 93–102. IEEE Computer Society, 2002. doi: 10.1109/SFCS.2002.1181886. URL `https://lucatrevisan.github.io/pubs/bot.pdf`.

[4] Binh-Minh Bui-Xuan, Jan Arne Telle, and Martin Vatshelle. H-join decomposable graphs and algorithms with runtime single exponential in rankwidth. *Discrete Applied Mathematics*, 158 (7):809–819, 2010. doi: 10.1016/j.dam.2009.09.009. URL `https://perso.lip6.fr/binh-minh.bui-xuan/files/BTV.pdf`. Our Proposition 3.6 locator follows the author-hosted manuscript.

[5] Irit Dinur. The PCP theorem by gap amplification. *Journal of the ACM*, 54(3), June 2007. doi: 10.1145/1236457.1236459. URL `https://www.wisdom.weizmann.ac.il/~dinuri/mypapers/combpcp.pdf`. Article 12. Our theorem numbering refers to the author manuscript dated 13 February 2007.

[6] Serge Gaspers, Tao Zixu He, and Simon Mackenzie. NP-completeness of deterministic communication complexity via relaxed interlacing, 2026. URL `https://arxiv.org/abs/2508.05597v4`. Earlier version of the present paper, version 4, 13 April 2026.

[7] Serge Gaspers, Tao Zixu He, and Simon Mackenzie. NP-completeness and ETH-based additive inapproximability of deterministic communication complexity via relaxed interlacing, 2026. URL `https://arxiv.org/abs/2508.05597v5`. Earlier ETH stage of the present paper, version 5, 22 October 2026.

[8] Venkatesan Guruswami and Euiwoong Lee. Strong inapproximability results on balanced rainbow-colorable hypergraphs. *Combinatorica*, 38(3):547–599, 2018. doi: 10.1007/s00493-016-3383-0. URL https://eccc.weizmann.ac.il/report/2014/043/. Theorem numbering in our source discussion refers to ECCC TR14-043.

[9] Shuichi Hirahara, Rahul Ilango, and Bruno Loff. Hardness of constant-round communication complexity. In *36th Computational Complexity Conference (CCC 2021)*, volume 200 of *Leibniz International Proceedings in Informatics (LIPIcs)*, pages 31:1–31:30, 2021. doi: 10.4230/LIPIcs.CCC.2021.31. URL https://drops.dagstuhl.de/entities/document/10.4230/LIPIcs.CCC.2021.31.

[10] Shuichi Hirahara, Rahul Ilango, and Bruno Loff. Communication complexity is NP-hard, 2025. URL https://arxiv.org/abs/2507.10426v1. Preprint, version 1, 14 July 2025.

[11] Pavel Hrubeš. Hard submatrices for non-negative rank and communication complexity. *Computational Complexity*, 34, 2025. doi: 10.1007/s00037-025-00269-4. URL https://doi.org/10.1007/s00037-025-00269-4. Article 9.

[12] Rahul Ilango, Bruno Loff, and Igor Carboni Oliveira. NP-hardness of circuit minimization for multi-output functions. In *35th Computational Complexity Conference (CCC 2020)*, volume 169 of *Leibniz International Proceedings in Informatics (LIPIcs)*, pages 22:1–22:36, 2020. doi: 10.4230/LIPIcs.CCC.2020.22. URL https://drops.dagstuhl.de/entities/document/10.4230/LIPIcs.CCC.2020.22.

[13] Interlacing additive Lean. Interlacing and additive inapproximability: Lean companion. GitHub repository, 2026. URL https://github.com/SimonWMackenzie/interlacing-additive-lean. Companion formalization and proof resources for this paper.

[14] Interlacing additive proof inspector. Interlacing and additive inapproximability: Paper and Lean proof inspector. Interactive proof website, 2026. URL https://simonwmackenzie.github.io/interlacing-additive-lean/inspector/. Companion formalization and proof resources for this paper.

[15] Daphne Koller and Nimrod Megiddo. Constructing small sample spaces satisfying given constraints. *SIAM Journal on Discrete Mathematics*, 7(2):260–274, 1994. doi: 10.1137/S0895480192228140. URL https://theory.stanford.edu/~megiddo/pdf/stoc93.pdf. First appeared in STOC 1993. Algorithm and theorem numbering here refers to the author-hosted conference version.

[16] Eyal Kushilevitz and Enav Weinreb. On the complexity of communication complexity. In *Proceedings of the 41st Annual ACM Symposium on Theory of Computing*, pages 465–474. Association for Computing Machinery, 2009. doi: 10.1145/1536414.1536479. URL https://doi.org/10.1145/1536414.1536479.

[17] Simon Mackenzie. Lossless hardness condensation in deterministic communication complexity, 2026. URL https://arxiv.org/abs/2609.28691v1. Preprint, version 1, 23 September 2026.

[18] Simon Mackenzie and Abdallah Saffidine. Refuting the direct sum conjecture for total functions in deterministic communication complexity, 2026. URL https://arxiv.org/abs/2411.19003v2. Corrected full version, 14 April 2026; conference version appeared in STOC 2025.

[19] Noam Nisan and Avi Wigderson. Rounds in communication complexity revisited. *SIAM Journal on Computing*, 22(1):211–219, 1993. doi: 10.1137/0222016. URL `https://doi.org/10.1137/0222016`.

[20] Sang-il Oum, Sigve Hortemo Sæther, and Martin Vatshelle. Faster algorithms for vertex partitioning problems parameterized by clique-width. *Theoretical Computer Science*, 535: 16–24, 2014. doi: 10.1016/j.tcs.2014.03.024. URL `https://arxiv.org/abs/1311.0224v2`. Our theorem numbering refers to arXiv version 2.

[21] Arshia Soltani Moakhar, Iman Gholami, Max Springer, Mahdi JafariRaviz, and Mohammad-Taghi Hajiaghayi. Beyond the Library: An agentic framework for autoformalizing research mathematics, 2026. URL `https://arxiv.org/abs/2606.31134v3`. Version 3, 5 August 2026; direct-sum formalization in Section 3.2.2 and Appendix C.3.

[22] Andrew Chi-Chih Yao. Some complexity questions related to distributive computing (preliminary report). In *Proceedings of the Eleventh Annual ACM Symposium on Theory of Computing*, pages 209–213. Association for Computing Machinery, 1979. doi: 10.1145/800135.804414. URL `https://doi.org/10.1145/800135.804414`.